

PML-AI Body of Knowledge

The reference for the PCI Project Management Leader – AI
Leadership · delivery systems · governance · the governed use
of AI

FIRST EDITION — PHASE 1 PROTOTYPE (NOT FOR RELEASE)

PROJECT CONTROLS INSTITUTE GLOBAL

Finance intelligently. Control predictively. Deliver successfully.

AI proposes; the professional verifies, decides and remains accountable.

— Prototype notice

This document is a **Phase 1 production prototype** of one domain, built to validate the book's pattern, calculation-verification harness, figures and typesetting pipeline before scaled authorship. It is not a released PCI publication and carries no entitlement, syllabus or examination status. The released edition will carry the full copyright, disclaimer and trademark notices of the PCI book family; all examples here are original and fictional, all calculations pass the golden-answer suite ([verify_formulas.py](#)), and no standard's text is reproduced.

Contents

The Project Leadership Profession	7
Why this domain exists	7
Knowledge Area 1.1 Projects, programmes and portfolios	8
Knowledge Area 1.2 The project leader's accountability	10
Knowledge Area 1.3 Systems thinking and value	13
Knowledge Area 1.4 Ethics and the responsible use of AI	16
Advanced topics — Domain 1	19
Industry variations — Domain 1	19
Case study — Domain 1: Meridian under scrutiny (public health)	20
Case study B — Domain 1: the plan nobody could critique (financial services)	20
Executive perspective — Domain 1	21
Calculation exercises — Domain 1	21
Practitioner's toolkit — Domain 1	21
Exam preparation — Domain 1	22
Domain 1 summary	22
Strategy, Selection and Business Alignment	24
Why this domain exists	24
Knowledge Area 2.1 Strategic context and alignment	25
Knowledge Area 2.2 The business case and selection	27
Knowledge Area 2.3 Benefits, value and sustainability	32
Knowledge Area 2.4 Strategic termination	36
Advanced topics — Domain 2	39
Industry variations — Domain 2	39
Case study — Domain 2: the Meridian case that should have been written (public health)	40
Case study B — Domain 2: the platform that could not be stopped (financial services)	40
Executive perspective — Domain 2	41
Calculation exercises — Domain 2	41
Practitioner's toolkit — Domain 2	42
Exam preparation — Domain 2	43
Domain 2 summary	43
Governance, Organization and Decision Rights	44
Why this domain exists	44
Knowledge Area 3.1 Governance models	45
Knowledge Area 3.2 Sponsorship and steering	49
Knowledge Area 3.3 Assurance, gates and escalation	55
Advanced topics — Domain 3	62
Industry variations — Domain 3	63
Case study — Domain 3: the four-week month (health, Meridian)	64
Case study B — Domain 3: the decision nobody made (financial services)	65
Executive perspective — Domain 3	65
Calculation exercises — Domain 3	66

Practitioner's toolkit — Domain 3	67
Exam preparation — Domain 3	67
Summary — Domain 3	68
Planning, Scheduling and Delivery Flow (quantitative flagship)	71
Why this domain exists	71
Knowledge Area 6.1 Planning levels and logic networks	72
Knowledge Area 6.2 The critical path and float	75
Knowledge Area 6.3 Resources, constraints, milestones and rolling wave	78
Knowledge Area 6.4 Delivery flow: predictive, agile and hybrid; recovery and forecasting	82
Advanced topics — Domain 6	87
Industry variations — Domain 6	88
Case study — Domain 6: recovering Auriga (utilities / technology)	88
Case study B — Domain 6: the milestone that was typed, not computed (public programme)	89
Executive perspective — Domain 6	89
Calculation exercises — Domain 6	90
Practitioner's toolkit — Domain 6	90
Exam preparation — Domain 6	91
Domain 6 summary	91
Cost, Resources and Commercial Awareness (quantitative flagship)	93
Why this domain exists	93
Knowledge Area 7.1 Estimating and budgeting	94
Knowledge Area 7.2 The cost baseline, actuals and forecasting	97
Knowledge Area 7.3 Earned value: measurement, variances and forecasting	98
Knowledge Area 7.4 Resource economics, procurement strategy and cash	103
Advanced topics — Domain 7	106
Industry variations — Domain 7	106
Case study — Domain 7: the forecast the board actually needed (utilities)	107
Case study B — Domain 7: past the point of total assumption (technology)	107
Executive perspective — Domain 7	108
Calculation exercises — Domain 7	108
Practitioner's toolkit — Domain 7	109
Exam preparation — Domain 7	110
Domain 7 summary	110
Risk, Uncertainty and Resilience (quantitative)	111
Why this domain exists	111
Knowledge Area 8.1 Threats, opportunities and identification	112
Knowledge Area 8.2 Analysis: from screening to quantification	114
Knowledge Area 8.3 Responses, reserves and governance	119
Knowledge Area 8.4 Resilience, bias and crisis leadership	121
Advanced topics — Domain 8	123
Industry variations — Domain 8	124
Case study — Domain 8: the survey Auriga did not commission (utilities)	124
Case study B — Domain 8: the diversified portfolio that wasn't (technology programme)	125
Executive perspective — Domain 8	125

Calculation exercises — Domain 8..... 126

Practitioner's toolkit — Domain 8 126

Exam preparation — Domain 8 127

Domain 8 summary 127

01

PART ONE

Leading projects

Domains 1-4 — the profession, strategy and selection, governance and decision rights, and delivery architecture: what a project leader is for and answerable for.



01

DOMAIN 1

The Project Leadership Profession

Group: Leading projects (Domain 1 of 4 in Part One). **Target:** ~68 pages. **Binds to:** the PCI Book Pattern Specification and the shared registries ([docs/books/registries/](#)). This domain fixes the book's core vocabulary — accountability, outputs/outcomes/benefits, systems thinking, the responsible-AI principle — that every later domain assumes. British English; USD (+SAR where useful, indicative USD 1 ≈ SAR 3.75).

— Why this domain exists

Every technique in this book — the networks of Domain 6, the earned value of Domain 7, the risk quantification of Domain 8 — is instrumentation. This domain is about the person holding the instruments: what a project leader is for, what they are answerable for, and how they think. It establishes the delivery landscape and how leading temporary work differs from running a steady-state function (KA 1.1); defines accountability precisely, because the word is used loosely and the distinction decides who answers when things go wrong (KA 1.2); builds the two habits of mind that separate a leader from a coordinator — systems thinking, and the discipline of outputs versus outcomes versus benefits (KA 1.3); and grounds professional ethics and the governed use of AI (KA 1.4). A reader who finishes only this domain should already reason like the profession: know what success actually means, know who owns it, and stay answerable for the judgment even when a machine produced the analysis.

Learning objectives. After this domain a candidate can: distinguish projects, programmes and portfolios and explain what each optimises; contrast project leadership with operational leadership; explain why a temporary organisation makes trust and clarity urgent; distinguish accountability from responsibility and apply the distinction to a delegation; state the leader's obligations to sponsor, team, users and the public, and where they conflict; explain the professional standard of care; analyse a project as a system with feedback and delay; distinguish outputs, outcomes and benefits and demonstrate why output-based claims overstate value; compute the value of time in benefit terms; describe the ethical obligations of the role; and apply the PCI responsible-AI principle to a realistic delivery decision.

The master case. One programme runs through this domain and returns in Domains 2 and 16: **Meridian Care Records**, a fictional public-health programme rolling a shared clinical-records system out to **40 clinics**. Its numbers are used in KA 1.3 to make the outputs/outcomes/benefits distinction arithmetic rather than rhetorical.

KNOWLEDGE AREA 1.1

Projects, programmes and portfolios

Topics: 1.1.1 the delivery landscape · 1.1.2 project versus operational leadership · 1.1.3 the temporary organisation.

1.1.1 The delivery landscape

Definitions. A **project** is a temporary endeavour undertaken to produce a defined result. A **programme** is a group of related projects and change activities managed together to achieve **outcomes** no single project could deliver alone. A **portfolio** is the set of projects and programmes an organisation chooses to fund, selected and balanced against strategy and capacity. The three are not sizes; they are different objects of management:

Level	Optimises	Primary success test	This book
Project	Delivery of a defined output to time, cost, quality	Was the thing delivered, fit for purpose?	Domains 4–13
Programme	Coherent outcomes and benefits across components	Did the change actually land?	Domain 15
Portfolio	Value and balance of the whole investment set	Are we doing the right things at all?	Domains 2, 15

The practical consequence is that **the same event can be a project success and a programme failure**: Meridian delivering all 40 clinic installations on time is a project triumph and worth nothing if clinicians do not use the system. KA 1.3 makes that arithmetic.

1.1.2 Project versus operational leadership

Operations optimise a repeating process; projects deliver a novel result once. The differences that change a leader's behaviour:

	Operational leadership	Project leadership
Work	Repeating, refined over time	Novel, done once
Team	Stable, known, yours	Assembled, borrowed, temporary
Authority	Usually line authority	Often influence without authority
Time	Continuous	Bounded, with an end that must be planned for
Learning curve	Amortised over many cycles	Paid once, mid-delivery
Failure mode	Drift, inefficiency	Late discovery, then irreversibility

The last row is the one that shapes everything downstream. Operational problems are usually recoverable at similar cost whenever they are found; project problems get monotonically more expensive as commitments harden. That is why this book spends so much on early definition (Domain 5), honest schedules (Domain 6) and early warning (Domains 7–8) — **in temporary work, lead time is the cheapest resource a leader has, and it is spent whether used or not.**

1.1.3 The temporary organisation

A project team is a temporary organisation: people arrive with divided loyalties, no shared history, and no accumulated trust to draw on. Three consequences a leader must design for. **Trust has no time to accrue naturally**, so it must be built deliberately and early — the team charter (Domain 12) exists because the informal norms that operations grow over years must here be stated in week one. **Clarity substitutes for familiarity**: roles, decision rights and done-definitions have to be explicit precisely because nobody can infer them (Domains 3-5). **Endings must be engineered**: a temporary organisation dissolves, and the knowledge, the relationships and the operational readiness either transfer deliberately or evaporate (Domain 16).

AI in this domain — first statement of the principle

AI now sits inside the delivery environment itself: drafting plans, summarising status, generating risk lists, answering questions the team used to bring to a senior colleague. The suite's governing principle, stated once here and applied in every later domain:

AI proposes; the professional verifies, decides and remains accountable.

Three things it means concretely for a leader, before any technique in this book is applied. **Verification is a duty, not a preference** — machine output entering a plan, report or decision is checked by a named human against evidence. **Accountability is not transferable** — "the model said so" is never a defence, because accountability cannot be delegated to software (KA 1.2 makes the reason precise). **Disclosure is part of honesty** — material AI use in a deliverable is stated, so colleagues can weigh it. Domain 14 builds the full governance architecture; this domain establishes that the obligation attaches to the person, not the tool.

■ KEY TERMS — KA 1.1

Term	Meaning
Project	Temporary endeavour producing a defined result.
Programme	Related projects managed together for outcomes no one project delivers.
Portfolio	The funded, balanced set of projects and programmes.
Temporary organisation	A team with no accumulated trust or shared norms; both must be designed.
Irreversibility	The rising cost of correction as commitments harden — projects' defining failure mode.
Responsible-AI principle	AI proposes; the professional verifies, decides and remains accountable.

■ SAMPLE MCQS — KA 1.1

MCQ 1.1-A [1.1.1 · Analysis] A programme's four projects each deliver on time and to budget, but the intended service change does not materialise. The most accurate statement is: - A. the programme succeeded, since all its projects succeeded - B. the projects succeeded on their own test while the programme failed on its own test — outcomes are not the sum of outputs - C. the projects must therefore have failed - D. programme success cannot be assessed until every project closes

Rationale: Projects are judged on delivered outputs, programmes on outcomes; the two tests are different, which is exactly why the distinction exists. C rewrites history to preserve a single verdict, and D denies the measurement programmes exist to make.

MCQ 1.1-B [1.1.2 · Analysis] Why does project leadership place more weight on early definition than operational leadership does? - A. project teams are less capable - B. project problems grow monotonically more expensive as commitments harden, so lead time is the cheapest resource available - C. operations have no need for planning - D. projects have larger budgets

Rationale: Irreversibility is the structural difference (1.1.2). A and D are unfounded; C misstates operational practice, which plans continuously — just with the ability to correct cheaply next cycle.

MCQ 1.1-C [1.1.3 · Recall] A team charter is written in week one primarily because: - A. governance requires the document - B. a temporary organisation has no accumulated norms, so what operations grow informally must here be made explicit - C. it replaces the need for a project plan - D. it fixes the team's membership for the duration

Rationale: Explicit clarity substitutes for familiarity the team has had no time to build. A is a by-product, C confuses artefacts, D is false — membership usually changes.

■ SELF-CHECK — KA 1.1

1. Can a project succeed while its programme fails? — Yes: outputs delivered, outcomes not realised. The tests differ (1.1.1).
2. Name the defining failure mode of temporary work. — Late discovery followed by irreversibility; correction cost rises as commitments harden.
3. State the responsible-AI principle. — AI proposes; the professional verifies, decides and remains accountable.

KNOWLEDGE AREA 1.2

The project leader's accountability

Topics: 1.2.1 accountability and responsibility · 1.2.2 the obligation set · 1.2.3 the professional standard of care.

1.2.1 Accountability and responsibility

Definitions. Responsibility is the obligation to do the work — it can be shared and delegated. **Accountability** is the obligation to answer for the outcome — it sits with exactly one named person and **cannot be delegated**, only reassigned by whoever holds the authority to do so. A leader may hand a task to a specialist, an outsourcer or an AI tool; the answer for the result stays where it was.

Two working rules follow.

Delegation transfers responsibility, never accountability. The leader who delegates keeps the duty to specify, to verify at a proportionate depth, and to notice when the work is failing. "I gave it to the vendor" is a description of responsibility, not a defence of accountability.

One accountable person per outcome. Where two people are accountable for the same thing, nobody is: each defers, and the gap is discovered late. This is why Domain 3 builds explicit decision rights and why RACI-style matrices insist on a single "A" per row.

The AI corollary is now immediate rather than novel. A tool has no standing to be accountable — it cannot be asked to answer, cannot be sanctioned, cannot hold a duty. So every use of AI in delivery leaves accountability exactly where it was: with the professional who used the output. That is not a policy choice PCI made; it follows from what accountability is.

1.2.2 The obligation set

A project leader owes duties in four directions, and the skill is holding them together honestly rather than pretending they never conflict:

To	Owes	Typical conflict
Sponsor / organisation	Honest forecasts, stewardship of funds, no surprises	Pressure to report reassurance rather than status
Team	Clarity, safety, achievable asks, fair credit	Absorbing schedule pressure by quietly overloading people
Users / customers	A result that is genuinely fit for purpose	Delivering on time by descoping what users needed most
Public / third parties	Safety, legality, honest claims, environmental care	Cost pressure meeting a safety or compliance margin

The escalation duty. When these conflict beyond the leader's authority to resolve, the professional act is to escalate with options — not to absorb the conflict silently and hope. An unescalated conflict becomes a late surprise, and late surprises are the failure mode of §1.1.2.

The honesty asymmetry. Bad news travels badly: it costs the messenger and helps the project. A leader's most consequential cultural act is making early bad news safe — because a schedule or cost problem reported five weeks early is usually recoverable (Domain 6's case study turns on exactly this) and the same problem reported at the deadline is not.

1.2.3 The professional standard of care

Professionals are judged not on outcomes alone but on whether they exercised the **care a competent practitioner would** in the circumstances. Practically, four things constitute care in this discipline: **method appropriate to the stakes** (a two-week internal task and a safety-critical programme do not warrant the same rigour); **evidence proportionate to the claim** (a forecast presented to a board carries its basis — Domain 7's estimate basis sheet); **records that let others check the work** (decision logs, assumption registers, audit trails); and **candour about uncertainty** (ranges, not false single numbers). A leader who did all four and still missed can defend the work. One who hit the date by luck and cannot show any of them has no defence available when the next project does not get lucky.

■ KEY TERMS — KA 1.2

Term	Meaning
Responsibility	The obligation to do; shareable and delegable.
Accountability	The obligation to answer; single-holder, non-delegable.
Escalation duty	The obligation to raise, with options, a conflict beyond one's authority.
Honesty asymmetry	Bad news costs the messenger and helps the project; leaders must make it safe.
Standard of care	The care a competent practitioner would exercise in the circumstances.

■ SAMPLE MCQS — KA 1.2

MCQ 1.2-A [1.2.1 · Application] A leader outsources the migration design to a specialist vendor under a fixed-price contract. Accountability for the migration's success now rests with: - A. the vendor, under the contract - B. the leader, who retains the duty to specify, verify proportionately and act on failure - C. jointly and equally with the vendor - D. the sponsor, who approved the contract

Rationale: The contract transfers responsibility (and commercial risk); accountability is non-delegable. C describes the arrangement that guarantees nobody answers (1.2.1); D confuses approving a route with owning the delivery.

MCQ 1.2-B [1.2.1 · Analysis] An AI planning tool generates a schedule that the leader submits unchanged; it contains an infeasible dependency that causes a two-month slip. The accountability position is: - A. shared with the tool's vendor, who supplied a defective product - B. the leader's: a tool cannot answer for an outcome, and submitting unverified output is a failure of the verification duty - C. nobody's — the failure was technological - D. the team's, for not catching it

Rationale: Accountability requires a party who can be asked to answer, so it cannot attach to software (1.2.1); the specific lapse is verification (1.1's principle). Vendor recourse (A) is a commercial question and does not relocate the professional answer; D inverts delegation.

MCQ 1.2-C [1.2.3 · Analysis] A leader's project overran despite a documented method, ranged forecasts, a decision log and escalated risks. Under the standard of care this is: - A. professional negligence, because the project overran - B. defensible practice — care is judged on the exercise of competent method, not on outcome alone - C. irrelevant, since standards of care apply only to regulated professions - D. defensible only if the overrun was under 10 %

Rationale: The standard is conduct-based (1.2.3). A collapses care into outcome; C is false in substance — the expectation attaches to professional practice generally; D invents a threshold.

■ SELF-CHECK — KA 1.2

1. What can be delegated and what cannot? — Responsibility can; accountability cannot.
2. Why can accountability never attach to an AI tool? — Accountability is the obligation to answer; a tool cannot be asked to answer, sanctioned, or hold a duty.
3. Name the four constituents of care in this discipline. — Proportionate method, evidence matching the claim, checkable records, candour about uncertainty.

KNOWLEDGE AREA 1.3**Systems thinking and value**

Topics: 1.3.1 projects as systems · 1.3.2 outputs, outcomes, benefits, value · 1.3.3 leading under uncertainty.

1.3.1 Projects as systems

A project is a system: parts that interact, with **feedback** and **delay** between cause and visible effect. Three systems behaviours a leader must recognise, because each defeats common-sense management.

Delay hides consequence. Adding people to a late project consumes the productivity of those already on it (onboarding, communication paths) before adding any, so the intervention looks harmful before it looks helpful — and leaders often reverse it in the interval.

Local optimisation degrades the whole. A discipline lead maximising their own efficiency levels their team's workload smoothly and delivers the interface documents late, stalling three other teams. Every function can be efficient while the project is slow — the constraint logic of Domain 6 (6.A.1) is the schedule expression of this.

Pressure moves, it does not vanish. Squeeze the schedule and the pressure emerges as quality defects (Domain 9), scope disputes (Domain 5), team attrition (Domain 12) or supplier claims (Domain 10). A leader's question on receiving any "we absorbed it" report is where did it go?


1.3.2 Outputs, outcomes, benefits and value

Definitions. An **output** is what the project delivers (a system installed, a road built). An **outcome** is the change in how things work as a result (clinicians actually using the records; journeys actually faster). A **benefit** is the measurable improvement that the outcome produces (hours released, cost avoided, harm reduced). **Value** is benefit weighed against cost and risk, for a stated stakeholder. The chain has a hard property: **each link can fail independently**, so an output claimed as a benefit is an overstatement by construction.

WORKED EXAMPLE**Worked example 1.3.2 — what Meridian actually delivers.**

1. **Setup.** Meridian Care Records installs the shared records system in **40 clinics** (the output). Realistic adoption at 12 months is **70 %** of clinics using it in daily practice. In an adopting clinic the system releases **6 clinician-hours per week**, valued at **USD 85 per hour**, over a **48-week** operating year.
2. **Formula.** Annual benefit = clinics × adoption × hours/week × rate × weeks. Contrast with the output-based claim, which counts all 40 clinics.
3. **Substitution.** Adopting clinics $40 \times 0.70 = 28$; benefit $28 \times 6 \times 85 \times 48$. Output-based claim $40 \times 6 \times 85 \times 48$.
4. **Result.** Benefit **USD 685,440 per year** ($\approx$ SAR 2.57 million indicatively). The output-based claim gives **USD 979,200** — an overstatement of **USD 293,760**, exactly **30.0 %**, which is the non-adoption rate reappearing as fictitious value.
5. **Interpretation.** The 30 % gap is not an estimating error; it is the outcome link being skipped. This is why a benefits register (Domain 16) tracks adoption as a measure in its own right, and why "40 clinics live" is a milestone, not a benefit. Sensitivity is the leader's real lever: at 50 % adoption the benefit is USD 489,600; at 90 %, USD 881,280. **A programme of this shape creates more value by moving adoption than by moving its installation date** — which should change where the leader spends personal attention (Domain 11's engagement work, not Domain 6's schedule compression).

Each link can fail independently — an output claimed as a benefit overstates by construction



Adoption is the outcome measure that links an output to any benefit at all — so it needs its own owner (Toolkit 1.T.1).

Fig 1.3.1 — The value chain, and where it leaks

1.3.3 Leading under uncertainty

Uncertainty is the condition, not an exception to be eliminated before acting. Three disciplines constitute leading well inside it.

Decide at the right moment. Deciding early buys momentum and forecloses options; deciding late preserves options and costs time. The professional question is not "do we have enough information?" but what is the last responsible moment, and what does waiting cost?

Price the time. In benefit-generating work, delay has an arithmetic cost.

WORKED EXAMPLE

Worked example 1.3.3 — what a week of Meridian delay costs.

1. **Setup.** Meridian's benefit runs at the KA 1.3.2 rate once adopted. The team can compress the rollout by **8 weeks** at an incremental cost of **USD 60,000**. Should the leader buy it?
2. **Formula.** Benefit per week = adopting clinics × hours × rate. Net = (weeks × benefit per week) – cost. Breakeven weeks = cost ÷ benefit per week.
3. **Substitution.** Per week $28 \times 6 \times 85 = 14,280$; eight weeks **114,240**; net **114,240 – 60,000**.
4. **Result.** Benefit per week **USD 14,280**; eight weeks **USD 114,240**; **net +USD 54,240**. Breakeven at **4.20 weeks** — beyond which acceleration pays.
5. **Interpretation.** The decision is not "is 60,000 a lot?" but "does it buy more than it costs?" — and here it clearly does, provided adoption materialises. Note the dependency the arithmetic makes visible: every figure rests on the 70 % adoption of 1.3.2, so **spending 60,000 to arrive sooner at an unadopted system buys nothing**. Sequencing engagement before acceleration is not a preference; it is what the numbers say.

Act reversibly where you can. Prefer decisions that can be unwound, pilots over commitments, and options kept open at low cost — then commit hard once uncertainty has genuinely reduced. This is the leadership behaviour behind rolling-wave planning (Domain 6, KA 6.3.3) and staged investment (Domain 2).

AI in this KA

Systems and value reasoning is where AI assistance is weakest and most confidently wrong: a model will produce a fluent benefits case that silently treats outputs as benefits, because that is what most documents in its training data do. The governed use is narrow and real — stress-test a benefits chain by asking what would have to be true, generate adoption scenarios, surface second-order effects a team has not considered — followed by human judgment on which link is actually fragile. The check that catches the common failure is one question: for each claimed benefit, which measured outcome produces it, and who owns that measure? (Domain 16 turns that into the benefits register.)

■ KEY TERMS — KA 1.3

Term	Meaning
Output / outcome / benefit / value	Delivered thing · change in how things work · measurable improvement · benefit against cost and risk.
Adoption	The outcome measure linking an output to any benefit at all.
Feedback and delay	Systems properties that make interventions look wrong before they look right.
Local optimisation	Function-level efficiency that degrades whole-system performance.
Last responsible moment	The latest point a decision can be taken without foreclosing a needed option.
Cost of delay	The benefit forgone per period of lateness.

■ SAMPLE MCQS — KA 1.3

MCQ 1.3-A [1.3.2 · Application] 40 clinics installed; 70 % adoption; 6 hours/week released per adopting clinic at USD 85/hour over 48 weeks. The defensible annual benefit is: - A. USD 979,200 - B. USD 685,440 - C. USD 293,760 - D. USD 14,280

Rationale: $40 \times 0.70 \times 6 \times 85 \times 48 = 685,440$. A is the output-based claim that skips adoption; C is the overstatement itself; D is the weekly benefit.

MCQ 1.3-B [1.3.2 · Analysis] A programme board is told "40 clinics live — benefits delivered". The soundest challenge is: - A. ask for the installation evidence - B. ask which measured outcome produces the benefit and who owns that measure, since installation is an output and adoption is unverified - C. accept it, since the output target was met - D. ask for the cost variance instead

Rationale: The claim skips the outcome link, which is precisely where 30 % of the value leaked in 1.3.2. A verifies the wrong thing, C repeats the error, D changes the subject.

MCQ 1.3-C [1.3.3 · Application] Benefit accrues at USD 14,280 per week once adopted. Compressing the rollout by 8 weeks costs USD 60,000. The decision and its breakeven are: - A. reject — 60,000 exceeds the weekly benefit - B. accept — net +USD 54,240, breakeven at 4.20 weeks - C. accept — net +USD 114,240 - D. indifferent — the two are equal

Rationale: $8 \times 14,280 = 114,240$, less the 60,000 cost, nets **+54,240**; breakeven $60,000/14,280 = 4.20$ weeks. A compares a total against a weekly rate; C forgets to deduct the cost; D asserts an equality the arithmetic denies.

MCQ 1.3-D [1.3.1 · Analysis] A discipline lead reports having "absorbed" a two-week schedule squeeze with no impact. The systems-literate response is: - A. record the recovery and move on - B. ask where the pressure went — into quality, scope, people or suppliers — because pressure moves rather than vanishing - C. squeeze the remaining teams equally - D. treat it as evidence the original estimate was padded

Rationale: Pressure relocates (1.3.1); the leader's job is to find where. A accepts an unverified claim, C compounds it, D leaps to a conclusion the report does not support.

■ SELF-CHECK — KA 1.3

1. Why is an output claimed as a benefit an overstatement by construction? — Because the outcome link can fail independently; unadopted outputs produce no benefit (Meridian's 30 %).
2. What is the leader's question on hearing "we absorbed it"? — Where did the pressure go?
3. What must be true before paying to accelerate Meridian? — That adoption materialises; arriving sooner at an unadopted system buys nothing.

KNOWLEDGE AREA 1.4

Ethics and the responsible use of AI

Topics: 1.4.1 professional ethics in delivery · 1.4.2 the responsible-AI principle applied · 1.4.3 the leader's AI accountability.

1.4.1 Professional ethics in delivery

Ethics in this discipline is rarely dramatic; it is a series of small reporting and allocation decisions taken under pressure. Four recurring tests:

- **Report what is, not what is wanted.** Presenting an optimistic case as the base case is a misrepresentation whatever the spreadsheet says (the same rule PFL-AI applies to forecasts).
- **Do not spend other people's margins silently.** Absorbing pressure into team overtime, a supplier's contingency or a safety margin is a decision someone is entitled to be told about.
- **Claim only what is true.** Milestone and benefit claims are statements of fact to people relying on them (1.3.2).
- **Compete and procure honestly.** Conflicts declared, evaluations run as published, suppliers not used as free consultancies (Domain 10's ethical sourcing).

Conflicts of interest are handled by disclosure and separation: declared before engagement, managed or declined — and tested by daylight, would every party, seeing the full picture, still regard this as impartial? Where the answer wavers, impartiality has already failed.

1.4.2 The responsible-AI principle applied

The principle from KA 1.1 becomes operational through four obligations that recur in every later domain:

Obligation	What it requires of a leader
Appropriate use	Match the tool to the task; prohibit it where the stakes or data forbid (Domain 14)
Verification	Machine output checked against evidence by a named human before reliance
Transparency	Material AI use disclosed in deliverables and decisions
Human decision	The judgment, and its record, belongs to a person

The failure modes to expect, stated honestly: **hallucination** (fluent, false specifics — the invented precedent, the plausible citation); **silent staleness** (confident answers from out-of-date data); **bias** inherited from training data, especially in anything touching people; **confidentiality leakage** (project data entering a tool is a disclosure); and **over-trust through fluency** — the deepest one, because polished output suppresses the scrutiny that rough output invites. None of these is a reason to refuse the technology; each is a reason the verification duty is not optional.

1.4.3 The leader's AI accountability

Three concrete leadership acts, which Domain 14 then systematises:

Name the owner. Every AI-assisted artefact that informs a decision has a human owner who verified it. Anonymous machine output has no place in a governed decision record (Domain 3, KA 3.3.4).

Make verification proportionate and real. A summary for internal reading needs a glance; a forecast in a board paper needs recomputation and a source trace; a safety-relevant analysis needs independent review. Uniform verification is either wasteful or negligent — usually both, in different places.

Protect the team's judgment. The subtler risk is not a wrong output but the atrophy of the skills needed to notice one. A leader who accepts AI-drafted plans no one on the team can critique has traded capability for speed, and will discover the price during the first recovery (Domain 6, KA 6.4).

■ KEY TERMS — KA 1.4

Term	Meaning
Daylight test	Would every party, seeing the full picture, still regard this as impartial?
Hallucination	Fluent, confidently-stated content that is false.
Over-trust through fluency	Polished output suppressing the scrutiny rough output would attract.
Verification proportionality	Depth of checking matched to the stakes of reliance.
Named owner	The human who verified an AI-assisted artefact and answers for it.

■ SAMPLE MCQS — KA 1.4

MCQ 1.4-A [1.4.1 · Analysis] A sponsor asks that the optimistic scenario be presented as the base case "to keep the board confident". The professional response is: - A. comply — scenario labels are presentational - B. decline, and offer the honest base case with sensitivities and the evidence that supports confidence - C. comply while noting the true base case privately - D. present both without indicating which is the base case

Rationale: Candour about status is a duty, not a style (1.4.1). C documents the misrepresentation without preventing it; D abdicates the professional judgment the board is relying on.

MCQ 1.4-B [1.4.2 · Application] Which AI failure mode most directly explains an authoritative-sounding schedule that encodes assumptions nobody made? - A. confidentiality leakage - B. over-trust through fluency, compounded by absent verification - C. model bias - D. hardware limitation

Rationale: Polished output invites less scrutiny (1.4.2), and the missing control is verification. A concerns data egress, C concerns systematic skew (real, but not this symptom), D is irrelevant.

MCQ 1.4-C [1.4.3 · Analysis] A leader requires the same verification depth for every AI-assisted artefact, from internal summaries to safety analyses. This is: - A. best practice — consistency is the point - B. simultaneously wasteful and negligent: proportionality means depth matched to the stakes of reliance - C. acceptable only if the team is small - D. unnecessary, since approved tools need no verification

Rationale: Uniform depth over-checks the trivial and under-checks the critical (1.4.3). A mistakes uniformity for rigour; D abandons the duty entirely.

■ SELF-CHECK — KA 1.4

1. State the daylight test. — Would every party, seeing the full picture, still regard this as impartial?
2. Which AI failure mode is hardest to counter, and why? — Over-trust through fluency: polish suppresses the scrutiny that rough work attracts.
3. What does "protect the team's judgment" mean in practice? — Do not accept AI output the team cannot critique; capability traded for speed is repaid during recovery.

— Advanced topics — Domain 1

1.A.1 Authority, influence and the borrowed team

Most project leaders hold less authority than accountability — the defining asymmetry of the role. The sources of influence that actually work in temporary organisations: **competence** (visible command of the work), **reciprocity** (having helped others deliver), **information** (being the person with a true picture), **coalition** (sponsor and peer support secured before it is needed), and **fairness** (a record of allocating credit and blame accurately). Positional authority is the weakest of the six in borrowed teams, because the borrowed people's careers are decided elsewhere. The practical implication: influence is built in advance, and a leader who begins building it during a crisis is too late.

1.A.2 Success criteria and the multiple-verdict problem

Projects are judged simultaneously against delivery (time, cost, quality), outcome and benefit realisation, and stakeholder satisfaction — which is why the same project is often described as a success and a failure by different parties, honestly. The professional response is not to pick a flattering verdict but to **agree the criteria in advance, with owners and measures** (Domain 2's benefits mapping, Domain 5's acceptance criteria), and to report against all of them. Unagreed success criteria are settled retrospectively by whoever is most disappointed.

1.A.3 The reviewer's leadership eye

Invariants an experienced reviewer tests in the first hour on any project: exactly one accountable name per outcome; a benefits chain where each claimed benefit traces to a measured outcome with an owner; decision records that show what was known when; escalations that arrived early enough to matter; forecasts with ranges rather than single numbers; and AI-assisted artefacts carrying a named verifier. Every one of these is cheap to check and expensive to have missing — the same posture Domains 6 and 7 apply to schedules and cost.

— Industry variations — Domain 1

- **Public programmes.** Accountability is doubled — professional and political — and the public obligation of 1.2.2 is enforceable through scrutiny bodies; benefits (not outputs) are the currency of accountability, and announced dates often precede plans (Domain 6's Case B).
- **Regulated industries (pharma, aviation, nuclear).** The standard of care is codified and auditable; the "care exercised" defence of 1.2.3 is a documented reality rather than an argument, and AI use faces explicit validation expectations.
- **Construction and engineering.** Authority is contractual as much as organisational; influence runs through commercial relationships, and the leader's obligation set includes site safety as a first-order duty.
- **Technology and product.** Outcome ownership frequently sits with a product role rather than the delivery leader, making the 1.3.2 chain a shared accountability that must be explicitly split — the commonest source of "who owns adoption?" disputes.
- **Professional services and consulting.** Duties to the client organisation and to the end-users can diverge, and independence (1.4.1) becomes a standing rather than occasional test.

— Case study — Domain 1: Meridian under scrutiny (public health)

Situation. Eighteen months in, Meridian Care Records reports 40 of 40 clinics installed, on time, 3 % under budget. The programme is publicly praised. Nine months later a health-service audit finds clinician hours released are running at roughly a third of the business case, and the programme is publicly described as a failure. Both accounts cite true facts.

Analysis. The delivery verdict was sound: the outputs landed. The benefits case had been written against **outputs** — all 40 clinics, no adoption term — so it claimed USD 979,200 per year where the honest figure at 70 % adoption was USD 685,440, and actual adoption reached only about 40 %, giving roughly USD 391,680. Two failures compounded: an arithmetic one (skipping the outcome link, 1.3.2) and an accountability one (**nobody owned adoption** — the delivery leader owned installations, the clinical directorate owned practice change, and no single name owned the measure that connected them, 1.2.1).

What was done. The recovery was not technical. Adoption was made a named accountability with a monthly measure; the benefits case was restated with an adoption term and a sensitivity range (50/70/90 % → 489,600 / 685,440 / 881,280); clinical champions were funded in the eleven lowest-adoption clinics; and the programme board's report changed from "clinics live" to "clinics using, and hours released". Adoption reached 68 % within a year.

What the domain teaches here. Every element of this domain appears: outputs are not benefits (1.3.2); one accountable name per outcome (1.2.1); success criteria agreed in advance or settled by the disappointed (1.A.2); and the leader's attention belongs where the value actually leaks — which the arithmetic, not instinct, identified.

— Case study B — Domain 1: the plan nobody could critique (financial services)

Situation. A platform-migration leader under time pressure used an AI planning assistant to generate the full delivery schedule and dependency network, reviewed it briefly, and issued it. It was detailed, internally consistent and professionally formatted. It also assumed the legacy data extract could run concurrently with the schema freeze — an assumption no one on the team had made, and which the two engineers who would have spotted it never saw.

What happened. The conflict surfaced eleven weeks later during rehearsal, costing a seven-week slip and a weekend of reputational damage with the regulator's supervisory team. The post-incident review found no fault in the tool: the schedule was a reasonable inference from an under-specified prompt. The failures were the leader's — no verification against the discipline leads (the 1.1 principle and 1.4.2 obligation), and no named verifier on the artefact (1.4.3). The compounding finding was cultural: the polish of the output had actively discouraged challenge, and two team members later said they had assumed the network "must have been checked".

What the domain teaches here. Fluency is not evidence. Accountability did not move to the vendor or the model — it stayed with the person who issued the plan (MCQ 1.2-B is this case in miniature). And the durable fix is structural, not exhortative: a named verifier per artefact, discipline leads walking their own links, and a team kept capable of critique.

— Executive perspective — Domain 1

What a project leader cannot delegate in this domain:

- **The single accountable name.** For every outcome, including the awkward ones like adoption. Where two names appear, the leader's job is to reduce it to one.
- **The benefits chain's honesty.** Each claimed benefit traced to a measured outcome with an owner. A leader who lets outputs be reported as benefits has authorised the failure of Meridian's first eighteen months.
- **The obligation set, held openly.** Sponsor, team, users, public — with conflicts escalated with options rather than absorbed in silence.
- **Bad news made safe.** The single highest-leverage cultural act available, because it converts irreversible surprises into recoverable problems.
- **The AI accountability line.** Named verifiers, proportionate depth, and a team that remains able to critique what the tools produce.

— Calculation exercises — Domain 1

Exercise 1.1 A programme installs a system in 60 sites; adoption is 65 %; each adopting site releases 5 hours/week at USD 90/hour over 46 weeks. Compute the defensible annual benefit and the overstatement in an output-based claim. Solution. Adopting sites $60 \times 0.65 = 39$; benefit $39 \times 5 \times 90 \times 46 = \text{USD } 807,300$. Output-based $60 \times 5 \times 90 \times 46 = \text{USD } 1,242,000$; overstatement $\text{USD } 434,700 = 35.0 \%$, exactly the non-adoption rate. Common error: applying adoption to the rate rather than the site count — the same product here, but it conceals which link failed and misdirects the fix.

Exercise 1.2 Using Meridian's figures (28 adopting clinics, 6 h/week, USD 85/h), compute the cost of a 5-week delay and the maximum justified acceleration spend for 5 weeks. Solution. Per week $28 \times 6 \times 85 = \text{USD } 14,280$; five weeks $\text{USD } 71,400$ — which is also the maximum justified spend, since above it acceleration destroys value. Common error: comparing the spend against annual benefit rather than the benefit for the weeks actually saved.

Exercise 1.3 Meridian's benefit at 50 %, 70 % and 90 % adoption (40 clinics, 6 h/week, USD 85/h, 48 weeks). What does the spread imply for where the leader spends attention? Solution. $\text{USD } 489,600 \cdot 685,440 \cdot 881,280$. The 40-point adoption swing moves annual benefit by $\text{USD } 391,680$ — far more than plausible schedule compression could add (Exercise 1.2 values 5 weeks at 71,400). Attention belongs on adoption (Domain 11), not acceleration.

— Practitioner's toolkit — Domain 1

Adoption-ready artefacts; adapt headings to your organisation, then keep them stable.

Toolkit 1.T.1 — Leader's accountability map (one page)

Per outcome: the outcome stated in outcome language (not output language) · **the single accountable name** · responsible parties · the measure that evidences it and its owner · the decision

authority and threshold (Domain 3) · review cadence. Rule: if any row has two accountable names or an output in the outcome column, it is not finished.

Toolkit 1.T.2 — Benefits-chain test (run on any benefits claim)

- [] Each claimed benefit names the measured **outcome** that produces it.
- [] Each outcome measure has an owner and a reporting cadence.
- [] Adoption (or the equivalent behavioural change) appears explicitly, with a value.
- [] The claim carries a sensitivity range, not a single number.
- [] Output milestones are labelled as milestones, never as benefits.
- [] Cost of delay per period is stated, so timing decisions can be priced (1.3.3).

Toolkit 1.T.3 — AI-use and verification record

Per AI-assisted artefact: tool and approved environment · data classification cleared · prompt/intent summary · **verification performed and its depth** · **named verifier** · disclosure status in the deliverable · date. Standing agenda item, not an archive: the register's purpose is that the next artefact is verified, not that the last one was.

— Exam preparation — Domain 1

The traps. Treating outputs as benefits (1.3.2 — the domain's central trap, and Meridian's whole case study) · saying accountability is "shared" (1.2.1) · assigning accountability to a vendor or a tool (MCQ 1.2-A, 1.2-B) · comparing an acceleration cost against annual rather than saved-period benefit (Exercise 1.2) · reading a delivery success as a programme success (1.1.1) · accepting "we absorbed it" (1.3.1) · uniform AI verification depth (1.4.3) · judging the standard of care by outcome alone (1.2.3).

Reflection questions. 1. Take your current project: write the single accountable name for its principal outcome. If you hesitated, or wrote two, what does that predict? 2. For each benefit in your business case, which measured outcome produces it and who owns that measure? What percentage of your claimed value survives the question? 3. Which AI-assisted artefact in your last month of work had no named verifier — and what would have caught it if it had been wrong?

— Domain 1 summary

Project leadership is a distinct profession because temporary work fails distinctively: problems discovered late become irreversible, teams have no accumulated trust to draw on, and lead time is the cheapest resource available and is spent whether used or not. The discipline rests on a precise view of accountability — responsibility is delegable, the obligation to answer is not, one name per outcome — held across four directions at once (sponsor, team, users, public) with conflicts escalated rather than absorbed, and judged by the care a competent practitioner would exercise rather than by outcome alone. Two habits of mind make a leader rather than a coordinator: systems thinking, which expects feedback, delay, local optimisation and pressure that relocates instead of vanishing; and the outputs-outcomes-benefits-value chain, which Meridian turns into arithmetic — 40 clinics installed produce USD 685,440 of annual benefit at 70 % adoption, not the USD 979,200 an output-based claim asserts, and the 30 % gap is the outcome link being skipped. Timing decisions are priceable (USD 14,280 per week; acceleration breaking even at 4.20 weeks), and attention belongs where value actually leaks. Around all of it sits professional ethics and the suite's governing principle — **AI**

proposes; the professional verifies, decides and remains accountable — with named verifiers, proportionate depth, and a team kept capable of critique. Domain 2 turns strategy into selected work; Domain 3 gives the accountability of this domain its governance machinery.

02

DOMAIN 2

Strategy, Selection and Business Alignment

Group: Leading projects (Domain 2 of 4 in Part One). **Target:** ~70 pages. **Binds to:** the PCI Book Pattern Specification and the shared registries (<docs/books/registries/>). This domain continues the **Meridian Care Records** programme from Domain 1 into its business case and benefits map, and bridges to the appraisal machinery the finance book builds in detail (PFL-AI, Domain 4). British English; USD (+SAR where useful, indicative **USD 1 ≈ SAR 3.75**).

— Why this domain exists

Domain 1 established that a project is judged on outcomes and benefits, not delivered outputs, and proved it arithmetically. That leaves the prior question unanswered: **how does work get chosen at all, and how does anyone know it was worth choosing?** This domain answers it. It sets the strategic context and treats alignment as a test that must be re-run rather than a box ticked at approval (KA 2.1); builds the business case as a decision instrument, with the options analysis and appraisal that make it one, and the selection models that rank competing candidates (KA 2.2); works benefits mapping, measurement baselines and the sustainability dimension, plus the assumption and dependency management on which every forecast rests (KA 2.3); and closes with the hardest leadership act in the discipline — **stopping** (KA 2.4). The through-line is that a business case is a promise about the future made to obtain money, and the professional question is always whether anyone will be held to it.

Learning objectives. After this domain a candidate can: explain how strategy becomes a portfolio and why alignment decays; describe the components of a decision-grade business case; construct and appraise a genuine options set including the do-nothing baseline; build a benefits map from outputs through outcomes to benefits with owners and measures; **build a benefits profile that ramps rather than assuming steady state from day one, and quantify the overstatement a flat assumption produces**; define measurement baselines and avoid double-counting; rank candidates with weighted scoring and under a binding constraint, and state each method's limits; manage assumptions and dependencies as live risks; apply sunk-cost-free reasoning to a continuation decision; and set kill criteria that make honest gates possible.

The master programme. Meridian Care Records continues from Domain 1 — the shared clinical-records rollout to **40 clinics**, whose verified benefit figures (full potential **USD 979,200** per year; **USD 685,440** at the realistic 70 % adoption) are now used to build the business case that should have been written for it.

KNOWLEDGE AREA 2.1

Strategic context and alignment

Topics: 2.1.1 from strategy to portfolio · 2.1.2 drivers, constraints and environment · 2.1.3 alignment as a repeated test.

2.1.1 From strategy to portfolio

Strategy states intent; a portfolio is what an organisation is actually doing about it. The gap between the two is where most strategic failure lives, and it is visible in three symptoms a leader should be able to name: **a portfolio that does not reflect the stated priorities** (the strategy says "digital first", the funded work is 80 % maintenance); **more work than capacity**, so everything is slow and nothing is finished (Domain 15's capacity management); and **no route for stopping**, so the portfolio only ever accretes.

The mechanism that closes the gap is a **portfolio process** with real teeth: candidate work described comparably, ranked against explicit criteria, funded to capacity rather than to appetite, and reviewed at intervals that allow re-decision. Where a project leader sits inside this, the obligation is specific — **describe your work honestly enough to be compared**, including its weaknesses, because a portfolio process fed by advocacy rather than evidence selects the best storyteller.

2.1.2 Drivers, constraints and environment

A business case that does not name its **driver** cannot be tested. The families, and what each implies:

Driver	Success looks like	Failure mode if misidentified
Compliance / regulatory	The obligation is met by the date	Gold-plating a mandatory minimum
Cost reduction	Measured, sustained cost decrease	Cost moved rather than removed
Revenue / growth	New or protected income	Cannibalised existing revenue counted as new
Capability / enabling	A later benefit becomes possible	Benefits claimed twice — here and downstream
Risk reduction	Exposure measurably lower	Unquantified "insurance" spending
Service / outcome improvement	Users measurably better served	Output delivered, outcome unmeasured (Domain 1)

Constraints — funding envelope, regulatory deadline, capacity, technical dependency — are not the same as objectives, and conflating them corrupts option generation: a "constraint" that is actually a preference silently eliminates the option that would have won. The professional habit is to write constraints down and mark each **hard** (physics, law, contract) or **soft** (preference, convention), because only soft ones can be traded.

2.1.3 Alignment as a repeated test

Alignment is granted at approval and **decays** thereafter: strategy moves, markets move, and the project's own understanding of what it can deliver moves. A project can be perfectly executed and strategically irrelevant by the time it lands — the most expensive form of success available.

Three practices keep alignment live. **Re-test at gates** — each stage gate (Domain 3, KA 3.3.1) asks not only "is delivery on track?" but "is this still the right thing?". **Track the assumption set** rather than the conclusion: business cases fail through their assumptions (KA 2.3.4), and an assumption that has been falsified invalidates the case whether or not delivery is green. **Name a benefits owner outside the project** — usually the operational leader who will run the changed service — because a project that owns its own benefits case will never be the one to report that it has stopped making sense (Domain 1's accountability rule; Domain 16's realisation).

AI in this KA

Portfolio-level analysis is a legitimate AI application: normalising heterogeneous candidate descriptions into comparable form, surfacing dependency clashes across dozens of cases, flagging where a stated driver and the proposed benefits do not match. The governed limits are two. **Strategic weighting is a governance judgment**, not an optimisation output — a model asked to "optimise the portfolio" will faithfully optimise whatever proxy it was given, and the proxy is the decision. And **advocacy is invisible to a model**: an AI reading business cases cannot tell an evidenced benefit from a confident one, so it amplifies whatever the strongest writer submitted. AI proposes; the professional verifies, decides and remains accountable.

■ KEY TERMS — KA 2.1

Term	Meaning
Portfolio process	Comparable description, explicit criteria, funding to capacity, periodic re-decision.
Driver	The reason the work exists; determines what success means and how it is measured.
Hard / soft constraint	Physics, law, contract vs preference and convention; only soft ones are tradeable.
Alignment decay	The erosion of strategic fit after approval, as strategy and understanding move.
Benefits owner	The accountable person outside the project who will realise the benefit.

■ SAMPLE MCQS — KA 2.1

MCQ 2.1-A [2.1.3 · Analysis] A project is delivering to plan, but the strategic priority it served has been superseded. The governance-sound response is: - A. continue — the project is on track and the case was approved - B. re-test the case at the next gate and decide on current strategy and remaining cost and benefit - C. continue but reduce scope proportionately - D. transfer the project to another portfolio

Rationale: Alignment decays and gates exist to re-decide (2.1.3), on remaining cost and benefit (KA 2.4.2). A treats approval as permanent; C is an arbitrary compromise that decides nothing; D relocates the question without answering it.

MCQ 2.1-B [2.1.2 · Application] A case states its constraint as "must use the existing platform", which reflects an architectural preference rather than a contractual or legal requirement. The correct treatment is: - A. accept it as a constraint, since the architects have stated it - B. record it as a **soft** constraint, so it can be traded in option generation and its cost made visible - C. remove it from the case entirely - D. re-classify it as an objective

Rationale: Only hard constraints (physics, law, contract) are untradeable; a soft one recorded as hard silently eliminates options (2.1.2). C loses real information; D confuses a limit with a goal.

■ SELF-CHECK — KA 2.1

1. Name the three symptoms of a strategy-portfolio gap. — Funded work not reflecting priorities; more work than capacity; no route for stopping.
2. Why must a benefits owner sit outside the project? — A project owning its own benefits case will not be the one to report that the case has stopped making sense.
3. What distinguishes a hard from a soft constraint? — Physics, law or contract versus preference; only soft constraints can be traded.

KNOWLEDGE AREA 2.2

The business case and selection

Topics: 2.2.1 the business case as a decision instrument · 2.2.2 options and appraisal · 2.2.3 selection and prioritisation models.

2.2.1 The business case as a decision instrument

A decision-grade business case answers six questions and nothing else: **why** (the driver and the problem, with evidence); **what options** exist, including doing nothing; **what it costs**, with a range and an accuracy class (Domain 7, KA 7.1.1); **what it delivers**, as outcomes and benefits with owners and measures (KA 2.3); **what could go wrong**, with the exposure quantified (Domain 8); and **how it will be judged**, with success criteria agreed in advance (Domain 1, KA 1.A.2).

Its two failure modes are worth naming precisely. A **advocacy document** answers only "why" and "what it delivers", omitting ranges, options and risk — designed to obtain approval rather than support a decision. A **compliance artefact** answers all six in a template nobody reads, produced after the decision has been taken informally. Both are common; the test that separates them from the real thing is whether the case **could have concluded "no"**.

The living-case principle. The case is not filed at approval; it is the instrument the gates re-test (2.1.3) and the source of the benefits register Domain 16 tracks. A business case whose figures no longer match the project is not a historical document — it is an unmanaged forecast.

2.2.2 Options and appraisal

A **genuine options set** includes the **do-nothing / do-minimum baseline** (against which all value is measured — without it, benefits are unmeasurable), at least one option that is materially cheaper than the preferred one, and at least one that is materially more ambitious. An options set constructed to make the preferred option look inevitable — the "straw options" pattern — is the advocacy failure mode in structural form, and reviewers spot it by checking whether any option could plausibly have won.

Appraisal applies the discounted machinery PFL-AI Domain 4 builds in full: NPV as the primary value measure, with the range and sensitivity that make it honest. This book uses the results; the finance book derives them. What belongs to a delivery leader is the **benefits profile** — because that is where most business cases are wrong, and the error is systematic.

WORKED EXAMPLE**Worked example 2.2.2 — the Meridian business case, twice.**

1. **Setup.** Meridian costs **USD 2,400,000** to deliver. Domain 1 established the arithmetic: full potential at 100 % adoption is **USD 979,200** per year; realistic steady-state adoption is **70 %**, worth **USD 685,440** per year. The benefits are appraised over **8 years** at a **7 %** discount rate. Compare the case as commonly written — full potential from year one — with an honest **ramped** profile: 40 % adoption in year 1, 60 % in year 2, 70 % thereafter.
2. **Formula.** Flat case: $PV = \text{annual benefit} \times AF(r, n)$. Ramped case: $PV = \sum (\text{potential} \times \text{adoption}_t) / (1 + r)^t$. $NPV = PV - \text{cost}$.
3. **Substitution.** $AF(0.07, 8) = 5.971299$. Flat: $979,200 \times 5.971299$. Ramped: $391,680/1.07 + 587,520/1.07^2 + 685,440 \times (\text{years 3–8 discounted})$.
4. **Result.**

	Year 1	Year 2	Years 3–8	PV of benefits	NPV
Flat, full potential	979,200	979,200	979,200	5,847,096	+3,447,096
Ramped, 70 % steady state	391,680	587,520	685,440	3,732,898	+1,332,898

The flat case overstates NPV by **USD 2,114,198 — 158.6 %** of the honest figure.

1. **Interpretation.** Both cases approve the programme, which is precisely why the error survives: the decision is unchanged, so nobody checks, and the promise is what gets remembered. Meridian was later judged a failure against benefits it never could have produced (Domain 1's case study), and this table is where that judgement was actually created — at approval, two years before anyone noticed. Two further points a leader should extract. The ramp is not pessimism; it is the adoption curve of the same 70 % figure, merely arriving when it actually arrives. And the case is robust: **breakeven flat-equivalent adoption is 41.05 %** — the programme creates value at any sustained adoption above about 41 %, which is a far more useful sentence for a board than either NPV.




Fig 2.2.1 — Two Meridian business cases from identical facts

2.2.3 Selection and prioritisation models

Weighted scoring ranks candidates against criteria carrying explicit weights. Its value is that it forces the criteria and weights into the open, where they can be argued; its limits are that scores are **ordinal judgments** (the arithmetic caution of Domain 8, KA 8.2.1 applies) and that weights are chosen by whoever runs the model — so the model can be steered by anyone who understands it.

WORKED EXAMPLE

Worked example 2.2.3 — four candidates, two rankings.

- Setup.** Four candidate programmes scored 1-5 against weighted criteria: strategic fit (0.35), benefit value (0.30), deliverability (0.20), risk — inverse, so higher is safer (0.15).
- Formula.** Weighted score = $\Sigma(\text{weight} \times \text{score})$. Then, under a binding constraint, rank by **NPV per unit of the scarce resource** instead.
- Substitution and result.**

Candidate	Fit	Benefit	Deliver.	Risk	Weighted score
Meridian	4	5	3	3	3.95
Beta	5	3	4	4	4.05
Gamma	3	4	5	4	3.85
Delta	2	2	5	5	3.05

Now suppose the binding constraint is **integration-team capacity**, of which only 3 units exist: Meridian needs 3 units for an NPV of 1,693,072; Beta 2 units for 1,200,000; Gamma 1 unit for 900,000.

Candidate	NPV	Capacity	NPV per unit
Meridian	1,693,072	3	564,357
Beta	1,200,000	2	600,000
Gamma	900,000	1	900,000

1. **Interpretation.** Two defensible methods give two different answers. The scoring model ranks **Beta first** (4.05 against Meridian's 3.95) on strategic fit and deliverability. Under the capacity constraint, **Beta + Gamma** together consume the 3 units for a combined NPV of **2,100,000**, beating Meridian's 1,693,072 alone — so the constrained answer is to run the two smaller programmes, which neither the raw NPV ranking nor the scoring model would have selected. This is the delivery-side twin of PFL-AI's capital rationing (its Domain 4, KA 4.3.2), and it carries the same warning: **greedy ranking by value per unit is a heuristic, not an optimum** — lumpy candidates require enumerating the feasible sets. The leadership content is that **the binding constraint decides the method**, and naming that constraint honestly (usually a scarce team, not money) is the whole game.

What no model does. Scoring and ratios cannot see option value, sequencing dependencies, or the strategic cost of not doing something. They order candidates; they do not decide. A portfolio board that treats a scoring output as the decision has automated its own accountability (Domain 1, KA 1.2.1).

AI in this KA

Business-case drafting is now routinely AI-assisted, and this KA identifies the specific hazard: a model will produce a **fluent, complete-looking case with a flat benefits profile**, because that is what most cases in its training data look like. It will also generate plausible options sets that are structurally straw. The verification duty is therefore concrete — check the benefits profile ramps, check at least one option could have won, check every benefit traces to a measured outcome with an owner (Domain 1's benefits-chain test), and check the range and class on the cost. An AI-drafted case that passes those four checks is genuinely useful work; one that has not been checked against them is the advocacy document of 2.2.1 with better formatting.

■ KEY TERMS — KA 2.2

Term	Meaning
Do-nothing baseline	The comparison against which all benefits are measured; without it value is unmeasurable.
Straw options	An options set built so the preferred option cannot lose.
Benefits profile	The time-phased benefit stream; ramps in reality, flat in most business cases.
Breakeven adoption	The sustained adoption level at which NPV is zero — a more useful board sentence than NPV.
Weighted scoring	Ranking by $\Sigma(\text{weight} \times \text{ordinal score})$; forces criteria into the open, steerable by whoever sets weights.
Value per unit of constraint	Ranking basis when a scarce resource binds; a heuristic, not an optimum.

■ SAMPLE MCQS — KA 2.2

MCQ 2.2-A [2.2.2 · Application] Full benefit potential is 979,200 per year; steady-state adoption is 70 %; the profile ramps 40 %/60 %/70 % and is appraised over 8 years at 7 % (AF = 5.971299). Against a flat full-potential case, the ramped NPV is lower by: - A. USD 979,200 - B. USD 2,114,198 - C. USD 1,332,898 - D. USD 293,760

Rationale: Flat PV 5,847,096 less ramped PV 3,732,898 = **2,114,198** of overstated present value. C is the honest NPV itself; A is one year's potential; D is Domain 1's single-year output-based overstatement, a different figure.

MCQ 2.2-B [2.2.2 · Analysis] Both the flat and ramped Meridian cases produce a positive NPV and the same approval decision. Why does the flat case still matter? - A. it does not — the decision is unchanged, so the error is harmless - B. because the case becomes the promise the programme is later judged against, and benefits that were never deliverable are recorded as failure - C. because the discount rate must be recalculated - D. because the flat case understates cost

Rationale: The approval is identical; the commitment is not (2.2.2, and Domain 1's case study where Meridian was publicly called a failure). A is the reasoning that let the error survive.

MCQ 2.2-C [2.2.3 · Analysis] Integration capacity of 3 units is the binding constraint. Meridian needs 3 units for NPV 1,693,072; Beta 2 for 1,200,000; Gamma 1 for 900,000. The value-maximising selection is: - A. Meridian alone — the highest single NPV - B. Beta + Gamma — combined NPV 2,100,000 within the 3-unit constraint - C. Meridian + Gamma — the two highest NPVs - D. all three, phased across two years

Rationale: Beta and Gamma together fit the constraint and beat Meridian's 1,693,072. C needs 4 units; D changes the premise rather than answering under it; A ignores the constraint's implications.

MCQ 2.2-D [2.2.1 · Recall] The clearest test that a business case is a decision instrument rather than an advocacy document is: - A. whether it follows the corporate template - B. whether it could have concluded "no" - C. whether it was approved - D. whether its NPV is positive

Rationale: An options set and evidence that permit a negative conclusion are what make it a decision (2.2.1). A is the compliance failure mode; C and D are outcomes, not tests.

■ SELF-CHECK — KA 2.2

1. What must an options set contain to be genuine? — A do-nothing baseline, a materially cheaper option, a materially more ambitious one — and at least one that could plausibly have won.
2. Why is breakeven adoption a better board sentence than NPV? — It states the condition the value depends on, which is what the board can actually influence and monitor.
3. When does ranking by value per unit of constraint fail? — With lumpy candidates: the feasible sets must be enumerated, as Beta + Gamma shows.

KNOWLEDGE AREA 2.3

Benefits, value and sustainability

Topics: 2.3.1 benefits mapping · 2.3.2 measures and baselines · 2.3.3 sustainability and ESG value · 2.3.4 assumptions and dependencies.

2.3.1 Benefits mapping

A **benefits map** is a chain, drawn explicitly, from what the project delivers to why anyone wanted it: **output → enabling change → outcome → benefit → strategic objective**. Domain 1 established the links; this KA draws them and assigns them.

The **enabling change** is the step most maps omit and most programmes fail on. Meridian's outputs (installed clinics) produce no benefit without training, workflow redesign, clinical-champion support and data migration — none of which are software, and several of which sit outside the project's authority. A benefits map that leaps from output to benefit has hidden exactly the work that determines whether anything is realised, which is why Meridian's adoption stalled at 40 % while installation ran to plan.

Every element on the map carries an **owner** and, for outcomes and benefits, a **measure**. The rule from Domain 1 (KA 1.2.1) applies unchanged: one accountable name per outcome, and the benefits owner sits outside the project (2.1.3).

Each link can fail independently — and the highlighted column is the one most maps omit




Fig 2.3.1 — Meridian's benefits map, with the enabling change restored

2.3.2 Measures and baselines

A benefit is only claimable against a **baseline** — the measured position before the change. Three disciplines make measurement honest:

- **Baseline before, not after.** Measure the current state before the change lands; a baseline reconstructed afterwards is an estimate shaped by the result it is used to justify.
- **Attribution.** Other things change too. Where a benefit could have other causes, say so and quantify what portion is plausibly attributable — over-claiming is the fastest way to lose the credibility that future cases depend on.
- **No double-counting.** The single commonest audit finding in benefits work: an enabling programme claims the benefit its dependent projects also claim, and the portfolio's total exceeds what the organisation could possibly realise. The portfolio benefits register (Domain 15) exists to catch this, and it only works if each benefit has exactly one claimant.

Cash-releasing versus non-cash-releasing benefits must be distinguished plainly. Meridian releases clinician hours; that becomes cash only if the organisation reduces cost or converts the hours into additional activity that has value. Reporting released hours as savings, when the headcount is unchanged and the hours are absorbed, is the error that makes finance directors distrust benefits cases — and it is avoided by stating the conversion explicitly, or by claiming the benefit in its true unit (capacity) rather than in money it never became.

2.3.3 Sustainability and ESG value

Environmental, social and governance considerations enter this domain in two distinct ways, and conflating them causes most of the confusion.

As constraints and obligations: emissions limits, environmental permits, accessibility law, labour standards in the supply chain (Domain 10's ethical sourcing). These are hard constraints (2.1.2) and belong in the case as requirements, not benefits.

As value: whole-life carbon reduction, energy cost avoided, social value created, resilience to climate risk. These belong in the benefits map with measures and owners like any other benefit — and with the same honesty about attribution and monetisation. Where a carbon reduction has a price (a compliance cost avoided, a levy, an internal carbon price), it can be valued directly; where it does not, it should be reported in its physical unit alongside the financial case rather than converted with an invented price. **Whole-life thinking** is the practical contribution: a cheaper capital solution with higher operating energy and shorter life frequently loses on a whole-life basis, and the appraisal horizon of KA 2.2.2 is what makes that visible (PFL-AI Domain 4's equivalent annual value handles unequal lives).

2.3.4 Assumptions and dependencies

Business cases fail through their assumptions, and the assumptions are usually stated once and never revisited. Meridian's case rested on an adoption assumption that was never tracked as a measure — the direct cause of everything Domain 1's case study describes.

An **assumption register** carries, per assumption: the statement, why it is believed, the impact if false, the owner, the date it will be tested, and the **trigger** that would falsify it. Its relationship to the risk register is exact: **every assumption is a risk in disguise** (Domain 8, KA 8.1.3), so material assumptions get register entries with **EMV** where quantifiable.

Dependencies — on other projects, on suppliers, on decisions not yet taken — are managed by naming them, naming the owner on the other side, and stating the date by which the dependency must be satisfied. Two failure modes recur: **assumed dependencies** (nobody on the other side knows they owe you anything) and **unmanaged reciprocals** (they are depending on you too, on a different date). Domain 15's programme dependency management handles the multi-project case; at project level the discipline is simply that a dependency without a named counterpart is a hope.

AI in this KA

Benefits mapping is judgment work about causation, which is where AI is weakest — a model will happily assert that an output produces a benefit because most documents assert exactly that (Domain 1's KA 1.3 warning). Two legitimate uses. **Challenge:** ask a model to argue that the benefit will not materialise, and the enabling changes it names are often the ones the map omitted. **Cross-checking at portfolio scale:** detecting the same benefit claimed by two cases is pattern-matching over documents, which machines do well and humans do badly across fifty cases. The conclusions stay human, with named benefit owners.

■ KEY TERMS — KA 2.3

Term	Meaning
Benefits map	Output → enabling change → outcome → benefit → objective, with owners and measures.
Enabling change	The non-project work (training, workflow, adoption support) without which no benefit occurs.
Baseline	The measured pre-change position; must be measured before, not reconstructed after.
Attribution	The portion of a measured improvement plausibly caused by this project.
Double-counting	The same benefit claimed by more than one case; one claimant per benefit.
Cash-releasing / non-cash-releasing	Benefits that reduce spend versus those that release capacity.
Assumption register	Statement, basis, impact-if-false, owner, test date, falsifying trigger.

■ SAMPLE MCQS — KA 2.3

MCQ 2.3-A [2.3.1 · Analysis] A benefits map runs directly from "system installed" to "USD 685,440 released per year". Its principal defect is: - A. the benefit figure is too precise - B. the **enabling change** is missing — training, workflow redesign and adoption support, largely owned outside the project, are what convert the output into the outcome - C. it should be expressed in hours rather than money - D. the map should start with the strategic objective

Rationale: Omitting enabling change hides the work that determines realisation (2.3.1) — exactly how Meridian stalled at 40 % adoption with installation on plan. C is a separate (also real) question of unit; D is presentational.

MCQ 2.3-B [2.3.2 · Application] A programme releases 6 clinician-hours per week per clinic; headcount is unchanged and the hours are absorbed by existing demand. Reporting this as a cash saving is: - A. correct — released time has a value - B. incorrect: it is a non-cash-releasing (capacity) benefit unless cost is reduced or the capacity is converted to valued activity - C. correct if the hourly rate is documented - D. acceptable provided it is discounted

Rationale: Value released is not cash released (2.3.2); claiming otherwise is the error that discredits benefits cases. A and C mistake a valuation rate for a cash effect; D discounts a figure that was never cash.

MCQ 2.3-C [2.3.4 · Analysis] An enabling platform programme and three dependent projects each claim the same USD 4m of downstream benefit. The portfolio total is: - A. USD 16m, since each case is individually valid - B. overstated by double-counting — one benefit, one claimant, with the enabler credited through the dependents or vice versa but not both - C. USD 4m, and the dependent projects have no benefits - D. indeterminate until delivery completes

Rationale: Double-counting is the standing audit finding (2.3.2); the fix is a single claimant per benefit in the portfolio register. C overcorrects by denying the dependents any case; D defers a question answerable now.

MCQ 2.3-D [2.3.3 · Application] A carbon reduction has no priced compliance consequence for the organisation. The soundest treatment in the case is: - A. omit it, since it cannot be valued - B. report it in its physical unit alongside the financial case, rather than monetising it with an invented price - C. monetise it using a rate found in a published study - D. record it as a constraint

Rationale: Unpriced benefits are reported honestly in their own unit (2.3.3). A discards real value; C imports a price the organisation does not face; D confuses value with obligation.

■ SELF-CHECK — KA 2.3

1. Which element do most benefits maps omit, and why does it matter most? — The enabling change; it is the work that converts outputs into outcomes and it usually sits outside the project.
2. What makes a baseline trustworthy? — Measurement before the change, not reconstruction after.
3. State the relationship between assumptions and risks. — Every assumption is a risk in disguise; material ones get register entries with quantified exposure.

KNOWLEDGE AREA 2.4

Strategic termination

Topics: 2.4.1 stopping as a strategic decision · 2.4.2 sunk cost and escalation of commitment · 2.4.3 kill criteria and honest gates.

2.4.1 Stopping as a strategic decision

Organisations that cannot stop projects cannot really select them either: a portfolio that only accretes has no capacity for anything new, and its selection process becomes a queue. Stopping is therefore a **portfolio capability**, not an admission of failure — and the leaders who can do it create the capacity that funds everything else.

The barriers are cultural and structural, and both are addressable. **Culturally**, cancellation is read as failure and punished, so nobody proposes it; the counter is to celebrate early stopping explicitly and to distinguish "we learned this will not work" from "we managed this badly". **Structurally**, no gate has authority to stop, or the decision is scheduled after the money is committed; the counter is Domain 3's decision rights with a real stop option and gates placed before irreversible commitments (Domain 1's irreversibility point).

2.4.2 Sunk cost and escalation of commitment

The principle. A continuation decision depends only on **remaining cost and remaining benefit**. Money already spent is irrelevant to it — it cannot be recovered by spending more, and it carries no information about the future beyond what is already in the forecast.

WORKED EXAMPLE**Worked example 2.4.2 — Meridian at the reset point.**

1. **Setup.** A hypothetical reset: USD 1,800,000 of the 2,400,000 has been spent. Completion requires a further **USD 900,000** (the original 600,000 plus 300,000 of newly discovered integration work). The benefits case has been re-based on measured adoption, giving a remaining benefit present value of **USD 780,000**. Continue or stop?
2. **Formula.** Forward-looking NPV = remaining benefit PV – remaining cost. Sunk cost excluded.
3. **Substitution.** 780,000 – 900,000.
4. **Result. Forward NPV = (USD 120,000).** The correct decision is to **stop** (or to descope to a variant whose remaining benefit exceeds its remaining cost).
5. **Interpretation.** The pull to continue is powerful and entirely arithmetical in appearance: "we have spent 1,800,000 — stopping wastes it." That reasoning adds a sunk 1,800,000 back into a decision it cannot affect, and it converts a 120,000 loss into a 900,000 one. The 1,800,000 is already gone under **either** choice; the only question is whether the next 900,000 buys 780,000 of value. It does not. Note what stopping preserves: the 900,000, redeployable to candidates the portfolio could not previously fund (KA 2.2.3's constrained selection) — which is why stopping is a strategic act and not merely a loss.

Escalation of commitment is the behavioural pattern that makes this hard: personal identification with the project, reputational exposure for the approver, and the asymmetry that continuing defers the reckoning while stopping realises it now. Its countermeasures are structural, not exhortative — pre-set kill criteria (2.4.3), decision-makers who did not author the case, and forward-only framing in every gate paper (Domain 8's bias table).

2.4.3 Kill criteria and honest gates

Kill criteria are conditions, agreed **in advance**, under which the project will stop — "if adoption in the pilot cohort is below 50 % at month six", "if the regulatory approval is not granted by Q3", "if remaining cost exceeds remaining benefit at any gate". Their power is entirely in the advance agreement: the same evidence assessed against a pre-agreed threshold is a decision, while assessed against no threshold it is a negotiation, and negotiations are won by whoever has most invested.

An **honest gate** has four properties: a **real stop option** with someone authorised to exercise it; **evidence prepared to a standard set beforehand**, not assembled to support a conclusion; **decision-makers who did not write the case**; and a **recorded decision with reasons** (Domain 3's auditability). A gate that has never stopped anything is not a control; it is a milestone with a meeting attached.

Stopping well is its own competence: capture what was learned and make it findable (Domain 9's lessons); harvest anything reusable (code, designs, cleaned data, supplier relationships, qualified people); close contracts deliberately rather than by abandonment (Domain 10); redeploy the team with their reputations intact — because how a cancellation is handled determines whether anyone proposes the next one; and tell stakeholders honestly, early (Domain 11).

AI in this KA

The relevant AI use here is monitoring: tracking assumption triggers and kill-criteria thresholds across a portfolio, and flagging when a project's own reported data has crossed one. That is genuinely valuable because the crossing is usually visible in data nobody is reading. What must

remain human is the decision and the framing: a model asked whether to continue will optimise whatever objective it was handed, and the escalation-of-commitment bias lives in the humans around it, not in the arithmetic. The leader's use of AI here is to make the trigger impossible to miss — not to outsource the courage.

■ KEY TERMS — KA 2.4

Term	Meaning
Forward-looking NPV	Remaining benefit PV less remaining cost; the only basis for continuation.
Sunk cost	Unrecoverable spend, irrelevant to a continuation decision.
Escalation of commitment	Continuing because of what is invested; countered structurally.
Kill criteria	Pre-agreed stop conditions; powerful only because agreed in advance.
Honest gate	Real stop option, pre-set evidence standard, independent deciders, recorded reasons.

■ SAMPLE MCQS — KA 2.4

MCQ 2.4-A [2.4.2 · Application] 1,800,000 is spent; completion needs a further 900,000; remaining benefit PV is 780,000. The correct decision and its basis are: - A. continue — stopping wastes the 1,800,000 already spent - B. stop — forward NPV is (120,000); the sunk 1,800,000 is irrelevant to the remaining choice - C. continue — total spend 2,700,000 against total benefits exceeds the original case - D. continue at reduced pace to spread the cost

Rationale: Only remaining cost and benefit bear on the decision (2.4.2). A is the sunk-cost fallacy stated plainly; C reintroduces sunk cost as "total"; D changes the schedule without changing the negative economics.

MCQ 2.4-B [2.4.3 · Analysis] A portfolio's gates have never stopped a project in four years. The soundest inference is: - A. selection is excellent, so no project has needed stopping - B. the gates are not functioning as controls — no real stop option, evidence assembled to support continuation, or deciders too close to the cases - C. the gate criteria are too lenient and should be tightened numerically - D. stopping is unnecessary if delivery is well managed

Rationale: A control that never fires is not demonstrably a control (2.4.3). A is implausible across a whole portfolio; C addresses thresholds when the defect is process and authority; D mistakes delivery quality for strategic validity.

MCQ 2.4-C [2.4.3 · Recall] Kill criteria derive their power from: - A. the severity of the thresholds - B. being agreed in advance, so the same evidence produces a decision rather than a negotiation - C. being set by the project manager - D. being confidential until invoked

Rationale: Advance agreement converts assessment into decision (2.4.3); without it, whoever has most invested wins the argument. C removes independence; D prevents the behavioural effect entirely.

■ SELF-CHECK — KA 2.4

1. What is the only correct basis for a continuation decision? — Remaining benefit against remaining cost; sunk cost excluded.

2. Why is stopping a strategic act rather than a loss? — It releases capacity and funding for candidates the portfolio could not otherwise support.
3. Name two properties of an honest gate. — A real stop option with authority; deciders who did not write the case (also: pre-set evidence standards; recorded reasons).

— Advanced topics — Domain 2

2.A.1 Real options thinking in selection

Static NPV values a committed plan; much project value lies in **flexibility** — the option to scale if a pilot succeeds, to defer while uncertainty resolves, to abandon cheaply, to switch technology. Treating these as options changes selection behaviour in a specific way: a small, staged first phase with a negative standalone NPV can be the highest-value choice if it buys the right to a large second phase at a decision point where far more is known. The practical rule for a delivery leader is **structure the work so that decision points exist** — phases sized to reach a genuine learning milestone, contracts that permit exit (Domain 10), architectures that permit switching. Formal option valuation is specialist; the thinking is not, and its absence is why organisations routinely commit fully to things they could have tested for a fraction.

2.A.2 Portfolio balance, not just portfolio ranking

Ranking selects the best individual candidates; **balance** asks whether the resulting set is survivable and coherent. Four balance questions: **risk profile** (is everything high-risk, or everything safe and incremental?), **time-to-benefit** (does anything land this year, or is all value in years 3–5?), **capability spread** (does the whole set depend on one scarce team — the correlation problem of Domain 8, KA 8.A.1, at portfolio level?), and **strategic coverage** (are some stated objectives served by nothing at all?). A portfolio of individually optimal projects can be collectively unbalanced, and no scoring model detects it — which is why Domain 15 treats portfolio balancing as its own discipline.

2.A.3 The reviewer's business-case eye

Invariants testable in an hour: a do-nothing baseline exists and is quantified; at least one option could plausibly have won; costs carry a range and an accuracy class; **the benefits profile ramps**; every benefit traces to a measured outcome with an owner outside the project; enabling changes are named with their owners; no benefit is claimed by another case in the portfolio; assumptions have test dates and falsifying triggers; risk exposure is quantified and reconciles to the register; kill criteria exist and are numeric; and success criteria were agreed in advance. The single highest-yield check is the benefits profile — it is wrong in most cases, and it is wrong in the direction that flatters.

— Industry variations — Domain 2

- **Public sector.** Appraisal is often mandated (published guidance, prescribed discount rates, optimism-bias uplifts), social value and distributional effects sit beside financial NPV, and political commitment can precede the case — making 2.4's stopping problem acutely hard and 2.1.3's re-testing acutely necessary.

- **Regulated utilities.** Investment cases are made to a regulator on a periodic cycle; benefits are defined by the regulatory framework, and the "customer" for the case is external.
- **Technology and product.** Real options thinking (2.A.1) is the native mode — staged funding, pilots, kill criteria on usage metrics — and benefits are frequently non-cash-releasing capacity or optionality, so 2.3.2's honesty about units matters most here.
- **Energy and infrastructure.** Whole-life and carbon considerations are first-order value (2.3.3), horizons are decades, and PFL-AI's bankability tests (its Domain 5) run alongside the internal case because external capital must also be convinced.
- **Health and social programmes.** Benefits are largely non-financial and attribution is genuinely hard; the professional response is to measure outcomes in their own units with pre-change baselines rather than to force monetisation.

— Case study — Domain 2: the Meridian case that should have been written (public health)

Situation. Meridian's approved business case claimed **USD 979,200** of annual benefit from year one — the full-potential figure — supporting an NPV of **+3,447,096** over eight years at 7 % against a delivery cost of 2,400,000. Every number in it was arithmetically correct. No adoption term appeared anywhere, and no benefit had an owner outside the programme.

What the honest case looks like. Same cost, same horizon, same rate, same 70 % steady-state adoption Domain 1 establishes — but profiled as it would actually arrive (40 % / 60 % / 70 %): PV of benefits **3,732,898**, NPV **+1,332,898**. The approved case overstated present value by **USD 2,114,198**, or **158.6 %** of the honest NPV. It also omitted the sentence that would have mattered most: **breakeven sustained adoption is 41.05 %**, so the programme creates value at any adoption above roughly 41 % — a threshold a board can monitor, unlike an NPV.

What followed, and why it was predictable. Both cases approve the programme, so the flat profile changed no decision — and was therefore never challenged. Two years later, delivery complete and adoption at 40 %, the programme was measured against a promise of 979,200 and publicly judged a failure (Domain 1's case study). Actual benefit at 40 % adoption was **391,680** — above the 41.05 % breakeven only marginally, and far below a promise that had never been achievable.

What was done differently on the successor programme. The case carried a ramped profile with adoption as a tracked measure; a named benefits owner in the clinical directorate; the enabling changes (training, workflow redesign, champions) itemised with their own owners and costs; the breakeven adoption stated on the front page; and a kill criterion at 50 % adoption by month six. The board's monthly question changed from "are clinics live?" to "what is adoption, against 41 %?".

What the domain teaches here. A business case is a promise, and the profile is the promise's shape. An overstatement that changes no decision still changes the standard the programme will be held to — which makes the flat-profile error uniquely dangerous, because nothing at approval time creates any pressure to catch it.

— Case study B — Domain 2: the platform that could not be stopped (financial services)

Situation. A core-platform replacement, approved at USD 40m over three years, reached year four having spent USD 62m with roughly 55 % of scope delivered. A re-based assessment put remaining cost at USD 21m and remaining benefit present value at USD 14m — a forward NPV of **(USD 7m)**.

What happened at the gate. The paper presented total investment (62m spent, 21m to go) against total lifetime benefits, showing a positive "programme return", and recommended continuation. Three structural features made that recommendation almost inevitable: the gate had **no stop option** defined; the paper was written by the programme's own director; and there were **no kill criteria** from the original approval to test against. The sponsor who had championed the original case chaired the gate.

How it was resolved. An independent review, commissioned only after a regulator asked about delivery timescales, reframed the decision as forward-only: 21m to buy 14m. The programme was descoped to a variant whose remaining benefit (9m) exceeded its remaining cost (6m), the balance of capacity was redeployed, and the write-off was taken and disclosed. The descoped variant delivered in eleven months.

What the domain teaches here. Escalation of commitment is structural before it is personal: absent kill criteria, an independent decider and a real stop option, a rational actor presenting a "total investment" frame will continue. Note that the resolution was not "stop everything" — descoping to the part where remaining benefit exceeds remaining cost is the same arithmetic applied with more imagination.

— Executive perspective — Domain 2

What a project leader cannot delegate in this domain:

- **The benefits profile.** Not the NPV — the shape. The leader who lets a flat full-potential profile through has set the standard their programme will be judged against, and Meridian is what that costs.
- **The breakeven sentence.** Stating the condition on which value depends (41.05 % adoption), in a form the board can monitor monthly.
- **The enabling change and who owns it.** Naming the non-project work that converts outputs into outcomes, and securing owners for it outside the project, before approval rather than after.
- **The honest options set.** Ensuring at least one option could have won — including doing nothing.
- **Forward-only framing.** Refusing every "total investment" frame at every gate, and personally insisting on remaining cost against remaining benefit.
- **Kill criteria at approval.** The leader is the only person who can get them agreed while goodwill is high; after the first slip, nobody will.

— Calculation exercises — Domain 2

Exercise 2.1 A programme's full benefit potential is USD 600,000 per year, steady-state adoption 75 %, profiled 30 % / 60 % / 75 % thereafter, appraised over 6 years at 8 %. Delivery costs USD 1,300,000. Compute the ramped NPV and the flat full-potential NPV, and the overstatement.
 Solution. $AF(0.08, 6) = 4.622880$. Flat: $600,000 \times 4.622880 = 2,773,728$; NPV **+1,473,728**. Ramped: $180,000/1.08 + 360,000/1.08^2 + 450,000 \times (\text{years 3-6 factors } 0.793832 + 0.735030 + 0.680583 + 0.630170) = 2.839615 = 166,667 + 308,642 + 1,277,827 = 1,753,136$; NPV **+453,136**.
 Overstatement **USD 1,020,592**. Common error: applying the steady-state percentage to all years including the ramp period.

Exercise 2.2 Using Exercise 2.1's flat form, find the sustained adoption at which NPV is zero. Solution. Required annual benefit = $1,300,000 / 4.622880 = 281,210$; as a share of the 600,000 potential, **46.87 %**. Common error: dividing cost by total undiscounted benefits ($1,300,000 / 3,600,000 = 36.1 \%$), which ignores discounting and understates the threshold.

Exercise 2.3 Candidates scored 1–5 on strategic fit (0.40), benefit (0.30), deliverability (0.20), risk-inverse (0.10): P = 4/5/2/2; Q = 5/3/4/3. Rank them. Solution. P: $1.60 + 1.50 + 0.40 + 0.20 = 3.70$. Q: $2.00 + 0.90 + 0.80 + 0.30 = 4.00$. **Q ranks first**. Common error: unweighted summation (P = 13, Q = 15 — same order here, but it is coincidence, and the method loses the weights that were argued over).

Exercise 2.4 Spent USD 4,200,000; remaining cost USD 1,600,000; remaining benefit PV USD 1,950,000. Decide, and state what the sunk-cost frame would wrongly add. Solution. Forward NPV $1,950,000 - 1,600,000 = +\text{USD } 350,000 \rightarrow \text{continue}$. The sunk-cost frame would add the irrelevant 4,200,000 to the comparison (as "total spend of 5,800,000 against benefits"), which here happens to reach the same decision — the point being that it reaches it for the wrong reason and would reach the wrong one whenever forward NPV is negative.

— Practitioner's toolkit — Domain 2

Adoption-ready artefacts; adapt headings to your organisation, then keep them stable.

Toolkit 2.T.1 — Business-case skeleton (the six questions)

Why (driver named from 2.1.2, evidence, do-nothing consequence) · **Options** (including do-nothing and one materially cheaper; why each was rejected) · **Cost** (range, accuracy class, whole-life not just capital) · **Benefits** (map per 2.T.2; **profiled, not flat**; owners outside the project; measures and baselines; cash-releasing or not, stated) · **Risk** (quantified exposure reconciling to the register) · **Judgment** (success criteria agreed now, kill criteria numeric, gate schedule). Front page carries the **breakeven sentence**, not only the NPV.

Toolkit 2.T.2 — Benefits map and register row

Per benefit: description · **the outcome measure that produces it** and its owner (outside the project) · **enabling changes required**, each with an owner and whether funded · baseline value, measurement method and date measured · profile by year (**ramp explicit**) · cash-releasing or capacity · attribution note · single-claimant confirmation against the portfolio register · review cadence.

Toolkit 2.T.3 — Gate decision pack rules

One page of forward-only economics (**remaining cost, remaining benefit, forward NPV** — total spend may be reported for information but never as the decision basis) · kill criteria with current values against thresholds · assumption register extract showing which assumptions have been tested and which falsified · alignment re-test against current strategy · options at this gate, including stop and descope · decider names (none of whom authored the case) · recorded decision and reasons.

— Exam preparation — Domain 2

The traps. A flat full-potential benefits profile (2.2.2 — the domain's central error) · applying steady-state adoption to the ramp years (Exercise 2.1) · computing breakeven on undiscounted benefits (Exercise 2.2) · reading sunk cost into a continuation decision (2.4.2, Exercise 2.4) · reporting non-cash-releasing benefits as savings (2.3.2) · double-counting a benefit across an enabler and its dependents (MCQ 2.3-C) · treating a soft constraint as hard (2.1.2) · ranking by raw NPV when a resource constraint binds (MCQ 2.2-C) · mistaking a gate that has never stopped anything for a functioning control (2.4.3).

Reflection questions. 1. Take your current business case: is the benefits profile flat? What is the breakeven condition, and could you state it in one sentence to a board? 2. Which enabling changes does your case depend on, and does each have a funded owner outside your project? What happens to the benefits if they do not? 3. What are your project's kill criteria — and if there are none, who would have to agree them, and would they agree today?

— Domain 2 summary

Work gets chosen through a portfolio process, and the quality of that choice depends on candidates being described honestly enough to be compared — which makes a business case a decision instrument rather than an advocacy document only when it could have concluded "no". Alignment granted at approval decays, so gates re-test the case and not merely delivery, and the assumption set is what they test. The domain's central arithmetic is the benefits profile: Meridian's approved case claimed full potential from year one for an NPV of +3,447,096, where the same facts profiled as adoption actually arrives (40 % / 60 % / 70 % of a 979,200 potential) give **+1,332,898** — an overstatement of **USD 2,114,198**, 158.6 % of the honest figure, which changed no approval decision and set the standard the programme was later judged to have failed. The more useful board sentence is the breakeven: **41.05 % sustained adoption**. Selection methods each have limits — weighted scoring exposes criteria and can be steered; ranking by value per unit of a binding constraint beats raw NPV (Beta + Gamma's 2,100,000 over Meridian's 1,693,072) but remains a heuristic for lumpy candidates. Benefits require maps that include the **enabling change** most omit, baselines measured before the change, single claimants, honesty about cash-releasing versus capacity, and sustainability treated as constraint or as value but never confused between them. And the discipline closes with stopping: continuation depends only on remaining cost against remaining benefit — 780,000 of value for 900,000 of spend is a stop, whatever the 1,800,000 already gone suggests — made possible by kill criteria agreed in advance and gates that have a real stop option. Domain 3 supplies the governance and decision rights this domain assumes; Domain 16 measures what it promised.

03

DOMAIN 3

Governance, Organization and Decision Rights

Group: Leading projects (Domain 3 of 4 in Part One). **Target:** ~70 pages. **Binds to:** the PCI Book Pattern Specification and the shared registries ([docs/books/registries/](#)). This domain continues the **Meridian Care Records** programme from Domains 1 and 2 into the structures that decide things, and supplies the governance latency formula used later by Domain 6 (schedule), Domain 7 (change control) and Domain 8 (risk escalation). British English; USD (+SAR where useful, indicative [USD 1 ≈ SAR 3.75](#)).

— Why this domain exists

Domain 1 established what a project leader is answerable for; Domain 2 established how work gets chosen and whether the promise was honest. Both assumed something this domain has to supply: **that someone, somewhere, can actually decide.**

Governance is where most delivery failures are finally located, and almost never in the terms practitioners use for it. Post-project reviews record "poor governance" as though it were a culture problem. It is usually a design problem, and a measurable one: authority set at the wrong level, committees meeting on a cadence that cannot keep up with the work, escalation paths whose total latency nobody has ever added up, and decisions that are made but not recorded, so they are made again. Each of those has an arithmetic consequence, and this domain computes it. That is the domain's central claim: **governance is a delivery variable with a price, not a compliance overhead** — and a leader who cannot price it will be asked to accept a structure that guarantees delay.

The domain proceeds from purpose to mechanism. KA 3.1 asks what governance is for, and separates it from management, assurance and administration — the three things it is most often confused with — then works structures across organisational forms and the specific problem of governing iterative delivery. KA 3.2 covers the two roles that do the deciding: the sponsor, whose obligations are almost always understated, and the steering body, whose failure modes are predictable; it closes with decision authorities and thresholds, and prices the cost of setting them too low. KA 3.3 builds the assurance and gate machinery, designs escalation as a timed pathway rather than an organisation chart, and insists on the decision record — because an unrecorded decision is not a decision, it is a memory.

Learning objectives. After this domain a candidate can: state what governance is for and distinguish it from management and assurance; describe governance structures across functional, matrix, projectised and multi-party organisational forms, and identify the accountability weakness of each; govern iterative and hybrid delivery without either abandoning control or destroying cadence; specify the sponsor role as a set of testable obligations; diagnose the standard steering-committee failure modes; **compute the expected latency of a committee decision from its meeting**

interval and paper lead time, and use it to price a governance design; set delegation thresholds and demonstrate arithmetically when a threshold is too low; design stage gates, compute whether a gate is worth its elapsed time, and state its breakeven detection rate; distinguish the three lines of assurance and stop them duplicating; design an escalation path with a stated total latency and an out-of-cycle mechanism; maintain an auditable decision record; and govern AI-assisted governance analysis without delegating any decision to it.

The master programme. Meridian Care Records continues from Domains 1 and 2 — the clinical-records rollout to **40 clinics**, benefits **USD 685,440** per year at the realistic 70 % adoption, and the **cost of delay of USD 14,280 per week** derived in Domain 1. That last figure is what makes this domain quantitative: every week a decision waits has a price, and Meridian's governance design spends it without ever appearing in a budget line.

KNOWLEDGE AREA 3.1

Governance models

Topics: 3.1.1 what governance is for · 3.1.2 structures across organisational forms · 3.1.3 governance in agile and hybrid environments.

3.1.1 What governance is for

Definition. Project governance is the set of **decision rights, accountabilities and information flows** through which an organisation directs and controls a project: who may decide what, on whose authority, on what information, by when, and how the decision is recorded.

Every word of that carries weight, and the definition earns its keep mainly by exclusion. Governance is **not management** — management runs the work inside the authority granted; governance grants, bounds and withdraws that authority. It is **not assurance** — assurance forms an independent opinion on whether the work is likely to succeed; governance decides what to do about the opinion. And it is **not administration** — a portal, a report template and a monthly slide pack are governance artefacts, and an organisation can have all of them and no governance at all, because no one in the chain can actually decide anything. The confusion is not academic: the commonest remedy applied to a struggling project is more reporting, which adds administration, consumes the scarce attention of the people who could decide, and leaves the decision rights exactly where they were.

What governance exists to produce. Four things, and a governance design should be testable against each:

1. **Decidability.** Every decision the project will face has exactly one accountable decision-maker with sufficient authority. A decision with two owners is undecidable; one with none is unowned. This is a countable property, and KA 3.2 counts it.
2. **Timeliness.** The decision arrives while it can still change the outcome. This is the property most governance designs never examine, and KA 3.2 makes it computable.
3. **Legitimacy.** The decision is made by someone whose authority the organisation recognises, on information it can rely on, through a process it accepts — so the decision sticks rather than being relitigated by whoever dislikes it.
4. **Traceability.** The decision, its basis and its date can be reconstructed afterwards. This is the property that converts a decision from a memory into an institutional fact (3.3.4).

The two failure directions. Governance fails by being too heavy or too light, and both are common in the same organisation at once — heavy on the small decisions that are easy to control and light on the large ones that are uncomfortable. Heavy governance shows as latency: escalation for decisions that could have been delegated, agendas too full to reach the difficult item, and a reporting burden that consumes the delivery capacity it was meant to protect. Light governance shows as drift: decisions made by whoever is present, scope and budget changing without a decision-maker, and a project that no one can stop because no one is entitled to.

The professional discipline is therefore **proportionality, designed deliberately**. A governance structure is a design artefact with inputs (the project's value, novelty, risk, reversibility and external exposure) and outputs (tiers, thresholds, cadences, gates). Copying the structure used for the last programme is not proportionality; it is inheritance.

Governance and the leader's accountability. Domain 1's distinction returns here with force. Governance can allocate responsibility — the doing — freely, and should. It cannot dissolve the project leader's **obligation to answer**. A leader who says "the steering committee decided" has described a fact, not a defence: the obligations to have framed the decision honestly, to have surfaced the material information, to have stated the recommendation and the risk of the alternative, and to have escalated in time all remain the leader's, and every one of them is assessable afterwards from the decision record.

3.1.2 Structures across organisational forms

The same governance intent produces different structures depending on where authority sits in the host organisation, and each form has a characteristic weakness a leader should assume is present.

Functional. Resources and authority stay with functional line managers; the project leader coordinates. Characteristic weakness: the leader has accountability without authority over the people doing the work, so every resource conflict is an escalation. The governance countermeasure is explicit, written resource commitments with named individuals and dated availability, agreed by the functional manager and treated as breached when unilaterally changed — not a request, a commitment.

Matrix (weak, balanced, strong). Authority is shared on a spectrum. Characteristic weakness: **dual reporting without a stated precedence rule**, which does not split authority evenly but transfers it to whichever manager escalates harder. The countermeasure is a precedence rule written in advance for the specific conflicts that will arise — who wins on priority of work, on technical approach, on performance assessment — because a matrix without one is a functional organisation that holds project meetings.

Projectised. The project owns its resources and the leader holds full authority. Characteristic weakness: isolation from the enterprise — locally optimal decisions, divergent standards, and a benefits chain that ends where the project ends, precisely the failure Domain 2's benefits map addresses. The countermeasure is enterprise representation in the governance body with a real veto on the standards that matter, and a receiving-organisation owner for each benefit.

Multi-party (joint ventures, consortia, public-private). Two or more organisations with distinct interests, and often distinct legal duties. Characteristic weakness: **governance by unanimity**, which is indistinguishable from an inability to decide, particularly under stress when interests diverge. The countermeasures are structural and belong in the agreement, not in a terms of reference: a defined majority for defined classes of decision, a reserved-matters list that genuinely requires unanimity and is kept short, a deadlock-breaking mechanism with a deadline, and a single integrating authority for day-to-day direction. PFL-AI Domain 11 treats the risk-allocation side of the same problem; the

governance side is that a structure where any party can stop everything will eventually be used that way.

Programme and portfolio layers. Where a project sits inside a programme inside a portfolio, each layer adds a tier of potential escalation — and, unmanaged, adds its latency to every decision that travels up. KA 3.3 computes what that costs and why the design rule is **the fewest tiers a decision must legitimately pass through**, with the others informed rather than consulted.

A note on outsourced delivery. Where a supplier delivers, governance has to span the contract boundary, and the boundary is where decision rights are most often left undefined. The specific questions to settle before mobilisation: who may approve a change (and at what value on each side — the two thresholds are rarely equal and the mismatch is where change control fails); whose governance body is authoritative when they disagree; what the supplier is contractually obliged to escalate and within what period; and what information the client is entitled to, in what form, on what cadence. Domain 7 handles the commercial mechanics; the governance point is that a contract that does not allocate decision rights has not allocated the work.

3.1.3 Governance in agile and hybrid environments

Iterative delivery does not remove the need for governance; it changes what governance is asked to decide, and a structure designed for sequential delivery applied unchanged to iterative delivery fails in a specific, predictable way.

What changes. In sequential delivery, governance approves a plan and then controls variance against it. In iterative delivery the increment is the unit of decision, the backlog is the plan, and the governance question is not "is the project on plan?" but **"is the value being produced worth the next increment's cost, and is the direction still right?"** That converts governance from variance control to **continuation decisions at a cadence** — which is a more honest posture, and also a more demanding one, because it requires a decision-maker who can say "stop" repeatedly rather than once at a gate.

The characteristic failure. A monthly steering committee governing two-week sprints is always behind, and worse, its latency (computed in KA 3.2) exceeds the cycle it is governing — so the team either waits, destroying flow, or proceeds and seeks retrospective approval, destroying the governance. Both outcomes are attributed to the team. Neither is the team's doing.

The design response, in four moves:

- **Delegate inside a bounded envelope.** The team decides freely within stated bounds — scope within the agreed increment, technical approach within the architecture, spend within a period budget — and the governance body sets and reviews the envelope rather than the decisions inside it. This is the single highest-value governance decision in iterative delivery, and KA 3.2 prices it.
- **Match cadence to the work.** Governance interaction at the increment's rhythm — increment reviews as the decision point, with the committee attending rather than being reported to.
- **Govern outcomes, not activity.** Working software or a released service demonstrated, with the benefit measure attached (Domain 2's map), rather than percentage-complete against a plan that the method deliberately does not have.
- **Keep the gates that carry real optionality.** Funding tranches, go-live authorisation and regulatory approvals remain genuine gates because they are genuinely irreversible. What should be removed are the gates that merely re-approve a decision already taken.

Hybrid honestly. Most real programmes are hybrid — an iterative build inside a sequential infrastructure or regulatory frame — and the governance failure is to apply one model to both parts. Meridian is exactly this shape: an iteratively developed records application, a sequential clinic rollout

with immovable estate and training dependencies, and a regulatory approval that is a genuine gate. The workable design governs each part on its own terms and holds a **single integrated view of the whole** for the decisions that span them — which is the topic of Domain 4.

AI in this KA

AI is useful in governance analysis and dangerous in governance decisions, and the line is bright.

Where it earns its place. Reading a set of terms of reference, delegation schedules and contract schedules against each other and listing the decisions with no named owner, two owners, or conflicting thresholds — a document-comparison task at which it is fast and thorough, and one humans do badly because it is tedious. Extracting the decision log from meeting minutes into a structured register, flagging items recorded without an owner or a date. Summarising a decision's history for a board paper. Modelling the latency of a proposed governance design (KA 3.2's arithmetic) across alternative cadences, which is deterministic and verifiable.

Where it must not go. No decision, and no recommendation presented as a decision. Governance authority is conferred on accountable people; it cannot be exercised by a system that cannot be answerable, and an organisation that lets a model's output stand unchallenged as a governance conclusion has created accountability without a holder — the exact defect Domain 1 identified.

The governing principle, applied. AI proposes; the professional verifies, decides and remains accountable. Concretely: every AI-produced governance analysis is reviewed by the named accountable person before it reaches a decision body, its inputs are stated, its conclusions are independently checked on the material items, and the decision record shows a human decision-maker — never a tool — as the author of the decision.

■ KEY TERMS — KA 3.1

Term	Meaning
Governance	The decision rights, accountabilities and information flows through which an organisation directs and controls a project.
Decision right	The authority to make a specified class of decision, bounded by value, scope and reversibility.
Decidability	The property that every decision has exactly one accountable decision-maker with sufficient authority.
Reserved matter	A decision class that requires the highest (often unanimous) authority; kept deliberately short.
Bounded envelope	Stated limits within which a delivery team decides freely, reviewed rather than approved decision by decision.
Precedence rule	The pre-agreed rule for which authority prevails in a matrix conflict.
Governance artefact	A report, portal or template — evidence of governance, never a substitute for decision rights.

■ SAMPLE MCQS — KA 3.1

MCQ 3.1-A [3.1.1 · [Comprehension](#)] Which statement best distinguishes governance from management? - A. governance is performed by senior people and management by junior people - B. governance grants, bounds and withdraws the authority within which management runs the work

- C. governance is concerned with reporting and management with delivery - D. governance applies to programmes and management to projects

Rationale: The distinction is about authority, not seniority, reporting or scale (3.1.1). C describes governance artefacts, which an organisation can have in full while having no governance.

MCQ 3.1-B [3.1.2 · Analysis] A balanced matrix has no written precedence rule for conflicts between project and functional priorities. The most likely consequence is that authority: - A. is shared evenly, as intended - B. transfers to whichever manager escalates harder - C. defaults to the project leader - D. defaults to the steering committee automatically

Rationale: Undefined precedence does not split authority; it awards it to escalation behaviour (3.1.2). The countermeasure is a precedence rule written before the conflict arises.

MCQ 3.1-C [3.1.3 · Application] A monthly steering committee governs two-week sprints. The predictable failure is that: - A. the team will produce lower-quality increments - B. governance latency exceeds the cycle being governed, so the team either waits or proceeds and seeks retrospective approval - C. the committee will meet too often - D. the sprints must be lengthened to a month

Rationale: The mismatch is one of latency against cycle time (3.1.3); the design response is a bounded envelope and cadence matching, not slower delivery.

MCQ 3.1-D [3.1.2 · Analysis] In a three-party consortium, the governance provision most likely to produce paralysis under stress is: - A. a defined majority for defined decision classes - B. a short reserved-matters list - C. unanimity for all substantive decisions - D. a deadlock-breaking mechanism with a deadline

Rationale: Unanimity is indistinguishable from an inability to decide once interests diverge (3.1.2); the other three are the countermeasures.

■ SELF-CHECK — KA 3.1

1. Name the four things a governance design must produce. — Decidability, timeliness, legitimacy, traceability.
2. Why is "more reporting" the wrong remedy for a struggling project's governance? — It adds administration and consumes decision-makers' attention while leaving decision rights unchanged.
3. What does a project leader still owe after a steering committee decides badly? — The obligations to have framed the decision honestly, surfaced the material information, stated the recommendation, and escalated in time — all assessable from the record.

KNOWLEDGE AREA 3.2

Sponsorship and steering

Topics: 3.2.1 the sponsor role · 3.2.2 steering committees that work · 3.2.3 decision authorities and thresholds.

3.2.1 The sponsor role

Definition. The sponsor is the individual accountable for the project's **business outcome** — for its being worth doing, remaining worth doing, and being adopted by the organisation that must use it.

The project leader is accountable for delivery within the mandate; the sponsor is accountable for the mandate.

That is a demanding role and it is routinely treated as an honorific. The remedy is to state it as **testable obligations**, each of which can be evidenced or found absent:

Obligation	Evidence it is being met
Own the business case	The sponsor can state the benefit, its measure and its owner without notes (Domain 2).
Secure and defend funding	Funding decisions carry the sponsor's name; funding challenges are answered by the sponsor, not the project leader.
Decide within authority, promptly	Decisions escalated to the sponsor have recorded dates in and out; the interval is monitored.
Own the receiving organisation's readiness	The enabling change is resourced and progressing (Domain 2's omitted column).
Resolve conflicts between functions	Escalations that are genuinely inter-functional stop at the sponsor rather than travelling upward.
Hold the project leader to account, and support them	Recorded, dated performance conversations; and visible backing when the project is under pressure.
Be willing to stop it	Kill criteria agreed in advance and revisited at gates (Domain 2, KA 2.4).

The sponsor failure modes, in descending order of frequency and each with its detection test. The absent sponsor — a name on a chart, no diary time; test: the count of sponsor decisions in the last quarter, and their turnaround. The delegating sponsor — a deputy attends everything, so authority is in the room but accountability is not; test: whether the last three material decisions were made by the sponsor or reported to them afterwards. The advocate sponsor — committed to the project rather than the outcome, and therefore unable to stop it; test: whether they can state a condition under which they would recommend stopping (Domain 2's escalation of commitment, at the governance level). The operational sponsor — drawn into managing delivery, which vacates the role that only they can hold; test: whether the sponsor's contributions in the last three meetings were about outcome and mandate, or about task sequence.

What a project leader does with a weak sponsor. This is a real and common situation, and the professional response is neither to complain nor to absorb the gap silently — absorbing it is attractive, because it feels like competence, and it is how a leader ends up accountable for a mandate they were never given. The workable response is procedural: write the decisions needed and their dates; send them with a recommendation and a stated consequence of non-decision; record the non-decision when it occurs; escalate the pattern rather than the individual instance once it recurs; and keep the record. That is not politics. It is the mechanism by which an organisation's governance failure becomes visible to the organisation while there is still time to fix it — and, incidentally, the only defensible position for the leader afterwards.

3.2.2 Steering committees that work

A steering committee exists to make the decisions that exceed the project leader's authority, in time for them to matter. Almost every observed dysfunction traces to one of five design faults.

Fault 1 — membership without authority. Attendees who must consult before agreeing. The committee's effective authority is the minimum of its members' authority on the decision in hand, not

the maximum, so one under-empowered member on a decision class blocks it. Test: for each decision class in the delegation schedule, name the member who can commit.

Fault 2 — membership too large. Beyond roughly six to eight decision-makers the body becomes a briefing audience: contributions become positional, dissent becomes private, and the difficult item gets deferred. Attendance is not membership — the separation of members (who decide) from attendees (who inform) is the cheapest available fix.

Fault 3 — the agenda consumed by reporting. Status occupies the meeting and the decisions arrive in the last ten minutes, which is why the reliable diagnostic of committee health is not attendance but the **share of agenda time spent on decisions**. Status is read in advance; the meeting decides.

Fault 4 — cadence mismatched to the work. Treated quantitatively in 3.2.3 — and it is the fault with the largest measurable price.

Fault 5 — no capacity model. A committee has finite throughput, and a design that routes more decisions to it than it can hear does not slow decisions down evenly: the items that get deferred are the contentious ones, which are the ones that needed the decision. The capacity arithmetic is elementary and almost never done.

WORKED EXAMPLE

Worked example 3.2.2 — does Meridian's steering committee have the capacity to govern?

1. **Setup.** Meridian's steering committee meets every **4 weeks** (**13** meetings a year) and can handle **8** substantive agenda items per meeting. Annual demand: **36** escalated change requests (the count above the current delegation threshold, from 3.2.3), **26** standing reports (two per meeting), and **15** gate and assurance decisions.
2. **Formula.** Capacity = meetings × items per meeting. Utilisation = demand ÷ capacity.
3. **Substitution.** Capacity $13 \times 8 = 104$. Demand $36 + 26 + 15 = 77$. $77/104$.
4. **Result.** Capacity **104** item-slots; demand **77**; utilisation **74.0 %**.
5. **Interpretation.** Seventy-four per cent looks comfortable and is not, for two reasons a leader should be able to state. First, demand is **not uniform** — changes and gate decisions cluster around phase boundaries, so the peak months exceed capacity while the average does not, and it is in the peak months that decisions are deferred. Second, the 26 standing reports consume a quarter of the committee's total capacity producing no decisions at all (Fault 3): moving them to pre-reading releases 26 slots and drops utilisation to **49.0 %**, which is the headroom a contentious item needs. Raising the delegation threshold as 3.2.3 recommends removes a further 12 items, taking demand to **65** and utilisation to **62.5 %** on the original agenda design. The professional point is that committee capacity is a designed quantity, and a governance structure proposed without this arithmetic has not been designed.

3.2.3 Decision authorities and thresholds

The structure. A delegation schedule states, for each decision class, the level of authority required — typically as a monetary threshold, but the better schedules use three dimensions: **value** (how much), **reversibility** (how hard to undo), and **externality** (who outside the project is affected). A cheap, irreversible, externally visible decision may deserve more authority than an expensive, reversible, internal one; a schedule that reads only on value cannot express that.

Governance latency: the formula this domain contributes. Committees meet periodically and require papers in advance, and those two facts alone determine how long a decision waits. Let

M = the meeting interval (e.g. 4 weeks for a monthly committee)
 L = the paper lead time – how far before the meeting submissions close

For a decision that arises at a uniformly random point in the cycle:

$$E[\text{wait}] = M/2 + L$$

The derivation is elementary. A decision arising more than L before the next meeting makes that meeting, and waits on average $M/2$ of the remaining interval; one arising inside the closed window misses it and waits for the following meeting. Averaging over arrival times gives half the meeting interval plus the whole paper lead time.

Why the formula matters more than it looks. It settles a question organisations habitually get wrong. The two available levers do **not** have equal effect: cutting the paper lead time by one week saves **one full week** of expected wait, while shortening the meeting interval by one week saves only **half a week**. Administrative deadlines, which feel unchangeable and cost nothing to change, are twice as powerful as meeting frequency, which is expensive and resisted. And a governance design's latency can be computed at design time, from two numbers, before anyone has waited.

Meridian's steering committee: $M = 4$, $L = 2$, so $E[\text{wait}] = 4/2 + 2 = 4$ **weeks**. Note what that says — a monthly committee with a two-week paper deadline imposes an expected wait of a full month, not the fortnight most people assume. At Domain 1's cost of delay of 14,280 per week, one escalated decision on the critical path costs **USD 57,120** in delay alone, before anyone in the room disagrees about anything.

The paper deadline is twice the lever the meeting interval is — and the cheaper one to move




Fig 3.2.1 — Governance latency and its two levers

Setting the threshold. The threshold question is an economic one and can be answered as one. Setting a threshold too high risks decisions taken with insufficient authority or scrutiny; setting it too low guarantees latency on decisions that did not need it. The second cost is certain and computable; the first is probabilistic. Comparing them is the whole of the analysis.

WORKED EXAMPLE

Worked example 3.2.3 — Meridian's delegation threshold, priced.

1. **Setup.** Meridian's project leader may approve changes up to **USD 10,000**; everything above goes to the steering committee ($E[\text{wait}] = 4$ weeks, cost of delay **14,280** per week). Year-one change requests, by value: $\leq 10,000 \rightarrow 24$; **10,001-25,000** $\rightarrow 12$ (average value **17,000**); **25,001-100,000** $\rightarrow 15$; **100,001-500,000** $\rightarrow 7$; $> 500,000 \rightarrow 2$. Total **60**. Reviewed history shows **25 %** of escalated changes sit on the critical path and therefore convert their wait into programme delay; the rest do not. Should the threshold move to 25,000?
2. **Formula.** Escalated count $\times$ critical-path share $\times E[\text{wait}] \times$ cost of delay, evaluated at each threshold. Then the worst-case cost of delegating the affected band.
3. **Substitution.** At 10,000: escalated $60 - 24 = 36$, delaying $36 \times 0.25 = 9$, each costing $4 \times 14,280 = 57,120$. At 25,000: escalated $36 - 12 = 24$, delaying $24 \times 0.25 = 6$.
4. **Result.** Annual governance delay cost **USD 514,080** at a 10,000 threshold and **USD 342,720** at 25,000 — a saving of **USD 171,360** a year. The band being delegated is 12 changes worth **204,000** in total. Even in the **worst imaginable case** — the delegate decides every one of the twelve wrongly and destroys **40 %** of the value each time — the loss is $12 \times 17,000 \times 0.40 = \text{USD } 81,600$, leaving the delegation ahead by **USD 89,760**. The threshold at which the two are equal is a value destruction of $171,360/204,000 = 84 \%$ of every delegated decision.
5. **Interpretation.** The escalation cannot be justified at any plausible error rate — the delegate would have to destroy 84 % of the value of every delegated decision for the committee's involvement to break even, and a delegate who did that would be removed for reasons unrelated to thresholds. That is the general shape of the result, and it is why over-centralised delegation schedules are so common and so expensive: **the cost of escalation is certain, recurring and invisible, while the cost of a delegated error is uncertain, occasional and highly visible.** Organisations optimise against the visible cost. Three professional cautions, however. The calculation prices delay, not scrutiny — where a decision class carries irreversibility or external exposure the value test is the wrong test, which is why the schedule reads on three dimensions and not one. The 25 % critical-path share is an assumption from history and must be stated as one, since the conclusion is proportional to it — though note that it would have to fall below $81,600/(12 \times 57,120) = 11.9 \%$ before the worst case even competed. And a raised threshold requires the delegate to have the information to decide, which is a real prerequisite: delegation without information is abdication.

AI in this KA

Where it earns its place. Modelling latency across candidate governance designs — the arithmetic above, over dozens of cadence and threshold combinations, is deterministic and a natural fit. Classifying a change-request history into value bands and testing whether a proposed threshold would have altered any actual outcome — a genuinely useful retrospective test, and one nobody has time to do by hand. Drafting a delegation schedule from an existing one plus a set of stated changes,

for human review. Detecting the decidability defects of 3.1.1 across a document set — decisions with no owner, two owners, or thresholds that conflict between the contract and the terms of reference.

Where it must not go. Setting a threshold, which is a risk-appetite decision belonging to the sponsor and the accountable body. Estimating the critical-path share or the error rate from plausibility rather than data — a model asked for these will produce confident numbers with no provenance, and they will then drive a real decision. Nor should its latency model be trusted unverified: the formula is two operations, and a leader who cannot reproduce $M/2 + L$ on paper should not be quoting it.

Verification, concretely. Reproduce every number by hand, state the assumptions with their sources, and put the sensitivity in the paper — the breakeven error rate and the breakeven critical-path share, not merely the point estimate, because those are what tell a board whether the recommendation is robust or fragile.

■ KEY TERMS — KA 3.2

Term	Meaning
Sponsor	The individual accountable for the project's business outcome and mandate.
Steering committee	The body that makes decisions exceeding the project leader's authority, in time for them to matter.
Delegation schedule	The statement of authority by decision class, on value, reversibility and externality.
Governance latency	The expected wait for a committee decision: $E[\text{wait}] = M/2 + L$.
Meeting interval (M)	The period between meetings of a decision body.
Paper lead time (L)	How far before a meeting submissions close.
Committee capacity	Meetings per year × substantive items per meeting; compared against decision demand.
Member vs attendee	Members decide; attendees inform. Conflating them enlarges the body and reduces its authority.

■ SAMPLE MCQS — KA 3.2

MCQ 3.2-A [3.2.3 · Application] A committee meets every 6 weeks and closes papers 2 weeks before each meeting. The expected wait for a decision arising at a random point in the cycle is: - A. 3 weeks - B. 5 weeks - C. 6 weeks - D. 8 weeks

Rationale: $E[\text{wait}] = M/2 + L = 6/2 + 2 = 5$ weeks (3.2.3). A counts only half the interval and omits the paper lead time; C is the interval itself; D adds the whole interval to the lead time.

MCQ 3.2-B [3.2.3 · Analysis] A committee meets every 4 weeks with a 2-week paper lead time. Two options are on the table: cut the paper lead time by one week, or cut the meeting interval by one week. Which reduces expected latency more, and by how much? - A. cutting the meeting interval, by 1.0 week - B. cutting the paper lead time, by 1.0 week — twice the 0.5-week saving from the interval - C. both, equally, by 1.0 week - D. both, equally, by 0.5 weeks

Rationale: $E[\text{wait}] = M/2 + L$ is 4.0 weeks initially; cutting L to 1 gives 3.0 (saves 1.0), cutting M to 3 gives 3.5 (saves 0.5) — a one-week cut in the paper lead time always saves a full week, a one-week cut in the meeting interval only half of one (3.2.3). The administrative deadline is also the cheaper lever to move.

MCQ 3.2-C [3.2.3 · Evaluation] Raising Meridian's threshold from 10,000 to 25,000 saves 171,360 a year in delay. The band delegated comprises 12 changes worth 204,000. The strongest argument for the change is that: - A. the project leader is experienced - B. the delegate would have to destroy 84 % of the value of every delegated decision for the escalation to break even - C. the committee is too busy - D. 12 changes is a small number

Rationale: The decisive argument is the breakeven value-destruction rate, which makes the comparison explicit and quantitative (3.2.3). A and D are assertions; C is a capacity argument, which is real but secondary.

MCQ 3.2-D [3.2.2 · Analysis] A steering committee's effective authority on a given decision class is best described as: - A. the authority of its chair - B. the highest authority among its members - C. the lowest authority among the members whose agreement is required - D. the authority delegated to the project leader

Rationale: One member who must consult before agreeing blocks the class, so the binding constraint is the minimum, not the maximum (3.2.2, Fault 1).

MCQ 3.2-E [3.2.1 · Analysis] The most reliable test of whether a sponsor is an advocate sponsor rather than an effective one is whether they: - A. attend every steering committee - B. can state a condition under which they would recommend stopping the project - C. defend the project's funding - D. know the project's schedule in detail

Rationale: An advocate sponsor is committed to the project rather than the outcome and therefore cannot stop it (3.2.1). C is an obligation of the role; D is closer to the operational failure mode.

■ SELF-CHECK — KA 3.2

1. State the governance latency formula and both of its levers. — $E[\text{wait}] = M/2 + L$; shorten the paper lead time (a full week per week cut) or the meeting interval (half a week per week cut).
2. Why do organisations systematically set delegation thresholds too low? — The cost of escalation is certain, recurring and invisible; the cost of a delegated error is uncertain, occasional and highly visible, so optimisation runs against the visible cost.
3. What must accompany a raised threshold for it to be delegation rather than abdication? — The information, criteria and support the delegate needs to decide, plus a record of what they decided.

KNOWLEDGE AREA 3.3

Assurance, gates and escalation

Topics: 3.3.1 stage gates · 3.3.2 assurance lines · 3.3.3 escalation design · 3.3.4 auditability and the decision record.

3.3.1 Stage gates

Definition. A stage gate is a point at which continuation requires an explicit decision by a named authority against stated criteria, with authority to stop, hold or redirect as well as proceed. A checkpoint without the authority to stop is a milestone with a meeting attached.

What a gate is for. Gates exist to convert **irreversibility into optionality**: before committing the next tranche of cost, an organisation buys the chance to stop or change direction while stopping is

still cheap. It follows that gates belong where irreversibility genuinely steps up — before major commitment, before build, before go-live, before regulatory submission — and not at regular calendar intervals, which is how they proliferate into re-approval ceremonies.

Gate criteria. Criteria must be set **in advance**, be **objectively assessable**, and cover both the work and the decision: is the deliverable adequate; is the business case still valid (Domain 2's benefits and assumptions, re-tested rather than restated); are the risks acceptable and owned; is the receiving organisation ready; and is the plan for the next stage credible? Criteria written after the evidence is available are not criteria.

The honest gate. A gate that has never held or stopped anything is not evidence of good delivery, it is evidence of a gate that does not function — and its cost is being paid for nothing. The pattern to watch for is the **conditional pass**: proceed subject to conditions, which are then neither tracked nor enforced. A conditional pass is a genuine and useful instrument, and it is only that if the conditions have owners, dates and a stated consequence of non-completion, verified at the next gate.

Gates cost time, and time has a price. This is the part of gate design that is almost never computed, and it decides whether the gate is worth having.

WORKED EXAMPLE

Worked example 3.3.1 — is Meridian's design gate worth its elapsed time?

- Setup.** Meridian's design-completion gate requires **USD 45,000** of review effort and **6 weeks** of elapsed time (cost of delay **14,280** per week). Historical evidence on comparable work: a material design defect is present with probability **0.30**; the gate detects it with probability **0.80**; a defect corrected at design costs **120,000**; the same defect found in build costs **900,000**.
- Formula.** Expected total cost with the gate = review + delay + $P(\text{defect}) \times [P(\text{detect}) \times \text{design-fix} + P(\text{miss}) \times \text{build-fix}]$. Without the gate = $P(\text{defect}) \times \text{build-fix}$.
- Substitution.** Without: $0.30 \times 900,000$. With: $45,000 + 6 \times 14,280 + 0.30 \times (0.80 \times 120,000 + 0.20 \times 900,000)$.
- Result.** Without the gate **USD 270,000**. With the gate $45,000 + 85,680 + 82,800 = \text{USD } 213,480$. The gate is worth **USD 56,520**.
- Interpretation.** The gate pays — and the useful output is not the 56,520 but the two **breakeven points** it implies, because those are what a leader negotiates with. Holding everything else, the gate stops paying once its elapsed time reaches $(270,000 - 45,000 - 82,800) / 14,280 = \mathbf{9.96 \text{ weeks}}$: a gate that takes ten weeks destroys the value it exists to protect, which is the arithmetic behind the familiar and usually unquantified complaint that assurance has become an obstacle. And at the actual 6 weeks, the gate needs a detection probability above **55.85 %** to be worth holding — so a review staffed by people who cannot competently detect the defect is worse than no review, because it costs the money and the delay and returns nothing. Both results generalise: **gate value is destroyed by elapsed time and by weak detection**, and a leader whose gate is under challenge should compute which of the two is the actual problem before defending or abandoning it.

3.3.2 Assurance lines

The three lines, and what each is for.

- **First line — management.** The project's own controls: reviews, testing, quality checks, reporting. Owned by the project leader, and the only line that can prevent a defect rather than detect it.
- **Second line — oversight function.** A PMO, risk or quality function providing independent challenge while remaining inside management's chain. Its value is comparability across the portfolio and pattern detection no single project can see.
- **Third line — internal audit (and external assurance).** Independent of management, reporting to the audit committee or equivalent, forming an opinion on whether the whole control system works.

The failure modes. Duplication — three lines asking the same questions, which multiplies cost and destroys the project team's willingness to engage. Gap — everyone assumes another line covers a risk. Capture — the second line drafts the plan it later assures, and so cannot challenge it; this is the most damaging and the least visible, because the assurance product still looks independent. Assurance as accountability transfer — a leader treats a favourable assurance opinion as a discharge of their own obligation; Domain 1's principle applies unchanged, an opinion is information and not a transfer of accountability.

The countermeasure is an **assurance map**: risks and controls down the side, lines across the top, and each cell marked as covered, not covered or duplicated — reviewed by the accountable authority. It is a half-day artefact that surfaces both gaps and duplication, and its absence is the usual reason a project is simultaneously over-assured and unassured.

Proportionality. Assurance effort should scale with novelty, value, irreversibility and external exposure — not uniformly, and not with organisational anxiety. A high-value, low-novelty, reversible project may need less assurance than a small, novel, irreversible, publicly visible one, and a regime that cannot express that will over-assure the first and under-assure the second.

3.3.3 Escalation design


The principle. Escalation is a **designed pathway with stated latency**, not an instruction to tell someone senior. A usable design states, for each escalation class: the trigger (objective, not "if concerned"), the destination, the decision required, the **time within which the decision will be made**, and the out-of-cycle mechanism if the ordinary cadence is too slow.

That fourth element is the one usually missing, and its absence has a computable cost.

WORKED EXAMPLE**Worked example 3.3.3 — the total latency of Meridian's escalation path.**

1. **Setup.** A decision requiring executive authority passes three tiers: the **project board** (meets every 2 weeks, papers close 1 week ahead), the **programme board** (every 4 weeks, papers 2 weeks ahead), and the **executive committee** (quarterly — every 13 weeks — papers 3 weeks ahead). Cost of delay **14,280** per week.
2. **Formula.** Total expected latency = Σ over tiers of $M/2 + L$. Cost = latency $\times$ cost of delay.
3. **Substitution.** Project board $2/2 + 1 = 2.0$; programme board $4/2 + 2 = 4.0$; executive committee $13/2 + 3 = 9.5$. Total $2.0 + 4.0 + 9.5$.
4. **Result.** **15.5 weeks** of expected latency for a single decision, costing **USD 221,340** in delay — before anyone in any of the three rooms disagrees with the recommendation.
5. **Interpretation.** Two hundred and twenty-one thousand dollars is the price of an organisation chart, and it is invisible: it appears in no budget, is attributed to no decision, and is generally described afterwards as the project having been slow. Notice the distribution — the quarterly committee alone accounts for **9.5 of the 15.5 weeks (61 %)** and **USD 135,660**, which is what makes "add the executive committee to the approval path" such an expensive sentence when spoken casually in a governance review. Two redesigns follow directly. **Reduce the tiers a decision must legitimately pass:** if the programme board can decide with the executive committee informed, latency falls to **4.0 weeks** and **USD 57,120** — a saving of **USD 164,220**, or **74.2 %**, from a change that removes no scrutiny that anyone can name. **Add an out-of-cycle mechanism:** a written-resolution procedure with a five-working-day turnaround makes the single-tier path **1.0 week** and **USD 14,280**, a **93.5 %** reduction against the original. The general rule this supports: **count the tiers, add the latency, price it, and then justify each tier against its cost** — which is a conversation governance reviews almost never have, because until the latency is added up there is nothing to weigh the tier against.

Count the tiers, add the latency, price it — then justify each tier against its cost



Each tier's wait is $M/2 + L$: $2/2+1 = 2.0$ · $4/2+2 = 4.0$ · $13/2+3 = 9.5$ — priced at USD 14,280 per week

Fig 3.3.1 — The price of an escalation path

Designing the trigger. Objective triggers only: a forecast variance beyond a stated tolerance, a risk exposure above a stated threshold, a decision required above the delegated authority, a dependency breached. "Escalate if concerned" produces two failure modes at once — late escalation by those who

fear it and constant escalation by those who fear the alternative — and both are then treated as individual judgement problems rather than as the design defect they are.

The escalation culture problem. Escalation is only a functioning mechanism if using it is safe. Where escalation is read as failure, it happens late, which is precisely when the decision can no longer help. The countermeasures are governance ones and belong to the sponsor and the committee: respond to early escalation visibly and well; separate the escalation of an issue from any judgement about the person raising it; and track the **lead time** of escalations — how far before the impact they arrive — because a shortening lead time is the earliest available signal that the mechanism is decaying, and it appears in no standard report.

3.3.4 Auditability and the decision record

The principle. A decision that is not recorded has not been made — it has been remembered, and memories diverge exactly when the stakes rise. The decision record is the mechanism that converts a meeting into an institutional fact.

What a decision record must contain, per decision, and it is short enough that its absence is never a resourcing problem: a unique reference; the date; the decision-maker by name and role (never a committee alone — a body records the decision, a person is accountable for it); the decision, in words that permit only one reading; the options considered and why the chosen one was chosen; the information relied on, referenced to its version; the conditions attached, with owners and dates; and the review date if the decision is provisional.

Why the "information relied on, referenced to its version" line matters more than it looks. When a decision is examined afterwards — in a lessons review, a dispute, an audit or an inquiry — the question is almost never "was the decision right?" but "**was it reasonable on what was known at the time?**" Only a versioned reference can answer that, and its absence converts a defensible decision into an indefensible one at exactly the moment the defence is needed.

Two governance defects the record exposes, and this is the practical case for keeping it well. The **re-decided decision:** the same question arriving at the same body a third time is a symptom of either an unrecorded decision or an illegitimate one, and the log makes the recurrence visible where minutes do not. And the **decision nobody made:** a change that took effect without an entry, which is how scope and budget move without anyone having decided — the most common finding of a governance audit and the hardest to see from inside.

RACI and its integrity check. Responsibility assignment matrices — Responsible, Accountable, Consulted, Informed — are the standard artefact for decision rights, and their standard failure is being drafted, circulated and never checked. The checks are countable, and a matrix that fails any of them is not a matrix but a document: **exactly one A** per decision (two make it undecidable, none make it unowned); the A holds sufficient authority under the delegation schedule; C and R are distinguished (consultation is not agreement); and the count of C's per decision is bounded, since consultation is where latency accumulates invisibly.

WORKED EXAMPLE**Worked example 3.3.4 — auditing Meridian's decision-rights matrix.**

1. **Setup.** Meridian's matrix covers **12** decision classes across 7 roles. On review: **9** classes carry exactly one Accountable, **2** carry two, and **1** carries none.
2. **Formula.** Defect rate = classes failing the single-A test ÷ total classes.
3. **Substitution.** $(2 + 1)/12$.
4. **Result.** 3 of 12 classes are defective — a **25.0 %** defect rate.
5. **Interpretation.** A quarter of the programme's decision classes cannot be decided as documented, and the two failure types behave differently under stress, which is why they are worth distinguishing rather than totalling. The **two-A** classes will be decided — by whichever holder acts first, or after a delay while they reconcile, and the delay is at 3.2.3's committee rates. The **zero-A** class will not be decided at all: it will drift until it becomes an escalation, which is the worst of the available outcomes because it arrives late and without preparation. The test costs an hour and it is the highest-yield governance check available before mobilisation.

AI in this KA

Where it earns its place. Extracting a structured decision register from meeting minutes and flagging entries missing an owner, a date or a versioned information reference — a genuinely tedious task with a clear right answer. Running the RACI integrity checks above across a large matrix. Detecting re-decided decisions by clustering decision text across a long log, which is exactly the pattern humans miss because the recurrences are months apart. Testing gate criteria for assessability and flagging the unmeasurable ones. Modelling gate and escalation economics across alternative designs, as computed above.

Where it must not go. No gate decision, no assurance opinion, and no authorship of a decision record entry — the record must show the accountable human, and a record generated wholesale by a tool cannot evidence that a person applied judgement. Nor should a model's summary of a decision's history be relied on for a dispute or an audit without verification against the source documents: the summary is a convenience, the versioned source is the evidence.

Verification, concretely. Any AI-produced register or audit is checked against the source minutes on a sampled basis with the sample size stated; every flagged defect is confirmed by a human before it is reported as a finding; and the gate and latency arithmetic is reproduced by hand, because all of it is a handful of operations and none of it should ever be taken on trust.

■ KEY TERMS — KA 3.3

Term	Meaning
Stage gate	A continuation decision by a named authority against pre-set criteria, with power to stop, hold or redirect.
Conditional pass	Proceeding subject to conditions — a real instrument only if the conditions have owners, dates and consequences.
Three lines of assurance	Management controls; independent oversight inside management; independent audit outside it.
Assurance map	Risks and controls against assurance lines, marking coverage, gaps and duplication.
Assurance capture	An assurance function assuring work it helped produce, and therefore unable to challenge it.
Escalation class	A defined trigger, destination, decision, latency and out-of-cycle route.
Out-of-cycle mechanism	A written-resolution or delegated-authority route used when the ordinary cadence is too slow.
Escalation lead time	How far before impact an escalation arrives; a shortening trend signals a decaying mechanism.
Decision record	The versioned, attributable log that converts a decision from a memory into an institutional fact.
Single-A test	The check that each decision class has exactly one Accountable role.

■ SAMPLE MCQS — KA 3.3

MCQ 3.3-A [3.3.1 · Application] A gate costs 45,000 in review effort and 6 weeks of elapsed time at a delay cost of 14,280 per week. Expected remediation cost with the gate is 82,800; without the gate it is 270,000. The gate's net value is: - A. USD 141,480 - B. USD 56,520 - C. USD 85,680 - D. USD 187,200

Rationale: $270,000 - (45,000 + 85,680 + 82,800) = 56,520$ (3.3.1). A omits the delay cost; C is the delay cost alone; D omits the remediation cost.

MCQ 3.3-B [3.3.1 · Evaluation] For the gate above, the elapsed time at which it stops adding value is closest to: - A. 6 weeks - B. 8 weeks - C. 10 weeks - D. 19 weeks

Rationale: $(270,000 - 45,000 - 82,800)/14,280 = 9.96$ weeks (3.3.1) — the arithmetic behind the complaint that assurance has become an obstacle.

MCQ 3.3-C [3.3.3 · Application] Three tiers must approve a decision: 2-weekly with 1-week papers, 4-weekly with 2-week papers, and 13-weekly with 3-week papers. Total expected latency is: - A. 9.5 weeks - B. 15.5 weeks - C. 19.0 weeks - D. 6.0 weeks

Rationale: $(2/2+1) + (4/2+2) + (13/2+3) = 2.0 + 4.0 + 9.5 = 15.5$ (3.3.3). A is the top tier alone; D counts only half the intervals and omits the paper lead times.

MCQ 3.3-D [3.3.2 · Analysis] A PMO drafts the delivery plan and later provides second-line assurance on it. The defect is: - A. duplication of first-line controls - B. assurance capture — the function cannot challenge what it produced - C. a proportionality failure - D. an assurance gap

Rationale: Capture is the most damaging and least visible line failure, because the product still looks independent (3.3.2).

MCQ 3.3-E [3.3.4 · Comprehension] The decision-record field most often missing and most consequential when a decision is later examined is: - A. the date - B. the versioned reference to the information relied on - C. the decision-maker's role - D. the decision reference number

Rationale: The retrospective question is whether the decision was reasonable on what was known at the time, which only a versioned reference can answer (3.3.4).

MCQ 3.3-F [3.3.4 · Analysis] In a 12-class decision-rights matrix, 2 classes carry two Accountable roles and 1 carries none. Which statement is most accurate? - A. the defect rate is 8.3 % and only the zero-A class matters - B. the defect rate is 25.0 %; the two-A classes will be decided late or twice, and the zero-A class will drift until it becomes an escalation - C. the matrix is acceptable because 9 of 12 classes are correct - D. the defect rate is 25.0 % and all three classes fail identically

Rationale: $3/12 = 25.0 \%$, and the two failure types behave differently under stress, which is why they are distinguished rather than totalled (3.3.4).

■ SELF-CHECK — KA 3.3

1. What does a gate buy, and what does that imply about where gates belong? — Optionality against irreversibility; so gates belong where irreversibility steps up, not at calendar intervals.
2. Name the assurance failure mode that is hardest to see, and why. — Capture: the assurance product still looks independent while the function cannot challenge what it helped produce.
3. Which element of an escalation class is most often missing? — The stated time within which the decision will be made, and with it the out-of-cycle mechanism.

— Advanced topics — Domain 3

3.A.1 Governance under stress, and the recovery structure

Governance designed for steady state is tested in crisis, and the failures are consistent: the cadence is too slow for the decision rate, authority fragments as senior figures intervene individually, information becomes contested at the moment it must be relied on, and the decision record — never robust — collapses first, so that a fortnight later nobody can reconstruct why a choice was made.

A **recovery governance** design is a legitimate and different structure, and it should be prepared in advance rather than improvised: a single accountable decision-maker with materially raised authority; a short daily or twice-weekly decision cadence with a standing agenda of decisions rather than reports; a narrowed membership, since crisis governance fails from too many participants before it fails from too few; a single authoritative information source, with disputes about the data resolved before decisions rather than during them; and a decision log kept in the room, in real time, because reconstruction afterwards is precisely what will not be possible.

Two invariants are worth stating because they are what distinguishes recovery governance from an abandonment of governance. Raised authority must remain **bounded and time-limited**, with a stated review date, or it does not revert. And the escalation of bad news must be made safer, not harder, because the mechanism's failure mode under stress is silence — and silence is the one condition under which no governance design of any shape can work.

3.A.2 Governance of AI-assisted delivery

Where a project uses AI in its own delivery — estimating, scheduling, code generation, risk analysis, document review — governance acquires obligations that most existing structures do not express, and the honest position is that they must be added deliberately.

Four are minimal. **A register of AI uses**, stating for each what the tool does, whose decision it informs, and who is accountable for the output — the same decidability test as 3.1.1, applied to tools. **A verification standard proportional to consequence**: an AI-drafted internal summary and an AI-produced estimate that will set a baseline are not the same object and must not carry the same review. **A prohibition on unattributable authorship** in anything that will be relied on — a decision record, an assurance opinion, a board paper conclusion — because Domain 1's accountability principle requires a person to answer for it, and a tool cannot be that person. And **a data and confidentiality boundary** stated before use, not after an incident.

The governance body's own duty here is the one most often skipped: to ask, of any AI-informed recommendation it receives, what was the model's input, who verified its output, and what would change the conclusion? A body that cannot get those three answers is not governing the recommendation; it is ratifying it.

3.A.3 The reviewer's governance eye

Invariants to test on any governance design, each cheap and each diagnostic:

Every decision class has **exactly one** accountable role, and that role's authority is sufficient under the delegation schedule. The delegation schedule reads on **value, reversibility and externality**, not value alone. The **latency of every escalation path is stated**, computed as $\Sigma (M/2 + L)$, and each tier can be justified against its cost. Every decision body has a **capacity number** and a demand estimate. Committee agendas show the **share of time on decisions** rather than reports. Every gate has criteria set in advance, the authority to stop, and a **computed value** against its elapsed time and detection rate. Conditional passes have owners, dates and consequences, verified at the following gate. The assurance map shows **no gaps and no duplication**. Every escalation class has a stated decision deadline and an out-of-cycle route. The decision log's entries carry a named decision-maker and a **versioned information reference**. And — the one test that subsumes many of the others — the **re-decision count** is tracked, because a question arriving at the same body a third time is proof that something upstream of it is broken.

— Industry variations — Domain 3

- **Public sector and government.** Statutory decision rights, published assurance regimes and mandatory gates that are genuinely non-negotiable; latency is high by design and the leader's lever is almost entirely the paper lead time and the out-of-cycle route, not the cadence.
- **Regulated industries (pharmaceutical, nuclear, aviation, financial services).** Some gates are external and unmovable, and the governance design must place the internal decision before the external submission with enough margin that a failed internal test does not force a submission.
- **Construction and infrastructure.** Multi-party governance is the norm; decision rights follow the contract structure, so a governance design inconsistent with the contract loses to the contract every time (PFL-AI Domain 11).

- **Technology and product organisations.** Bounded envelopes and continuous funding decisions dominate; the characteristic risk is light governance on genuinely irreversible choices — architecture, data model, third-party dependency — because they do not present as gates.
- **Healthcare.** Clinical governance runs in parallel with project governance and clinical authority is not delegable to a project body; Meridian's design must place clinical sign-off with clinical authority and integrate rather than absorb it.
- **Energy and resources.** Stage-gated capital processes with substantial assurance at each gate, and the standard pathology is gate proliferation — re-approval ceremonies whose elapsed time exceeds 3.3.1's breakeven.

— Case study — Domain 3: the four-week month (health, Meridian)

Situation. Twelve weeks into the clinic rollout, Meridian's steering committee was widely described as "the bottleneck". The committee met monthly, was well attended, and had never refused a request. The programme was six weeks behind, and the delay was attributed in the quarterly report to "slower than expected clinic readiness".

What the arithmetic showed. The project leader computed three numbers before the next governance review. Governance latency: $M/2 + L = 4/2 + 2 = 4$ **weeks** — the committee that everyone called monthly imposed a full month's expected wait, and no one in the organisation had believed it was more than a fortnight. Escalation volume: with a **10,000** delegation threshold, **36** of the year's **60** change requests required the committee, of which about a quarter sat on the critical path, costing $9 \times 4 \times 14,280 = \text{USD } 514,080$ a year in delay. Committee capacity: **104** item-slots against **77** items of demand, of which **26** were standing reports producing no decision — so a quarter of the scarcest resource in the programme was being spent on status.

What changed. Three changes, none of which removed any scrutiny anyone could name. The paper lead time went from 2 weeks to 1, taking expected latency to **3.0 weeks** — a full week saved from a change to an administrative deadline. The delegation threshold went from 10,000 to **25,000**, worth **USD 171,360** a year against a worst case of 81,600 even if every delegated decision were decided wrongly and destroyed 40 % of its value. And the standing reports moved to pre-reading, releasing 26 slots and taking utilisation from 74.0 % to **49.0 %**, which is what created room for the contentious item to be reached.

The outcome, and the part that mattered. Escalated decisions began arriving in about three weeks rather than four, and the peak-month deferrals stopped. But the durable change was to the conversation: "the committee is a bottleneck" is an accusation, while "our governance design costs 514,080 a year in delay, and here are three changes worth most of it" is a proposal. The second was approved in one meeting. The first had been raised, in various forms, for two quarters.

What the domain teaches here. Governance is a delivery variable with a computable price. Until the price is computed, governance complaints are cultural and go nowhere; once computed, they are engineering, and engineering gets approved. And note which lever was largest per unit of organisational pain: the paper deadline — free, unglamorous, and twice as effective per week as meeting more often.

— Case study B — Domain 3: the decision nobody made (financial services)

Situation. A payments-platform programme completed build 5 % over budget with no approved change above 50,000 — a clean change-control record. A routine internal audit then found that the delivered scope differed materially from the approved baseline in four respects, none of which appeared in any decision record.

What had happened. Each of the four changes had been discussed in a working group, agreed in substance, and implemented. None had been raised as a change, because each had been individually below the project leader's **50,000** authority and had therefore, in the team's understanding, required no decision at all. Cumulatively they were worth approximately 700,000 and one of them altered a control the second line had assured on the basis of the original design. The programme's governance was not weak in the ordinary sense: the committee functioned, the gates were held, the delegation schedule existed. The defect was narrower and more instructive — **the schedule read on individual value only**, and the decision log recorded decisions taken by bodies rather than decisions taken.

How it resolved. Three corrections, each mapping to a specific defect. The delegation schedule acquired **reversibility and externality** dimensions, so a change touching an assured control required second-line agreement regardless of value. It acquired a **cumulative test**: related changes aggregating above a threshold within a period required the authority appropriate to the aggregate, which is the provision that would have caught all four. And the decision log was made the **register of record for every change to the baseline**, whoever decided it and at whatever value, so that a decision below a threshold still generated an entry. The audit finding was closed; the platform went live four weeks late.

What the domain teaches here. A decision below a threshold is still a decision and still needs a record. A delegation schedule that reads only on value cannot see irreversibility, externality or accumulation — the three ways a small decision becomes a large one — and a decision log that records only what committees decided will systematically miss the changes that matter most, because those are exactly the ones that never reached a committee.

— Executive perspective — Domain 3

What a programme director cannot delegate in this domain:

- **The latency of your own governance.** Compute $\Sigma (M/2 + L)$ for every escalation path you own, price it at your cost of delay, and be able to state it. An unquantified governance design is one you are not managing (3.2.3, 3.3.3).
 - **The delegation schedule, on three dimensions.** Value, reversibility, externality — and a cumulative test. Read on value alone, it will miss exactly the decisions that matter (Case study B).
 - **The sponsor's obligations, in writing.** Not the title. The seven testable obligations, agreed and evidenced, with turnaround on sponsor decisions monitored like any other lead time (3.2.1).
 - **Whether your gates function.** A gate that has never held anything is not working, and its elapsed time may already exceed the point at which it destroys value (3.3.1).
 - **The decision record's integrity.** Named decision-maker, versioned information reference, conditions with owners. This is what makes a defensible decision defensible a year later (3.3.4).
-

- **Escalation lead time as a leading indicator.** Track how far before impact escalations arrive. A shortening trend is the earliest signal that people no longer believe escalation is safe — and it appears in no standard report (3.3.3).

— Calculation exercises — Domain 3

Exercise 3.1 A committee meets every 8 weeks and closes papers 2 weeks before each meeting; cost of delay is 14,280 per week. Compute the expected latency and its cost, then compute the saving from (a) halving the meeting interval to 4 weeks and (b) halving the paper lead time to 1 week. Solution. $E[\text{wait}] = 8/2 + 2 = \mathbf{6.0 \text{ weeks}}$, costing $6 \times 14,280 = \mathbf{USD 85,680}$. (a) $4/2 + 2 = \mathbf{4.0 \text{ weeks}}$ — saves 2.0 weeks, **USD 28,560**. (b) $8/2 + 1 = \mathbf{5.0 \text{ weeks}}$ — saves 1.0 week, **USD 14,280**. Note the general rule the two results illustrate: a cut of x in L always saves x ; a cut of x in M saves $x/2$. Halving looks like the bigger intervention here only because M is four times L . Common error: assuming the expected wait is half the meeting interval, which omits the paper lead time entirely.

Exercise 3.2 An organisation logs 80 change requests a year: $\leq 5,000 \rightarrow 30$; $5,001-20,000 \rightarrow 20$ (average value 11,000); $20,001-75,000 \rightarrow 22$; $> 75,000 \rightarrow 8$. The committee meets every 4 weeks with a 1-week paper lead time; cost of delay is 9,500 per week; 30 % of escalated changes sit on the critical path. Compute the annual governance delay cost at a 5,000 threshold and at a 20,000 threshold, and the breakeven value-destruction rate on the delegated band. Solution. $E[\text{wait}] = 4/2 + 1 = \mathbf{3.0 \text{ weeks}}$, so each delaying change costs $3 \times 9,500 = \mathbf{28,500}$. At 5,000: escalated $80 - 30 = 50$, delaying $50 \times 0.30 = 15$, cost $15 \times 28,500 = \mathbf{USD 427,500}$. At 20,000: escalated $50 - 20 = 30$, delaying 9 , cost **USD 256,500**. Saving **USD 171,000**. The delegated band is worth $20 \times 11,000 = \mathbf{220,000}$, so the escalation breaks even only if the delegate destroys $171,000/220,000 = \mathbf{77.7 \%}$ of the value of every delegated decision. Common error: comparing the saving with the total value of the delegated band rather than with a plausible loss on it, which understates the case for delegating by an order of magnitude.

Exercise 3.3 A gate costs 30,000 in review effort and 4 weeks of elapsed time at a delay cost of 9,500 per week. A material defect is present with probability 0.25; the gate detects it with probability 0.75; correction at design costs 80,000 and in build 600,000. Compute the gate's net value, its breakeven elapsed time and its breakeven detection probability. Solution. Without the gate: $0.25 \times 600,000 = \mathbf{150,000}$. With: review 30,000 + delay $4 \times 9,500 = 38,000$ + expected remediation $0.25 \times (0.75 \times 80,000 + 0.25 \times 600,000) = 0.25 \times 210,000 = 52,500$, total **120,500**. Net value **USD 29,500**. Breakeven elapsed time $(150,000 - 30,000 - 52,500)/9,500 = \mathbf{7.105 \text{ weeks}}$. Breakeven detection probability: solve $30,000 + 38,000 + 0.25 \times (600,000 - 520,000d) = 150,000 \rightarrow d = \mathbf{52.31 \%}$. Common error: omitting the elapsed-time cost, which makes every gate look worthwhile and is the reason gates proliferate.

Exercise 3.4 A decision must pass a delivery board (every 2 weeks, 1-week papers), a portfolio board (every 6 weeks, 2-week papers) and an investment committee (every 12 weeks, 4-week papers). Cost of delay is 9,500 per week. Compute the total expected latency and cost, and the saving from a single-tier path at the portfolio board with a written-resolution procedure of one week. Solution. $2/2 + 1 = \mathbf{2.0}$; $6/2 + 2 = \mathbf{5.0}$; $12/2 + 4 = \mathbf{10.0}$. Total **17.0 weeks**, costing **USD 161,500** — of which the investment committee alone is 10.0 weeks and **95,000**. A one-week written-resolution path costs **USD 9,500**, a saving of **USD 152,000 (94.1 %)**. For comparison,

retaining the portfolio board on its ordinary cadence as the sole tier gives 5.0 weeks and **47,500**, saving 114,000. Common error: adding only the meeting intervals and ignoring the paper lead times, which understates latency here by 7 of the 17 weeks.

— Practitioner's toolkit — Domain 3

Adoption-ready artefacts; adapt headings to your organisation, then keep them stable.

Toolkit 3.T.1 — Governance design sheet

One page, completed before mobilisation and reviewed at every gate. Rows: each decision body, with its purpose, members (distinguished from attendees), decision classes it owns, meeting interval M , paper lead time L , computed $E[\text{wait}]$, capacity (meetings $\times$ items) and estimated demand. Below it, each escalation path with its tiers, the summed latency $\Sigma (M/2 + L)$, that latency priced at the project's cost of delay, and the out-of-cycle mechanism. The sheet's purpose is to make latency and capacity **visible at design time**; a governance design that cannot fill it in has not been designed.

Toolkit 3.T.2 — Delegation schedule with a cumulative test

Columns: decision class $\cdot$ value threshold $\cdot$ reversibility rating (reversible / costly to reverse / irreversible) $\cdot$ externality rating (internal / cross-functional / external or public) $\cdot$ authority required $\cdot$ **cumulative rule** (related decisions aggregating above X within period P require the authority appropriate to the aggregate) $\cdot$ information the delegate must hold to decide. The three dimensions prevent the value-only failure of Case study B; the cumulative rule prevents its specific mechanism; and the final column is what distinguishes delegation from abdication.

Toolkit 3.T.3 — Decision record entry and its integrity checks

Per-decision fields: reference $\cdot$ date $\cdot$ decision-maker (**named person and role**) $\cdot$ decision (one reading only) $\cdot$ options considered and why this one $\cdot$ information relied on **with versions** $\cdot$ conditions with owners and dates $\cdot$ review date if provisional. Monthly integrity checks, each a count: entries missing a named decision-maker; entries missing a versioned reference; conditions past their date; decision classes with zero or multiple Accountable roles in the RACI (the single-A test); and **re-decisions** — the same question arriving at the same body more than twice. Every one of these is a number, and a governance function that reports them monthly will find its defects before an auditor does.

— Exam preparation — Domain 3

What is assessed. Governance purpose and its separation from management and assurance; structural weaknesses by organisational form; governing iterative and hybrid delivery; the sponsor's obligations and failure modes; steering-committee design faults; **the governance latency formula and its application**; delegation thresholds and their economics; gate purpose, criteria and value; the three assurance lines and their failure modes; escalation design with stated latency; and the decision record with its integrity checks.

The calculations to be able to do under time pressure. $E[\text{wait}] = M/2 + L$ for a single body and $\Sigma (M/2 + L)$ for a multi-tier path, priced at a cost of delay. Escalation volume and delay cost at alternative thresholds, and the breakeven value-destruction rate on a delegated band. Gate net value, breakeven elapsed time and breakeven detection probability. Committee capacity and utilisation. RACI single-A defect rate.

The traps. Taking expected wait as half the meeting interval and forgetting the paper lead time (Exercise 3.1) · adding meeting intervals without paper lead times in a multi-tier path (Exercise 3.4) · treating governance artefacts as governance (3.1.1) · omitting elapsed-time cost from gate value, which makes every gate look worthwhile (Exercise 3.3) · comparing a delegation saving with the total value of the delegated band rather than a plausible loss on it (Exercise 3.2) · assuming a committee's authority is its highest member's rather than its binding minimum (3.2.2) · treating a favourable assurance opinion as a transfer of accountability (3.3.2) · recording only committee decisions and missing the ones taken below a threshold (Case study B) · reading a delegation schedule on value alone.

How the domain connects. Domain 1 supplies the accountability principle governance must respect and the cost of delay every calculation here is priced at. Domain 2 supplies the business case and kill criteria that gates re-test. Domain 4 integrates the governance of parts into a whole. Domain 6 consumes governance latency as schedule input — a decision path of 15.5 weeks is a 15.5-week predecessor, whatever the plan says. Domain 7's change control is the delegation schedule in operation. Domain 8's risk escalation uses the escalation classes designed here. And PFL-AI Domain 11 handles the risk-allocation face of the multi-party governance problem.

— Summary — Domain 3

Governance is the decision rights, accountabilities and information flows through which an organisation directs a project — and it is not management, not assurance, and emphatically not the reporting apparatus that is usually offered in its place. A governance design must produce decidability, timeliness, legitimacy and traceability, and each of those is testable.

The domain's contribution is to make governance **computable**. The expected wait for a committee decision is $M/2 + L$: half the meeting interval plus the whole paper lead time. That single formula prices a governance design at design time, shows that the administrative deadline is twice the lever meeting frequency is, and adds across tiers to reveal what an escalation path actually costs — 15.5 weeks and **USD 221,340** for Meridian's three-tier path, of which the quarterly committee alone is 61 %. The same arithmetic prices a delegation threshold: Meridian's 10,000 threshold cost **USD 514,080** a year, and raising it to 25,000 saves **USD 171,360** against a worst case of 81,600 even if every delegated decision were decided wrongly — a breakeven value destruction of **84 %** per decision, which no plausible delegate approaches. And it prices a gate: Meridian's design gate is worth **USD 56,520**, stops paying beyond **9.96 weeks** of elapsed time, and requires a detection probability above **55.85 %** to be worth holding at all.

The sponsor's role is a set of testable obligations, not a title; steering committees fail in five predictable ways, four of them designed in; gates buy optionality against irreversibility and belong where irreversibility steps up; assurance has three lines whose worst failure is capture; escalation is a timed pathway with an out-of-cycle route, not an instruction to tell someone senior; and a decision that is not recorded, with a named decision-maker and a versioned information reference, has not been made — it has been remembered. Case study B's four changes were each individually within

authority and cumulatively worth 700,000, which is the whole argument for a delegation schedule that reads on reversibility and externality and aggregates.

The through-line: **governance has a price, the price is computable, and until it is computed governance complaints are cultural and go nowhere.** Compute it, and they become engineering.

02

PART TWO

Delivering the work

Domains 5–10 — scope, planning and flow, cost and commercial, risk and resilience, quality and assurance, and procurement: the machinery of delivery.



06

DOMAIN 6

Planning, Scheduling and Delivery Flow (quantitative flagship)

Group: Delivering the work (Domain 6 of 6 in Part Two). **Target:** ~76 pages. **Binds to:** the PCI Book Pattern Specification and the shared registries (<docs/books/registries/>). This domain is the home of the schedule symbols — **ES**, **EF**, **LS**, **LF**, **TF** (total float), **FF** (free float) — restated by every domain that touches time. British English; USD (+SAR where useful, indicative **USD 1 ≈ SAR 3.75**).

— Why this domain exists

A project leader's authority rests on one repeated act: making a credible statement about the future — when the work will finish, what could move that date, and what it would cost to move it back. This domain builds the machinery behind that credibility. It starts where planning actually starts — with levels of plan and the logic that binds work together (KA 6.1); builds the critical path method in full, forward and backward passes, and both kinds of float (KA 6.2); adds the realities the textbook network ignores — resources, constraints, milestones and rolling-wave elaboration (KA 6.3); and finishes with delivery flow across predictive, agile and hybrid worlds: recovery, compression economics, and forecasting under uncertainty (KA 6.4). Cost joins schedule in Domain 7 (earned value); risk quantification deepens in Domain 8. A leader who cannot read a network diagram is hostage to whoever can; this domain removes the hostage-taking.

Learning objectives. After this domain a candidate can: choose the right planning level for an audience and a decision; build a logic network with correct dependency types; run forward and backward passes and derive total and free float; identify the critical path and near-critical paths and explain why float is not spare time; plan resources against the schedule and level them deliberately; apply rolling-wave elaboration honestly; select and price schedule compression; distinguish predictive, agile and hybrid scheduling and govern each; and use three-point estimates and scenario analysis to forecast completion with stated uncertainty — with any AI-generated schedule subject to the family's verification rule.

The master worked project. One project runs through this domain and returns in Domains 7 and 8: **Project Auriga**, a control-systems upgrade for a regional utility, planned in weeks:

ID	Activity	Duration (wk)	Predecessors
A	Mobilise	2	—
B	Detailed design	6	A
C	Procure control hardware	8	B
D	Civil and cabling works	7	B
E	Installation	5	C, D
F	Testing and commissioning	4	E
G	Operator training and handover	2	D

Every calculation below uses this network.

KNOWLEDGE AREA 6.1

Planning levels and logic networks

Topics: 6.1.1 levels of plan · 6.1.2 dependencies and logic · 6.1.3 network rules and quality.

6.1.1 Levels of plan

The principle. One schedule cannot serve a board, a site supervisor and a lender at once. Mature delivery organisations maintain a **hierarchy of schedules**, each derived from the one below, each honest with the one above:

Level	Name	Audience & decision	Typical granularity
L1	Executive summary	Board, sponsor — go/no-go, milestones	10–30 bars
L2	Management schedule	Steering, PMO — stage gates, interfaces	50–300 activities
L3	Control schedule	Project team — the CPM network; progress and float live here	300–5,000 activities
L4	Execution/lookahead	Site and squads — 2–6 week detail	daily/shift detail

The **control schedule (L3)** is where this domain's machinery operates: it carries the logic, the floats and the critical path, and it is the version under change control (Domain 4, KA 4.3). An L1 milestone that cannot be traced to L3 logic is decoration, and steering committees should treat it as such (Domain 3, KA 3.2).

Fig 6.1.1 — The schedule hierarchy: one logic, four honest views

6.1.2 Dependencies and logic

Definitions. Activities link through four dependency types — finish-to-start (**FS**, the default: design finishes, then procurement starts), start-to-start (**SS**: cabling can start once civils have started, often with a lag), finish-to-finish (**FF**: commissioning documentation finishes when commissioning finishes),

and the rare start-to-finish (**SF**). A **lag** delays a successor (**FS+2** = start two weeks after the predecessor finishes); a **lead** (negative lag) overlaps it. Logic expresses physics and choice — what must follow, and what the team has chosen to sequence — and the two should be distinguishable in the schedule's notes, because only chosen logic can be re-chosen when recovery demands it (KA 6.4).

Common pitfall — the lag that hides work. A 3-week lag labelled "curing" is physics; a 3-week lag labelled nothing is often an activity somebody didn't model — unresourced, uncosted, invisible to progress measurement. Audit rule: every lag has a stated reason, and long lags (> 10 % of project duration) are activities in disguise.

WORKED EXAMPLE

Worked example 6.1.2 — re-choosing a dependency.

1. **Setup.** Survey work P (4 weeks) precedes report drafting Q (6 weeks), currently linked **FS+2** (two weeks of data cleaning after all surveying ends). The team proposes drafting in parallel as survey sections complete: **SS+1**. Compute Q's dates both ways (P starts week 0).
2. **Formula.** **FS+lag:** $ES(Q) = EF(P) + \text{lag}$. **SS+lag:** $ES(Q) = ES(P) + \text{lag}$.
3. **Substitution.** **FS+2:** $ES(Q) = 4 + 2 = 6$, $EF(Q) = 12$. **SS+1:** $ES(Q) = 0 + 1 = 1$, $EF(Q) = 7$.
4. **Result.** The chain finishes week **12** under **FS+2** and week **7** under **SS+1** — five weeks earlier from one logic choice.
5. **Interpretation.** Nothing was crashed and nobody works faster: the saving comes from overlapping, and its price is rework risk if early sections change after later surveying. This is why chosen logic must be visibly chosen (not fossilised as physics) — it is the cheapest recovery currency the project owns (KA 6.4.2), and spending it is a risk decision, not a scheduling trick.

6.1.3 Network rules and quality

A network the passes can trust obeys checkable rules: every activity except the first has a predecessor and except the last a successor (**no dangles**); no loops; durations estimated by the owning discipline, not the scheduler; no date constraints doing logic's job (a "must start on" pin that overrides logic converts the schedule from a model into a poster); and logic density in a healthy range — too few links and float is fiction, too many and nothing can move. These rules are the schedule-quality gate a leader can run without scheduling software: ask for the dangle count, the constraint count and the reasons for both.

AI in this KA

Scheduling tools now draft networks from scope statements and historical libraries. Useful — the draft logic is often 80 % right — and dangerous in a specific way: generated logic looks authoritative while encoding assumptions nobody made. The governed workflow: AI proposes the network; the discipline leads walk every link of their scope ("does B truly need all of A finished?"); the scheduler runs the quality rules above; and the leader signs the logic as a decision record (Domain 3, KA 3.3.4).

AI proposes; the professional verifies, decides and remains accountable.

■ KEY TERMS — KA 6.1

Term	Meaning
Schedule hierarchy (L1-L4)	Nested plans for different audiences, derived from one logic.
Control schedule	The L3 CPM network under change control; home of float and progress.
FS / SS / FF / SF	The four dependency types.
Lag / lead	Imposed delay / overlap on a dependency.
Dangle	An activity missing a predecessor or successor; a network defect.
Date constraint	A pinned date overriding logic; each one weakens the model.

■ SAMPLE MCQS — KA 6.1

MCQ 6.1-A [6.1.2 · Application] Cabling may begin one week after civil works begin. The correct dependency is: - A. FS+1 - B. SS+1 - C. FF+1 - D. FS-1

Rationale: The condition binds the two starts with a one-week lag: start-to-start plus one. A would wait for all civils to finish; C binds the finishes; D (a lead on FS) overlaps the finish, which is not what was stated.

MCQ 6.1-B [6.1.3 · Analysis] A schedule shows 14 "must finish on" constraints, all on milestones reported to the board. The most accurate reading is: - A. the schedule is well-governed, because board dates are protected - B. the milestones will be met, because the tool will honour the constraints - C. the schedule may no longer model reality — pinned dates can mask logic-driven slippage that float analysis would otherwise reveal - D. constraints are neutral scheduling hygiene with no analytical effect

Rationale: Pins override logic: the network can be slipping while pinned milestones hold still, hiding negative float until it is unrecoverable. A and B mistake suppression for control; D is false — every pin degrades the model's predictive value.

MCQ 6.1-E [6.1.1 · Recall] The schedule level that carries the logic, the floats and the critical path — and sits under formal change control — is: - A. L1, because the board owns the schedule - B. L2, because the PMO maintains it - C. L3, the control schedule - D. L4, because it has the most detail

Rationale: The L3 control schedule is the analytical model; L1/L2 summarise it and L4 elaborates near-term execution from it. Detail (D) is not the same as control — L4 churns weekly by design and is never the baselined network.

MCQ 6.1-C [6.1.2 · Application] Survey P (4 wk) feeds report Q (6 wk). Under FS+2 the chain completes week 12. Re-linked SS+1, it completes week: - A. 7 - B. 11 - C. 12 - D. 5

Rationale: $ES(Q) = ES(P) + 1 = 1$, $EF(Q) = 7$. B subtracts one lag week from 12; C assumes logic changes nothing; D forgets Q's own duration must still run.

MCQ 6.1-D [6.1.3 · Recall] A "dangle" in a schedule network is: - A. an activity with negative float - B. an activity missing a predecessor or successor link - C. a milestone with zero duration - D. a lag longer than its predecessor

Rationale: Dangles are unbound activity ends — the passes cannot constrain them, so their dates and float are unreliable. A describes constraint-driven lateness; C describes every milestone; D is unusual but legal when justified.

■ SELF-CHECK — KA 6.1

1. Why must chosen logic be distinguishable from physical logic? — Because recovery (KA 6.4) works by re-choosing choices; physics cannot be renegotiated.
2. What three numbers give a fast schedule-quality read? — Dangle count, constraint count, unexplained-lag count — each should be at or near zero, with reasons for the rest.

KNOWLEDGE AREA 6.2

The critical path and float

Topics: 6.2.1 the forward and backward passes · 6.2.2 total and free float · 6.2.3 reading the critical path.

6.2.1 The forward and backward passes

Definitions. For each activity: **ES/EF** — earliest start/finish (forward pass, left to right: **ES** = latest **EF** of predecessors; **EF** = **ES** + duration); **LS/LF** — latest start/finish that do not delay the project (backward pass, right to left: **LF** = earliest **LS** of successors; **LS** = **LF** - duration). The project duration is the largest **EF** in the network.

Worked example 6.2.1 — Auriga, both passes. The passes are inherently tabular; the labelled table replaces the five-step skeleton, and the interpretation follows (pattern rule).

Act	Dur	ES	EF	LS	LF	TF	FF
A	2	0	2	0	2	0	0
B	6	2	8	2	8	0	0
C	8	8	16	8	16	0	0
D	7	8	15	9	16	1	0
E	5	16	21	16	21	0	0
F	4	21	25	21	25	0	0
G	2	15	17	23	25	8	8

Interpretation. Project duration **25 weeks**; the **critical path is A-B-C-E-F** ($2+6+8+5+4 = 25$), every activity on it with zero float. D can slip one week without moving the end date; G can slip eight. The passes turn a drawing into an instrument: every schedule statement a leader makes in this domain — "we finish in week 25", "D matters more than G", "procurement drives the job" — is read directly off this table.




Fig 6.2.1 — The Auriga network with both passes

6.2.2 Total and free float

Definitions.

TF = LS - ES — how far an activity can slip without delaying the project
 FF = min(ES of successors) - EF — how far it can slip without delaying ANY successor

$FF \leq TF$ always. The Auriga table carries the domain's subtlest lesson: **D has TF = 1 but FF = 0** — D can slip a week without moving the end date, but the moment it slips at all, G's earliest start moves. Total float is a project-level commodity; free float is a courtesy to your successors. A leader who grants a subcontractor "the float" without saying which kind has just given away someone else's schedule.

Float is not spare time. Float belongs to the path, not the activity: consume D's week and every activity sharing that path inherits zero. Contract regimes differ on who owns float (Domain 10, KA 10.3, treats float ownership clauses); the delivery answer is simpler — float is a risk buffer under the leader's governance, spent deliberately or not at all.

Common pitfall — managing only the critical path. D sits one week from critical. A team watching only A-B-C-E-F will discover D's importance the week it becomes too late to matter. Near-critical analysis (all paths within, say, 10 % of project duration) is standing practice; KA 6.4's scenario work builds on it.

6.2.3 Reading the critical path

The critical path is the longest path and therefore the shortest possible duration — both statements are the same fact, and fluency means holding both. Three leader-grade readings of the Auriga table: **(1)** procurement (C) is the single most schedule-driving activity — an expediting conversation is worth more than any amount of site pressure; **(2)** the C/D pair converge on E, so E's start is hostage

to the later of two chains — convergence points are where slippage compounds and where progress review should focus (Domain 8, KA 8.1); **(3)** G's eight weeks of float make it the natural donor of resources in any recovery (KA 6.4).

AI in this KA

CPM arithmetic is deterministic — ideal for machine computation and instant human verification. The two checks that catch nearly every pass error: the critical path's durations must **sum exactly** to the project duration, and every critical activity must show $TF = 0$ with $ES = LS$. An AI-produced schedule that fails either is broken; one that passes both may still have wrong logic — which is why KA 6.1's link-walking discipline precedes trust in any pass, human or machine.

■ KEY TERMS — KA 6.2

Term	Meaning
Forward / backward pass	Left-to-right ES/EF computation; right-to-left LS/LF computation.
ES EF LS LF	Earliest/latest start and finish per activity.
Total float TF	$LS - ES$; slip available without delaying the project.
Free float FF	Slip available without delaying any successor; $FF \leq TF$.
Critical path	The zero-float longest path; the project's shortest possible duration.
Convergence point	A node where paths merge; slippage compounds there.

■ SAMPLE MCQS — KA 6.2

MCQ 6.2-A [6.2.1 · Application] In the Auriga network, activity D (duration 7, predecessor B finishing week 8, successors E and G) has $LS = 9$ and $ES = 8$. Its total float is: - A. 0 - B. 1 - C. 8 - D. 7

Rationale: $TF = LS - ES = 9 - 8 = 1$. A confuses D with the critical activities; C is G's float; D is the activity's duration, not its float.

MCQ 6.2-B [6.2.2 · Analysis] An activity shows $TF = 1$ and $FF = 0$. Delaying it by one week will: - A. delay the project by one week - B. delay nothing, because total float absorbs the slip - C. delay at least one successor's earliest start while leaving the project end date unchanged - D. be impossible — free float can never be below total float

Rationale: Positive TF protects the end date (not A); zero FF means some successor moves immediately (not B). D inverts the invariant — $FF \leq TF$ always.

MCQ 6.2-C [6.2.3 · Application] Auriga's procurement activity C is crashed from 8 weeks to 6. The new project duration is: - A. 23 weeks - B. 24 weeks - C. 25 weeks - D. 21 weeks

Rationale: After one week of crashing, path B-D-E ($8+7+5+4$ via D) becomes co-critical at 24; the second crashed week buys nothing — duration stays 24. A assumes both weeks convert to project weeks; C forgets the crash entirely; D subtracts from the wrong baseline. (The full economics: KA 6.4.)

MCQ 6.2-F [6.2.1 · Application] Auriga's procurement C slips from 8 weeks to 9 (all else unchanged). The new project duration is: - A. 25 weeks - B. 26 weeks - C. 27 weeks - D. 24 weeks

Rationale: C was critical with zero float, so its extra week passes straight through: E runs 17–22, F finishes week 26. A assumes float absorbed a critical activity's slip; C adds the week twice; D subtracts it.

MCQ 6.2-D [6.2.1 · Application] Auriga's training activity G runs ES 15–EF 17 with $LF = 25$. Its total float is: - A. 8 - B. 10 - C. 2 - D. 0

Rationale: $TF = LF - EF = 25 - 17 = 8$ (equivalently $LS - ES = 23 - 15$). B ignores G's duration (25 – 15); C is G's duration; D confuses G with the critical path.

MCQ 6.2-E [6.2.3 · Analysis] A programme's binding path shows $TF = -3$ after a constraint-honest pass. The professional reading is: - A. the software has malfunctioned — float cannot be negative - B. the plan, as constrained, is already three weeks late; the number sizes the recovery problem and must be reported now - C. delete the constraints so the float returns to zero - D. the project will finish three weeks early

Rationale: Negative float is the gap between what logic needs and what constraints allow — information, not error (A) and not good news (D). C is the pin-and-squeeze suppression this domain's Case B dissects; it hides the problem until it is unrecoverable.

■ SELF-CHECK — KA 6.2

1. State both definitions of the critical path and why they are equivalent. — Longest path through the network; shortest possible project duration — the longest chain of unavoidable work is precisely what bounds how soon the whole can finish.
2. Why is $FF \leq TF$ always? — Delaying any successor is one of the ways of delaying the project; the constraint set defining FF is stricter.
3. Auriga slips: D takes 10 weeks, not 7. New duration? — 27 weeks: D finishes week 18, E waits for it (18 > 16), E–F add 9 → the critical path re-routes through D (KA 6.4 works the recovery).

KNOWLEDGE AREA 6.3

Resources, constraints, milestones and rolling wave

Topics: 6.3.1 resource planning and levelling · 6.3.2 constraints and milestones · 6.3.3 rolling-wave planning.

6.3.1 Resource planning and levelling

The principle. The passes of KA 6.2 assume infinite resources; reality staffs the network. Loading resources onto the schedule produces a **histogram** — demand per period — and the histogram almost always spikes. Two responses: **resource smoothing** uses float to flatten demand without moving the end date (slip D and G inside their float to pull crews off the peak); **resource levelling** accepts a later end date if the resource cap is hard. The order of operations is a leadership decision in disguise: smoothing spends float — the project's risk buffer — to save money, and someone accountable should price that trade explicitly, not discover it in a progress meeting.

WORKED EXAMPLE**Worked example 6.3.1 — smoothing Auriga's field crews.**

1. **Setup.** Weeks 9-15: civil works D needs 3 crews; early installation staging inside C needs 2; the site cap is 4 crews. Overlap weeks 9-15 demand 5 — one over cap.
2. **Formula.** Smoothing test: can a demanding activity slip within **TF** (and, if successors matter, **FF**) to clear the peak?
3. **Substitution.** D has **TF** = 1, **FF** = 0: slipping D one week clears one peak week but consumes D's entire float and immediately moves G (**FF** = 0). The staging work inside C has discretionary timing across weeks 9-16.
4. **Result.** Re-sequence the staging within C to weeks 14-16 where D's demand tails off: peak demand falls to 4, the cap holds, the end date holds, and **no float is spent**.
5. **Interpretation.** The cheapest capacity is re-sequencing; the second cheapest is float; the expensive ones are overtime, second shifts and extension. A leader asks for the histogram with the options priced, not just the spike.

WORKED EXAMPLE**Worked example 6.3.1b — when the cap is hard: levelling, priced.**

1. **Setup.** Now suppose the site cap is **3 crews** (a permit condition, not a preference). D needs all 3 for weeks 9-15; the 4-week staging package (2 crews) must finish before E can start (week 16). Delay costs USD 45,000 per week (the client's bonus forgone); a second-shift waiver for the staging crew costs USD 20,000 per week. What does the leader do?
2. **Formula.** Compare total cost of each feasible plan: pure levelling (extend) vs paid capacity (second shift) vs any float-funded re-sequence.
3. **Substitution.** Levelling: staging cannot start until D releases crews (week 16); it runs weeks 16-19, E starts week 20 — project 25 → 29, cost $4 \times 45,000 = 180,000$. Second shift: staging runs weeks 12-15 off-shift, E starts on time — cost $4 \times 20,000 = 80,000$. No float-funded option exists: D's **TF** = 1 cannot host a 4-week move.
4. **Result.** **Buy the second shift: USD 80,000** against USD 180,000 of extension — and the end date holds.
5. **Interpretation.** A hard cap converts a scheduling problem into a procurement problem. The histogram's job is to surface that conversion early enough for the cheap option to exist — second shifts, pre-assembly, off-site staging all have lead times of their own. Levelling that silently extends the project is a decision someone should have been asked to make.




Fig 6.3.1 — Auriga field-crew histogram, before and after smoothing

6.3.2 Constraints and milestones

Constraints are externally imposed dates — an outage window, a regulatory deadline, a seasonal limit. Modelled honestly they are logic (an outage window is a predecessor with a date); modelled lazily they are pins that falsify float (KA 6.1.3). **Milestones** are zero-duration events that carry meaning: contractual obligations (Domain 10), gate decisions (Domain 3), integration points with other projects (Domain 15). A milestone's date is an output of the network, never an input — the moment a milestone is typed rather than computed, the schedule has stopped forecasting and started decorating.

Milestone hygiene: each has an owner, an unambiguous done-definition (Domain 5's acceptance criteria), traceable L3 logic, and — for contractual milestones — a stated float position reported honestly to the counterparty (Domain 11's reporting ethics).

6.3.3 Rolling-wave planning

The principle. Detail decays with distance: estimating month 18's tasks at daily granularity today manufactures precision, not knowledge. **Rolling wave** plans near-term work at execution detail (L4) and far work at planning packages (L3 summary), elaborating each wave as it approaches under change control.




Fig 6.3.2 — The rolling-wave horizon

AI in this KA

Resource optimisation is a genuine AI strength — levelling across thousands of activities and dozens of calendars is combinatorial work machines do better than planners. Governance holds the same shape: the optimiser proposes a levelled plan; the planner verifies the constraint set was real (crew interchangeability, shift rules, site logistics); the leader owns the trade the optimiser cannot see — whether spending float, money or scope is the right currency this month. An optimiser given a wrong constraint produces a beautifully efficient fiction.

■ KEY TERMS — KA 6.3

Term	Meaning
Resource histogram	Demand per period from the loaded schedule.
Smoothing / levelling	Flattening demand within float / accepting a later date under a hard cap.
Constraint	An externally imposed date; modelled as logic, not as a pin.
Milestone	Zero-duration event with an owner and a done-definition; its date is computed.
Rolling wave	Near-term detail, far-term packages, elaborated under change control.
Planning package	A far-wave summary activity with ranged duration awaiting elaboration.

■ SAMPLE MCQS — KA 6.3

MCQ 6.3-A [6.3.1 · Application] A planner clears a resource peak by delaying an activity with $TF = 1$, $FF = 0$ by one week. The immediate consequences are: - A. no schedule effect of any kind - B. project delay of one week - C. the end date holds, the activity's path float is exhausted, and at least one successor's earliest start moves - D. the resource peak worsens

Rationale: One week inside TF protects the end date (not B), but zero FF moves a successor and the path's buffer is now spent (not A) — the smoothing was legal but not free.

MCQ 6.3-B [6.3.3 · Analysis] A 30-month programme shows daily-level tasks throughout, including month 30. The strongest inference is: - A. the planning team is unusually diligent - B. far-horizon detail is manufactured precision that will churn on every update, obscuring real variance - C. the schedule will need no re-baselining - D. rolling-wave planning has been correctly applied

Rationale: Detail beyond the knowable horizon creates update churn and false confidence — the opposite of diligence in effect (A) and the opposite of rolling wave (D); constant churn makes re-baselining more likely, not less (C).

MCQ 6.3-E [6.3.1 · Application] D needs 3 crews in weeks 9–15 and concurrent staging needs 2; the site cap is 4. The excess demand the histogram must clear is: - A. 1 crew for 7 weeks - B. 5 crews for 7 weeks - C. 2 crews for 7 weeks - D. nothing — $3 + 2 = 5$ is within a 4-crew cap across two activities

Rationale: Peak demand 5 against cap 4 leaves one excess crew-week in each of the seven overlap weeks — the precise quantity smoothing must relocate. B is total demand, not excess; C is the staging demand itself; D misreads the cap as per-activity when it is per-site.

MCQ 6.3-C [6.3.1 · Application] A hard 3-crew cap forces either a 4-week project extension (delay cost USD 45,000/week) or a second-shift waiver at USD 20,000/week for the same 4 weeks (end date held). The value-maximising choice and its saving are: - A. extend: saves USD 100,000 - B. second shift: saves USD 100,000 - C. second shift: saves USD 25,000 - D. they cost the same

Rationale: Extension costs $4 \times 45,000 = 180,000$; the shift premium $4 \times 20,000 = 80,000$ — choosing the shift saves USD 100,000 and keeps the date. A picks the dearer plan; C compares one week only; D ignores the 65,000-per-week difference... which compounds four times.

MCQ 6.3-D [6.3.3 · Recall] A legitimate planning package in a rolling-wave schedule must carry: - A. daily-level task detail for its whole span - B. a ranged duration consistent with its estimate class and a dated elaboration event under change control - C. zero float, to keep pressure on the team - D. a pinned finish date agreed with the sponsor

Rationale: Far-wave honesty is ranged duration plus a governed elaboration point. A is the manufactured precision rolling wave exists to avoid; C fabricates criticality; D is the typed milestone anti-pattern (Case study B).

■ SELF-CHECK — KA 6.3

1. Why is smoothing "spending the risk buffer"? — It consumes float, and float is the project's absorption capacity for the unknown (KA 6.2.2).
2. What makes a milestone honest? — Computed date, named owner, unambiguous done-definition, traceable logic.
3. What must a planning package carry to be legitimate? — A ranged duration consistent with its estimate class and a dated elaboration event under change control.

KNOWLEDGE AREA 6.4

Delivery flow: predictive, agile and hybrid; recovery and forecasting

Topics: 6.4.1 flow across delivery modes · 6.4.2 schedule compression and recovery · 6.4.3 forecasting and scenario analysis.

6.4.1 Flow across delivery modes

Predictive scheduling (this domain so far) plans the whole network and measures variance against baseline. **Agile** delivery (Domain 13 in full) fixes cadence and capacity and lets scope flow: the schedule questions become throughput (items per iteration), cycle time and forecast by velocity. **Hybrid** programmes — the commonest reality — run engineered streams predictively and product streams by cadence, joined at **integration milestones** that appear in the CPM network as fixed-capacity deliveries. The leader's job is translation: a velocity forecast ("the reporting module completes in 4–6 sprints") enters the network as a ranged duration on the integration activity, and the network returns what the programme needs from the agile stream — the latest acceptable delivery, read from **LS**. Neither mode is senior; the network prices time, the cadence protects focus, and governance (Domain 3, KA 3.1.3) holds both to evidence.

6.4.2 Schedule compression and recovery

The two levers. **Crashing** buys duration with money (more resources, overtime, expediting); **fast-tracking** buys it with risk (overlapping activities that logic preferred in sequence). Both obey the same law: work the critical path, and re-run the passes after every move, because **the critical path migrates**.

WORKED EXAMPLE

Worked example 6.4.2 — the crash that stopped paying.

1. **Setup.** Auriga's client offers a bonus of USD 45,000 ($\approx$ SAR 168,750) per week saved. Expediting C costs USD 30,000 per week, up to two weeks. Should the leader buy one week, two, or none?
2. **Formula.** Crash only while (weeks actually saved $\times$ value per week) $>$ crash cost — checked against the re-run passes, not the original network.
3. **Substitution.** Crash C by 1 (8 $\rightarrow$ 7): duration 25 $\rightarrow$ 24; net gain $45,000 - 30,000 = +15,000$. Crash C by 2 (8 $\rightarrow$ 6): path B-D-E-F (2+6+7+5+4) is now co-critical at 24 — duration stays 24; second week's net $0 - 30,000 = -30,000$.
4. **Result. Crash one week only:** +USD 15,000. The second week is pure loss.
5. **Interpretation.** Compression is subject to sharply diminishing returns because every crash promotes the next-longest path. The disciplined sequence: identify the binding path $\rightarrow$ price the cheapest week on it $\rightarrow$ re-run the passes $\rightarrow$ repeat until the next week costs more than it is worth. Leaders who mandate "crash procurement by two weeks" from the original network pay for weeks that no longer exist.




Fig 6.4.1 — The economics of crashing Auriga's procurement

Recovery. When the network slips (self-check 6.2.3: D at 10 weeks → 27), recovery options rank by cost-of-time: re-sequencing and logic re-choice (cheapest — KA 6.1's chosen logic); float harvesting from non-critical paths (G's 8 weeks fund nothing on the new critical path, but its resources might); crashing the new critical path (now through D); fast-tracking E against D's tail with an explicit rework risk (Domain 8's risk register prices it); and scope or acceptance re-negotiation (Domain 5) as the honest last resort. A **recovery plan** states the target date, the moves, their costs and risks, and the decision authority spending them — the template is toolkit 6.T.2.

6.4.3 Forecasting and scenario analysis

Three-point estimates. Uncertain durations are ranges, not points. The PERT expected duration and spread for an activity estimated optimistic o , most-likely m , pessimistic p :

$$t_e = (o + 4m + p) / 6 \quad \sigma = (p - o) / 6$$

WORKED EXAMPLE**Worked example 6.4.3 — Auriga's installation risk.**

1. **Setup.** Installation E is estimated $o = 4$, $m = 5$, $p = 12$ weeks (the long tail: a legacy-system integration problem discovered late).
2. **Formula.** $t_e = (o + 4m + p)/6$; $\sigma = (p - o)/6$.
3. **Substitution.** $t_e = (4 + 20 + 12)/6 = 36/6$; $\sigma = (12 - 4)/6 = 8/6$.
4. **Result.** $t_e = 6.0$ weeks (not 5); $\sigma = 1.33$ weeks.
5. **Interpretation.** The deterministic network used 5 weeks; the risk-weighted expectation is 6 — the skewed tail alone adds a week to the honest forecast, pushing expected completion toward 26 weeks before any risk event "happens". Deterministic dates are the mode, not the mean; boards deserve the mean and the spread. Full quantitative schedule risk analysis — correlating activities, simulating paths, criticality indices — is Domain 8 (KA 8.2), built on exactly this input.

Scenario analysis. Between the single date and the full simulation sits the leader's workhorse: three coherent scenarios (base / threat / opportunity), each a complete re-run of the passes under stated assumptions, each with owners for the assumptions that differ. Scenarios are cheap, auditable, and — because each is a real network — they expose path migration that a percentage-confidence number hides.

The earned-schedule bridge. Once cost joins schedule (Domain 7), progress data yields a time-based forecast: **earned schedule ES** asks when the value now earned was planned to have been earned, and $SPI(t) = ES / AT$ (actual time) measures schedule efficiency in time units. If Auriga's week-22 status shows work that the baseline expected by week 20, $SPI(t) = 20/22 = 0.91$ — the programme is running at 91 % of planned tempo, a signal the currency-based **SPI** famously loses late in a project. Domain 7 (KA 7.3) builds the full machinery; it is flagged here because schedule forecasting belongs to whoever owns the network, not only to the cost engineer.

AI in this KA

Machine forecasting earns trust in exactly one way: calibration. A model that ingests progress data and produces completion forecasts must show its record — forecast vs actual across past periods — before its numbers enter a board pack (Domain 14, KA 14.3's verification workflow). The leader's questions are the auditor's: what data trained it, what does it do when the data thins, who re-ran its arithmetic, and what would make it wrong? Fluent confidence without a calibration record is the machine version of the site manager who is "sure it'll be fine".

■ KEY TERMS — KA 6.4

Term	Meaning
Crashing / fast-tracking	Buying duration with money / with overlap risk.
Path migration	The critical path moving as durations change; why passes are re-run.
Recovery plan	Target date, ranked moves, costs, risks, decision authority.
Three-point / PERT estimate	$t_e = (o+4m+p)/6$, $\sigma = (p-o)/6$.
Scenario analysis	Complete network re-runs under coherent stated assumptions.
Integration milestone	The CPM node where an agile stream's delivery enters the network.

■ SAMPLE MCQS — KA 6.4

MCQ 6.4-A [6.4.2 · Application] Saving a week is worth USD 45,000; crashing the critical activity costs USD 30,000 per week for up to two weeks; after one week of crashing, a parallel path becomes co-critical. The value-maximising decision is: - A. crash two weeks (net +USD 30,000) - B. crash one week (net +USD 15,000) - C. crash nothing (avoid all cost) - D. crash two weeks and fast-track the parallel path at no cost

Rationale: Week one converts to a project week: +15,000. Week two is absorbed by the co-critical path: -30,000. A prices weeks that don't materialise; C leaves +15,000 unclaimed; D invents a free fast-track — overlap always carries rework risk.

MCQ 6.4-B [6.4.3 · Application] An activity is estimated $o = 4$, $m = 5$, $p = 12$ weeks. Its PERT expected duration is: - A. 5.0 weeks - B. 7.0 weeks - C. 6.0 weeks - D. 6.5 weeks

Rationale: $(4 + 4 \times 5 + 12)/6 = 36/6 = 6.0$. A is the most-likely value (the mode); B is the unweighted mean of the three points; D miscounts the weighting.

MCQ 6.4-C [6.4.1 · Analysis] A hybrid programme needs a software module from an agile team for an integration milestone. The schedule-sound way to join the two worlds is: - A. impose the CPM date on the team as a sprint deadline - B. enter the team's velocity-based forecast as a ranged duration and read the latest acceptable delivery from the backward pass - C. exclude the module from the network since agile work cannot be scheduled - D. convert the agile team to predictive planning for the integration period

Rationale: B translates in both directions — evidence-based forecast in, latest-start requirement out. A dictates without evidence; C leaves the network blind at a convergence point; D destroys the team's delivery system to decorate the network.

MCQ 6.4-F [6.4.3 · Application] An activity is estimated optimistic 3, most-likely 4, pessimistic 8 weeks. Its PERT expected duration and standard deviation are: - A. $t_e = 4.5$, $\sigma = 0.83$ - B. $t_e = 4.0$, $\sigma = 0.83$ - C. $t_e = 5.0$, $\sigma = 1.67$ - D. $t_e = 4.5$, $\sigma = 5.0$

Rationale: $t_e = (3 + 16 + 8)/6 = 4.5$; $\sigma = (8 - 3)/6 = 0.83$. B reports the mode as the mean; C is the unweighted three-point average with a doubled spread; D confuses σ with the pessimistic-minus-optimistic range.

MCQ 6.4-D [6.4.3 · Application] At week 22, a programme has earned the value its baseline planned to earn by week 20. Its time-based schedule performance index $SPI(t)$ is: - A. 1.10 - B. 0.91 - C. 0.80 - D. 20.0

Rationale: $SPI(t) = ES/AT = 20/22 = 0.91$ — the programme delivers at 91 % of planned tempo. A inverts the ratio; C subtracts the two weeks from the wrong base; D reports earned schedule itself, not the index.

MCQ 6.4-E [6.4.3 · Analysis] Compared with quoting "78 % confidence of the date", giving the board three fully re-run schedule scenarios is stronger because: - A. three numbers always beat one - B. each scenario is an auditable network with named assumptions, showing how the date moves and where the path migrates - C. percentages are always statistically invalid - D. scenarios eliminate the need for risk analysis

Rationale: The scenario's power is auditability and mechanism — assumptions with owners, visible path migration. A is numerology; C overclaims (a calibrated percentage is legitimate, Domain 8); D reverses the relationship — scenarios are inputs to risk analysis, not substitutes.

■ SELF-CHECK — KA 6.4

1. Why does compression show diminishing returns? — Each crash promotes the next-longest path; saved weeks stop converting to project weeks (path migration).
2. What makes a scenario more useful to a board than a confidence percentage? — It is a complete, auditable network with named assumptions — it shows how the date moves, not just that it might.
3. When may a machine forecast enter a board pack? — With a calibration record, verified arithmetic, stated data lineage and a named human owner.

— Advanced topics — Domain 6

6.A.1 The drum: constraint-aware flow

Where one resource is the system's constraint (a single commissioning team, one test rig), throughput scheduling subordinates everything to that drum: buffer it, never starve it, and let non-constraint efficiency go. This reframes utilisation worship — a 100 %-busy non-constraint builds inventory, not progress. In schedule terms: protect the constraint's feed path with deliberate float placement rather than spreading buffer evenly.

6.A.2 Schedule baselines and the update cycle

The control schedule lives on a cycle: status (actuals and remaining durations in), re-run passes, variance vs baseline, forecast, and — rarely, formally — re-baseline under Domain 4's change control. Two integrity rules carried from the family's controls tradition: never edit history (actualised dates are records), and never re-baseline to hide variance — a re-baseline is a governance decision with an audit trail, not a cosmetic one.

6.A.3 PDM edge cases: float under SS and FF links

Precedence-diagram (PDM) links change how float behaves, and two edge cases bite reviewers. **(1) The SS-linked short successor.** P (10 weeks) **SS+2** Q (3 weeks): Q can run weeks 2–5, finishing seven weeks before its predecessor — legal, and correct if the dependency truly binds only the starts; but if Q also needs P's output to finish, the network is missing an FF link, and Q's float is fiction. Paired **SS+FF** links express "starts together, finishes together" honestly. **(2) Duration-driven float.** Under an FF link the successor's float depends on its own duration: lengthen Q and its **LS** moves earlier — planners who "add duration for safety" on FF-linked activities silently destroy their float. The audit habit: on every SS/FF link, ask which end of each activity the logic genuinely binds, and whether the unbound end needs its own link. A network is only as honest as its least-considered link type.

6.A.4 Negative float and the honest schedule

A constraint-bound network can compute $TF < 0$: the plan, as constrained, is late already. Negative float is information — the size of the recovery problem — and suppressing it (deleting constraints, shortening durations by fiat) is the scheduling equivalent of cooking the books. The professional response mirrors KA 6.4.2: quantify it, trace the binding path, price the recovery options, and escalate the decision to whoever owns the trade (Domain 3's escalation design).

WORKED EXAMPLE**Worked example 6.A.4 — buying back two weeks of negative float.**

1. **Setup.** A newly imposed outage constraint requires Auriga to finish by **week 23**; logic says 25 — so $TF = -2$ on A-B-C-E-F. Available buys: crash C at USD 30,000/week (max 2); crash E at USD 55,000/week (max 1). Find the cheapest feasible recovery to week 23.
2. **Formula.** Recover week by week on the currently binding path(s), re-running the passes after each buy (path migration, 6.4.2).
3. **Substitution.** Week one: crash C by 1 → duration 24, but B-D-E-F is now co-critical at 24. Week two must shorten both paths: the only shared activity available is E → crash E by 1 (USD 55,000) → duration 23. (C's second week alone would leave B-D-E-F at 24 — spent money, no schedule.)
4. **Result.** Feasible recovery: **crash C × 1 + crash E × 1 = USD 85,000**, finishing week 23.
5. **Interpretation.** Negative float is bought back on the binding path as it migrates — the second-cheapest week on the original critical path (C's second week, USD 30,000) is worthless, while the dearest activity (E, on both paths) is the only week-two purchase that works. This is why recovery plans list options in re-run order, not in unit-cost order (toolkit 6.T.2).

— Industry variations — Domain 6

The passes are universal; what counts as a schedule, and what discipline it needs, is sectoral:

- **Construction and EPC.** Deep L3/L4 hierarchies, thousand-activity networks, contractual float-ownership clauses (Domain 10) and delay-analysis protocols make float a legal quantity — the honesty settings of the Executive perspective are contract compliance here, not just culture.
- **Technology and product.** Hybrid is the default (KA 6.4.1): cadence-based streams joined to a thin predictive network at integration milestones; the commonest defect is a network so sparse that convergence points are invisible until they slip.
- **Shutdowns and turnarounds.** The network runs in hours; near-critical analysis widens to every path within minutes of critical, and the drum logic of 6.A.1 governs a single shared crane or permit desk. Compression economics (6.4.2) turn over in shifts, not weeks.
- **Pharmaceutical and regulated development.** Milestones are regulatory events with submission-window physics — genuine date constraints (6.3.2), modelled as logic, with the pin-and-squeeze temptation at its most dangerous because windows are unforgiving.
- **Public programmes.** Announced dates precede networks (Case study B); the leader's protection — milestones only after the backward pass — is hardest exactly where it matters most, and negative-float honesty (6.A.4) is the difference between an early re-plan and a public failure.

— Case study — Domain 6: recovering Auriga (utilities / technology)

Situation. End of week 13. Civil and cabling works D — planned at 7 weeks from week 8 — hits contaminated-ground remediation; the discipline lead's honest re-estimate takes D to 10 weeks (finish week 18). The client's outage window for final commissioning opens at week 25 and closing the window costs USD 45,000 per week of delay. The team re-runs the passes.

Analysis. With $D = 10$: E starts week 18 (D now later than C's week 16), F finishes week 27 — **two weeks into the penalty window**, and the critical path has migrated to **A-B-D-E-F**. Procurement C, critical all project, now carries $TF = 2$. G still holds float.

The recovery decision. Options priced by the team: (1) crash the old critical path — worthless, C is no longer binding; (2) crash D's remediation with a second civil crew — USD 35,000 for one week; (3) fast-track E against D's tail — start installation in the uncontaminated sections at week 16, accepted rework risk assessed at $20\% \times \text{USD } 60,000 = \text{USD } 12,000$ expected cost, one further week saved; (4) do nothing — $2 \times 45,000 = \text{USD } 90,000$ penalty exposure. The leader takes (2) + (3): cost $\approx \text{USD } 47,000$ expected against USD 90,000 — and finishes at week 25 with the window intact.

What the domain teaches here. Re-run the passes before spending a cent (the instinctive "crash procurement" would have bought nothing); price recovery in expected cost including risk; harvest overlap where the physics genuinely allows it; and report the new float positions to the client honestly — the schedule survived because the remediation estimate was honest a full five weeks before the window (Domain 11's reporting culture, applied to time).

— Case study B — Domain 6: the milestone that was typed, not computed (public programme)

Situation. A government services programme publishes a go-live milestone eighteen months out, set in a ministerial announcement before the L3 schedule existed. When the network is finally built, the backward pass puts the milestone's required start for user-acceptance testing three weeks before the environment that testing needs can exist — $TF = -3$ from day one.

What happened. For six reporting cycles the programme's L1 view showed the milestone green: the date was pinned, and each slip downstream was absorbed by silently shortening the testing window — the classic pin-and-squeeze. In cycle seven an assurance review (Domain 3, KA 3.3) ran the passes without the pin, surfaced the negative float, and forced the choice the pin had deferred: descope the first release (Domain 5), or move the date. The minister moved the date — at a political cost several times what a computed date would have cost eighteen months earlier.

What the domain teaches here. A typed milestone is a promise without a plan. Negative float is not embarrassment to be formatted away; it is the earliest, cheapest warning the programme will ever get (6.A.4). And the leader's protection is procedural, not heroic: milestones enter public commitments only after the backward pass says they exist.

— Executive perspective — Domain 6

What a project leader cannot delegate in this domain:

- **The logic sign-off.** Discipline leads own their links; the leader owns the assembled claim that this network is how the project actually works.
- **Float policy.** Who may spend float, in what increments, reported how — decided before the first progress update, not during the first crisis.
- **The honesty settings.** No pins doing logic's work; negative float reported the cycle it appears; re-baselining only through governance. Schedules fail morally before they fail mathematically.

- **The compression chequebook.** Crash and fast-track decisions are investments under uncertainty — the leader signs the expected-value arithmetic (6.4.2) and the risk acceptance (Domain 8), never a bare instruction to "go faster".
- **The translation duty.** Boards hear dates; teams live in passes, floats and scenarios. The leader is the honest converter between the two languages — both directions.

— Calculation exercises — Domain 6

Exercise 6.1 Using the Auriga network, compute **ES/EF** for every activity and the project duration, showing the convergence at E. Solution. A 0-2 · B 2-8 · C 8-16 · D 8-15 · E $\max(16,15)=16$ -21 · F 21-25 · G 15-17. Duration **25 weeks**; E's start is set by C (16), not D (15) — the convergence rule takes the later predecessor. Common error: taking the earlier predecessor at a merge (giving E 15-20 and a false 24-week project).

Exercise 6.2 Complete the backward pass and state **TF** and **FF** for D and G. Solution. From **LF** = 25: F 21/25 · E 16/21 · C 8/16 · G 23/25 · D **LF** = $\min(E \text{ LS } 16, G \text{ LS } 23) = 16$, **LS** = 9. **D: TF = 1, FF = $\min(16,15) - 15 = 0$. G: TF = 8, FF = 25 - 17 = 8**. Common error: computing D's **FF** against only E (giving 1) — free float tests every successor.

Exercise 6.3 The client offers USD 45,000 per week saved; crashing C costs USD 30,000/week (max 2 weeks); crashing E costs USD 55,000/week (max 1 week). Find the value-maximising plan. Solution. Crash C by 1: 25→24, +15,000. Second week of C: co-critical B-D-E-F holds 24 — -30,000, reject. Crash E by 1 (E is on both paths): 24→23, **45,000 - 55,000 = -10,000**, reject. **Plan: crash C by one week; net +USD 15,000**. Common error: crashing E first because it is "on more paths" — path count doesn't change its negative unit economics.

Exercise 6.4 Activity durations: o = 6, m = 8, p = 16 weeks. Compute **t_e** and **σ**, and state the deterministic bias. Solution. **t_e = (6 + 32 + 16)/6 = 9.0 weeks**; **σ = (16 - 6)/6 = 1.67 weeks**. The deterministic plan at m = 8 understates the expectation by a full week — right-skewed tails always pull the mean above the mode.

Exercise 6.5 A programme's network shows **TF** = -2 on the binding path. List, in cost order, the recovery families and the governance step each requires. Solution. (1) Logic re-choice/re-sequencing — planner's proposal, leader's sign-off; (2) float and resource harvesting from non-binding paths — float-policy authority; (3) crashing the binding path — expected-value case, budget authority; (4) fast-tracking — risk acceptance recorded in the register; (5) scope/acceptance renegotiation — sponsor and change control. Reporting the -2 honestly precedes all five (6.A.4).

— Practitioner's toolkit — Domain 6

Adoption-ready artefacts; adapt headings to your organisation, then keep them stable.

Toolkit 6.T.1 — Schedule quality gate (run before trusting any network)

- [] Zero dangles; zero loops; every lag has a stated reason.

- [] Date constraints listed, each justified as genuine external physics — none doing logic's job.
- [] Durations owned by the executing discipline, estimate class stated (far waves: ranged).
- [] Passes re-run after every change; critical-path durations sum exactly to project duration.
- [] Near-critical paths (within 10 %) listed with their floats.
- [] Float report distinguishes **TF** and **FF**; float policy names who may spend it.
- [] Milestones computed, owned, done-defined; no typed dates.
- [] AI-drafted logic or forecasts marked, verified, and owned by a named human.

Toolkit 6.T.2 — Recovery plan (one page)

Slip statement (activity, cause, size, date detected) · re-run pass results (new duration, new critical path, float table) · options priced in expected cost (re-sequence / harvest / crash / fast-track / renegotiate), each with risk and owner · selected plan and its decision authority · revised commitments and the stakeholder notice list · review date.

Toolkit 6.T.3 — Milestone register

Per milestone: computed date and current **TF** · owner · done-definition (acceptance reference) · contractual status and counterparty · L3 logic trace · reporting treatment (internal / contractual / public) · change history.

— Exam preparation — Domain 6

The calculation traps. Taking the earlier predecessor at a convergence (Exercise 6.1) · computing **FF** against one successor instead of all (Exercise 6.2) · quoting **TF** after it has been spent · crashing from the original network instead of re-running passes (path migration, 6.4.2) · treating the PERT mode as the mean (6.4.3) · reading a pinned milestone as a forecast · confusing project-level float with activity-level slack in reports.

Reflection questions. 1. Your programme's board pack shows every milestone green. What three questions establish whether that is information or formatting? (Computed or typed? Passes re-run this cycle? Where is the float table?) 2. A subcontractor asks for "the float" on their package. What must you settle before answering? (Which float — **TF** or **FF**; the float policy; contractual float ownership — Domain 10.) 3. When did you last see money spent on compression that the passes would have shown was worthless — and which governance step was missing? (6.4.2; toolkit 6.T.2.)

— Domain 6 summary

This domain turned "when will it finish?" from an opinion into an instrument. Planning levels give each audience a true view of one logic; dependency discipline makes the network a model rather than a drawing; the forward and backward passes yield the dates, the two floats and the critical path — with the master project showing why free float can vanish while total float survives, and why convergence points compound slippage. Resources, constraints, milestones and rolling waves connect the network to the real world's crews, windows and horizons. Delivery flow closes the loop: compression priced by expected value and re-run passes, recovery ranked from re-choice to renegotiation, PERT and scenarios replacing single dates with honest ranges, and agile streams joined to the network through ranged integration milestones. The leadership spine throughout: schedules fail morally before they fail mathematically — computed dates, visible float, reported

negative float, and machine output verified by named humans are what make a programme's word worth something.

07

DOMAIN 7

Cost, Resources and Commercial Awareness (quantitative flagship)

Group: Delivering the work (Domain 7 of 6 in Part Two — the cost counterpart to Domain 6's schedule).

Target: ~78 pages. **Binds to:** the PCI Book Pattern Specification and the shared registries ([docs/books/registries/](#)). This domain is this book's home of the earned-value symbols — **PV**, **EV**, **AC**, **BAC**, **CV**, **SV**, **CPI**, **SPI**, **EAC**, **ETC**, **VAC**, **TCPI** — carried unchanged from the PCI family master table. Note the family notation-clash rule: **PV** here is **Planned Value**; present value is written **PV(x)** or in words (PFL-AI, Domain 3). British English; USD (+SAR where useful, indicative **USD 1 ≈ SAR 3.75**).

— Why this domain exists

A leader who can defend a date but not a number is half-equipped. Domain 6 built the schedule; this domain builds the money that pays for it and the commercial arrangements that bind other people to deliver it. It starts with estimating and budgeting — how a credible number is constructed and what its accuracy claim actually means (KA 7.1); establishes the cost baseline and the forecasting that keeps it honest (KA 7.2); builds **earned value** in full, because it is the only technique that answers cost and schedule performance in one integrated language (KA 7.3); and closes with the commercial awareness a leader cannot delegate — resource economics, procurement strategy, contract models and cash flow (KA 7.4). Risk quantification deepens in Domain 8; contracts and supply networks get their own treatment in Domain 10. What belongs here is the leader's own numeracy: enough to know when a forecast is arithmetic and when it is hope.

Learning objectives. After this domain a candidate can: select an estimating method appropriate to the available definition and state its accuracy class; build a three-point cost estimate; assemble a cost baseline separating contingency from management reserve; compute and interpret the full earned-value set (**CV**, **SV**, **CPI**, **SPI**); forecast with each **EAC** method and choose the one matching the variance's cause; compute **VAC** and **TCPI** and explain what **TCPI** demands of the remaining work; read resource economics through blended rates; distinguish the main contract models by who carries cost risk; compute a cost-incentive fee and the point of total assumption; explain why cash flow and profit differ on a project; and govern AI-produced cost forecasts under the family verification rule.

The master worked project. Project Auriga continues from Domain 6 — the 25-week control-systems upgrade for a regional utility — now with money attached. Its approved cost baseline is **BAC = USD 4,000,000**. At the **data date, end of week 13**, the baseline says **PV = USD 2,080,000**; measurement gives **EV = USD 1,920,000** and **AC = USD 2,120,000**. Every calculation in KA 7.2–7.3 uses these four numbers.

KNOWLEDGE AREA 7.1

Estimating and budgeting

Topics: 7.1.1 estimating methods and accuracy classes · 7.1.2 three-point estimating · 7.1.3 from estimate to budget.

7.1.1 Estimating methods and accuracy classes

The principle. An estimate's accuracy is governed by how well the work is defined, not by how much effort went into the arithmetic. The three standard methods trade definition against speed:

Method	How it works	Needs	Typical use
Analogous (top-down)	Scale a comparable past project	A true analogue and a scaling basis	Concept, screening (Domain 2)
Parametric	Cost = rate × parameter (per metre, per point, per kW)	A calibrated rate from real history	Early design, repetitive work
Bottom-up	Estimate each work package, then sum	A WBS at package level (Domain 4)	Baseline setting, control

Accuracy classes. Mature cost practice publishes an estimate's class alongside its number — a range that narrows as definition matures (the AACE Total Cost Management framework's class-5-to-class-1 progression is the reference treatment, described here in this book's own words). The professional discipline is simple and widely broken: **an estimate is never a single number without a range and a class.** "USD 4 million" is not an estimate; "USD 4 million, −15 %/+30 %, at a class consistent with 30 % design completion" is.

Common pitfall — precision mistaken for accuracy. A bottom-up estimate summing 400 packages to **USD 4,183,662** looks authoritative and is no more accurate than its worst assumption. Leaders should be more suspicious of an over-precise number than a rounded one.

7.1.2 Three-point estimating

Cost, like duration (Domain 6, KA 6.4.3), is a distribution. The same PERT weighting applies:

$$C_e = (o + 4m + p) / 6 \quad \sigma = (p - o) / 6$$

WORKED EXAMPLE**Worked example 7.1.2 — Auriga's control-hardware package.**

1. **Setup.** Procurement of control hardware is estimated **optimistic USD 680,000, most-likely USD 750,000, pessimistic USD 1,000,000** (the tail: a single-source controller with volatile lead-time pricing).
2. **Formula.** $C_e = (o + 4m + p)/6$; $\sigma = (p - o)/6$.
3. **Substitution.** $C_e = (680,000 + 3,000,000 + 1,000,000)/6 = 4,680,000/6$; $\sigma = 320,000/6$.
4. **Result.** $C_e = \text{USD } 780,000$; $\sigma \approx \text{USD } 53,333$.
5. **Interpretation.** The most-likely figure is 750,000, but the expected cost is 780,000 — the right-skewed tail adds USD 30,000 before anything goes wrong. Budgeting at the mode systematically under-funds a portfolio of such packages; the difference between mode and mean, summed across a project, is a large part of what contingency exists to cover (7.1.3, and Domain 8's quantification).

7.1.3 From estimate to budget

The structure. An estimate becomes a controllable budget through a deliberate hierarchy:

work-package estimates	
→ control-account budgets	(WBS × organisation – the management-control points)
→ + contingency reserve	→ the COST BASELINE (BAC): the PM's authority
→ + management reserve	→ the total project funding requirement

Two rules carry the discipline. **Contingency is inside the baseline**, sized against identified risks (Domain 8) and spent by the project manager under a stated protocol. **Management reserve is outside the baseline**, held for unknown-unknowns, and released only by the sponsor or change authority (Domain 3's decision rights) — releasing it changes the baseline through change control (Domain 4, KA 4.4). Blurring the two is how projects appear to be "on budget" while consuming their own funding runway.

Time-phasing. The baseline is spread across the schedule to produce the **cumulative cost curve** (the S-curve) — and that curve is Planned Value (KA 7.3). A budget with no phasing cannot be measured against; the schedule (Domain 6) is therefore a precondition of cost control, not a parallel activity.

AI in this KA

Estimating is where AI assistance is most seductive and most in need of provenance. A model can produce a plausible parametric rate instantly, and it cannot tell you whether that rate came from comparable work, a different market, or nowhere at all. The governed workflow: **AI proposes** a method, a rate and a range; the estimator supplies the calibration evidence (which projects, which years, which escalation basis — the source discipline of the shared registry); the range and class are stated; and a named human owns the number. **AI proposes; the professional verifies, decides and remains accountable.** An estimate whose basis cannot be produced on request is not an estimate.

■ KEY TERMS — KA 7.1

Term	Meaning
Analogous / parametric / bottom-up	Scale an analogue · rate × parameter · sum the packages.
Accuracy class	The definition-linked range accompanying an estimate; never optional.
Control account	The WBS × organisation point where scope, budget and actuals integrate.
Contingency reserve	Inside the baseline; funds identified risks; PM-controlled.
Management reserve	Outside the baseline; unknown-unknowns; sponsor-controlled.
BAC	Budget at Completion — the time-phased cost baseline's total.

■ SAMPLE MCQS — KA 7.1

MCQ 7.1-A [7.1.2 · Application] A package is estimated o = 680,000, m = 750,000, p = 1,000,000. Its PERT expected cost is: - A. USD 750,000 - B. USD 780,000 - C. USD 810,000 - D. USD 840,000

Rationale: $(680,000 + 4 \times 750,000 + 1,000,000)/6 = 780,000$. A is the mode; C is the unweighted three-point mean; D over-weights the pessimistic value.

MCQ 7.1-B [7.1.3 · Analysis] A project reports "on budget" while having consumed 60 % of its management reserve at 40 % complete. The correct reading is: - A. genuinely on budget — management reserve exists to be spent - B. the baseline is intact but the project's total funding is eroding faster than progress; the trend belongs in the next report to the sponsor - C. a baseline breach requiring immediate re-baselining - D. an accounting error, since management reserve is inside the baseline

Rationale: Management reserve sits outside the baseline (so D is wrong and A is technically true but misleading), and its release is a sponsor-level signal, not a project-level convenience. It is not yet a breach (C) — it is the early warning that precedes one.

MCQ 7.1-C [7.1.1 · Recall] Which statement about a bottom-up estimate summing to USD 4,183,662 is soundest? - A. its precision indicates high accuracy - B. its accuracy is bounded by its assumptions and definition maturity, and it must still carry a range and class - C. rounding it would reduce its accuracy - D. bottom-up estimates do not need accuracy classes

Rationale: Precision (digits) and accuracy (closeness to outturn) are independent; definition maturity governs the latter. Rounding changes no information (C), and every estimate carries a class (D).

■ SELF-CHECK — KA 7.1

1. Where do contingency and management reserve sit, and who spends each? — Contingency inside the baseline, PM-controlled; management reserve outside it, sponsor-controlled.
2. Why is the time-phased baseline a precondition for cost control? — Because the phased cumulative curve is Planned Value; without it there is nothing to measure performance against.
3. What two things must accompany every estimate? — A range and an accuracy class tied to definition maturity.

KNOWLEDGE AREA 7.2

The cost baseline, actuals and forecasting

Topics: 7.2.1 measuring actual cost · 7.2.2 the forecasting question · 7.2.3 baseline integrity.

7.2.1 Measuring actual cost

The principle. AC must cover the same work as EV, in the same period, or every index built from them is fiction. Three mechanics decide whether it does:

- **Accruals.** Work received but not yet invoiced belongs in this period's AC. Omit accruals and cost performance looks excellent until the invoices land — the commonest cause of a "sudden" overrun that was months old.
- **Commitment vs actual.** A purchase order is a commitment, not a cost. Both matter: AC drives performance, commitments drive the funding forecast. Confusing them double-counts or hides money.
- **Open-commitment hygiene.** Stale purchase orders left open inflate the forecast; cleansing them is a standing month-end task.

7.2.2 The forecasting question

A forecast answers one question — what will this cost when it is done? — and the honest answer depends entirely on **why** the current variance exists. That judgment is the leader's; the arithmetic is KA 7.3.3's EAC family. The rule to internalise now: **a forecast is a statement about the remaining work, not an extrapolation ritual.** A variance caused by a closed, non-recurring event says nothing about what is left; a variance caused by a systemic productivity shortfall says everything.

7.2.3 Baseline integrity

The baseline's authority rests on its stability. Four standing rules: budgets of completed or open work packages are never retrospectively adjusted (except to correct an authorised error); re-baselining happens only through change control, with an audit trail (Domain 4, KA 4.3); transfers between control accounts are logged, not silent; and the baseline never absorbs a variance by moving budget to where the money went. A baseline edited to match reality has stopped measuring anything — the cost analogue of Domain 6's pinned milestone.

■ KEY TERMS — KA 7.2

Term	Meaning
AC	Actual Cost of the work performed, including period accruals.
Accrual	Cost of work received but not yet invoiced, recognised in the period.
Commitment	A contractual obligation (e.g. a PO); a funding fact, not yet a cost.
Re-baselining	A governed change to the baseline; never a variance-hiding device.

■ SAMPLE MCQS — KA 7.2

MCQ 7.2-A [7.2.1 · Analysis] A project's CPI has read 1.02 for four months; then one month it drops to 0.91 with no change in productivity. The likeliest explanation is: - A. genuine sudden

inefficiency - B. accruals were not being recognised, so earlier AC understated cost and this month absorbed the catch-up - C. the baseline was too generous - D. EV was over-claimed this month

Rationale: A step change in CPI without an operational change points at measurement, and missing accruals are the classic cause — earlier periods flattered, one period punished. D would raise, not lower, the earlier readings' credibility, and C would show as a stable, not stepped, pattern.

MCQ 7.2-B [7.2.3 · Recall] Which action preserves baseline integrity? - A. moving budget from an underspent control account to cover an overspend, without record - B. re-baselining through change control with an audit trail when scope genuinely changes - C. adjusting a completed package's budget to match its actual cost - D. reducing remaining budgets so the total still equals BAC

Rationale: Only governed change preserves the measurement. A is a silent transfer, C is retrospective adjustment, D is the classic "make the numbers add up" manoeuvre — each destroys comparability.

■ SELF-CHECK — KA 7.2

1. Why must AC include accruals? — So it covers the same work as EV in the same period; otherwise every index is misstated.
2. Is a purchase order a cost? — No: a commitment. It informs the funding forecast, not performance.
3. What single question does a forecast answer? — What the remaining work will cost, given why the variance exists.

KNOWLEDGE AREA 7.3

Earned value: measurement, variances and forecasting

Topics: 7.3.1 the three measures · 7.3.2 variances and indices · 7.3.3 the EAC family · 7.3.4 VAC and TCPI.

7.3.1 The three measures

Definitions. All three are expressed in the same budget currency so they are directly comparable:

- **Planned Value (PV)** — the budgeted cost of the work scheduled by the data date: the time-phased baseline of KA 7.1.3.
- **Earned Value (EV)** — the budgeted cost of the work actually performed: physical progress **valued at the budget rate**, never at what it cost.
- **Actual Cost (AC)** — the cost actually incurred for that work, accruals included (KA 7.2.1).

The single most important conceptual point: EV is measured at budget. That is what lets EV be compared with PV (both at budget → schedule progress) and with AC (budget vs actual for the same work → cost efficiency). Confusing "value earned" with "cost incurred" collapses the method.

Earning rules. How EV is claimed decides how much it can be gamed: **0/100** (nothing until complete — objective, good for short packages); **50/50; percent complete** (needs an objective basis); **units completed** (best where output is countable); **weighted milestones**; and **level of effort**, where EV is set equal to PV by the calendar, so it can never show a schedule variance. Level of effort dilutes whatever it is mixed into — formal practice segregates it and caps its share of a control account.

7.3.2 Variances and indices

$$\begin{array}{ll} CV = EV - AC & SV = EV - PV \\ CPI = EV / AC & SPI = EV / PV \end{array}$$

WORKED EXAMPLE

Worked example 7.3.2 — Auriga at week 13.

- Setup.** $BAC = 4,000,000$; at the data date $PV = 2,080,000$, $EV = 1,920,000$, $AC = 2,120,000$.
- Formula.** As above; variances in currency, indices as ratios (two decimals).
- Substitution.** $CV = 1,920,000 - 2,120,000$; $SV = 1,920,000 - 2,080,000$; $CPI = 1,920,000 / 2,120,000$; $SPI = 1,920,000 / 2,080,000$.
- Result.** $CV = (\text{USD } 200,000)$ · $SV = (\text{USD } 160,000)$ · $CPI = 0.91$ · $SPI = 0.92$. Auriga is **48.0 % complete** (EV/BAC) having spent **53.0 % of budget** (AC/BAC).
- Interpretation.** Behind and over: for every dollar spent the project is producing 91 cents of budgeted work, and it has delivered 92 % of what it planned to by now. The percent-complete versus percent-spent pair (48 vs 53) is the same story in the form a board absorbs fastest. Note this is the cost view of the very slippage Domain 6's case study recovered on the schedule side — the two domains are describing one project.



Fig 7.3.1 — Auriga's earned-value S-curves at week 13

7.3.3 The EAC family — forecasting

ETC = the remaining work's forecast cost	EAC = AC + ETC
(a) EAC = AC + (BAC - EV)	remaining work at the BUDGETED rate
(b) EAC = BAC / CPI	current cost efficiency PERSISTS
(c) EAC = AC + (BAC - EV) / (CPI × SPI)	cost AND schedule pressure compound
(d) EAC = AC + bottom-up ETC	re-estimate the remainder from scratch

WORKED EXAMPLE

Worked example 7.3.3 — the same project, four futures.


- Setup.** Auriga's week-13 figures, forecast to completion by each method.
- Formula.** (a)–(c) above; **CPI** = 0.905660, **SPI** = 0.923077 at full precision.
- Substitution.** (a) $2,120,000 + (4,000,000 - 1,920,000)$; (b) $4,000,000 / 0.905660$; (c) $2,120,000 + 2,080,000 / (0.905660 \times 0.923077)$.
- Result.** (a) **USD 4,200,000** · (b) **USD 4,416,667** · (c) **USD 4,608,056**.
- Interpretation.** The same four measured numbers support forecasts spanning **USD 408,056** — over 10 % of **BAC**. The spread is not uncertainty in the arithmetic; it is the assumption about the remaining work, made explicit. Choosing among them is the leader's judgment, and it must be stated: if Auriga's overrun was the contaminated-ground remediation of Domain 6 — a discrete, now-closed event — method (a) is right and the others over-forecast. If it reflects a systemic productivity shortfall that will persist, (b). If schedule pressure is now driving overtime and expediting, (c). A forecast presented without its assumption named is not a forecast; it is a number.

7.3.4 VAC and TCPI

VAC = BAC - EAC	the forecast overrun/underrun
TCPI = (BAC - EV) / (BAC - AC)	efficiency the REMAINING work must achieve to hit BAC
TCPI = (BAC - EV) / (EAC - AC)	...to hit a revised EAC

WORKED EXAMPLE**Worked example 7.3.4 — what recovery would actually require.**

1. **Setup.** Auriga's week-13 figures; test recovery to **BAC**, and to the method-(b) forecast.
2. **Formula.** $VAC = BAC - EAC$; **TCPI** as above.
3. **Substitution.** $VAC(a) = 4,000,000 - 4,200,000$; $VAC(b) = 4,000,000 - 4,416,667$. $TCPI(BAC) = 2,080,000 / (4,000,000 - 2,120,000) = 2,080,000/1,880,000$. $TCPI(EAC\ b) = 2,080,000 / (4,416,667 - 2,120,000)$.
4. **Result.** $VAC(a) = (\text{USD } 200,000)$ · $VAC(b) = (\text{USD } 416,667)$. **TCPI to BAC = 1.11**; **TCPI to the (b) forecast = 0.91**.
5. **Interpretation.** To finish at budget, the remaining 52 % of work must run at **1.11** — an 11 % efficiency gain by a team currently achieving 0.91. That is a 22-point swing, and the honest question in the room is what specifically would deliver it. Meanwhile **TCPI to the (b) forecast is 0.91** — exactly today's **CPI**, which is the arithmetic identity worth internalising: **forecasting with BAC/CPI is precisely the assumption that nothing changes**. **TCPI** is the reality check on every recovery promise: a required index far above demonstrated performance is a plan that needs a mechanism, not encouragement.

**Fig 7.3.2 — The 'EAC' fan and what 'TCPI' demands****AI in this KA**

Earned value is arithmetic on four numbers, so machine computation is safe and instant verification is cheap — which makes unverified AI output indefensible here rather than merely risky. The deterministic checks: EV/BAC must equal the claimed percent complete; $CV = EV - AC$ and $CPI = EV/AC$ recomputed independently; **TCPI to BAC/CPI** must equal **CPI** (7.3.4's identity); and every **EAC** must state its assumption. Where AI forecasts from trend data, the calibration rule of Domain 6 (KA 6.4) applies unchanged: show the record of past forecasts against outturn before the number enters a board pack, and name the human who owns it.

■ KEY TERMS — KA 7.3

Term	Meaning
PV EV AC	Budgeted cost of work scheduled · performed · the cost actually incurred.
CV SV	$EV - AC$ · $EV - PV$, in currency.
CPI SPI	EV/AC · EV/PV , as ratios.
EAC ETC	Estimate at / to complete; $EAC = AC + ETC$.
VAC	$BAC - EAC$ — forecast variance at completion.
TCPI	Efficiency the remaining work must achieve to hit a stated target.
Level of effort	Earning by calendar ($EV = PV$); can never show schedule variance.

■ SAMPLE MCQS — KA 7.3

MCQ 7.3-A [7.3.2 · Application] BAC 4,000,000; PV 2,080,000; EV 1,920,000; AC 2,120,000. CPI and SPI are: - A. 0.91 and 0.92 - B. 1.10 and 1.08 - C. 0.92 and 0.91 - D. 0.91 and 1.02

Rationale: $CPI = 1,920,000/2,120,000 = 0.91$; $SPI = 1,920,000/2,080,000 = 0.92$. B inverts both ratios; C swaps them (dividing EV by the wrong denominator); D miscomputes SPI against BAC-derived progress.

MCQ 7.3-B [7.3.3 · Analysis] The overrun was caused by a one-off ground-remediation event, now closed, and the remaining work is expected to run to budget. The appropriate EAC is: - A. $AC + (BAC - EV) = 4,200,000$ - B. $BAC/CPI = 4,416,667$ - C. $AC + (BAC - EV)/(CPI \times SPI) = 4,608,056$ - D. $BAC = 4,000,000$

Rationale: A discrete, closed cause makes the variance **atypical**, so remaining work is forecast at the budgeted rate. B assumes the inefficiency persists and C that it compounds with schedule pressure — both contradict the stated cause. D ignores money already spent above budget.

MCQ 7.3-C [7.3.4 · Application] With BAC 4,000,000, EV 1,920,000 and AC 2,120,000, the TCPI required to complete at BAC is: - A. 0.91 - B. 1.00 - C. 1.11 - D. 1.21

Rationale: $(4,000,000 - 1,920,000)/(4,000,000 - 2,120,000) = 2,080,000/1,880,000 = 1.11$. A is the demonstrated CPI; B assumes recovery needs only par performance; D uses PV in place of AC in the denominator.

MCQ 7.3-D [7.3.1 · Analysis] A control account is 70 % level-of-effort by budget. Its reported SPI of 1.00 most likely means: - A. the account is exactly on schedule - B. little about schedule: level of effort earns by the calendar, so $EV = PV$ for most of the account regardless of progress - C. the discrete work is ahead, offsetting a delay - D. the earning rules were misapplied

Rationale: LOE sets EV equal to PV by construction, so a heavily-LOE account reads 1.00 whatever happens — which is why practice segregates and caps it. C invents an offset the data cannot show; the rules may have been applied entirely correctly (D), and that is the problem.

■ SELF-CHECK — KA 7.3

1. Why is EV measured at budget, not at cost? — So it is comparable with PV (schedule) and with AC (efficiency); at cost it would collapse into AC.
2. What does TCPI = 1.11 against CPI = 0.91 tell a leader? — Recovery to budget requires a 22-point efficiency swing; the plan needs a named mechanism, not optimism.

3. State the identity linking *TCPI* and *BAC/CPI*. — *TCPI* to an *EAC* of *BAC/CPI* equals the current *CPI*: that forecast is the assumption that nothing changes.

KNOWLEDGE AREA 7.4

Resource economics, procurement strategy and cash

Topics: 7.4.1 resource economics and blended rates · 7.4.2 contract models and cost risk · 7.4.3 incentive fees and the point of total assumption · 7.4.4 cash flow versus profit.

7.4.1 Resource economics and blended rates

Most project cost is people. A leader who reasons in headcount rather than **cost per unit of capability** will mis-price every option in Domain 6's compression decisions.

WORKED EXAMPLE

Worked example 7.4.1 — Auriga's engineering blended rate.

- Setup.** The engineering pool is 40 technicians at USD 95/hour, 25 engineers at USD 140/hour, 15 senior specialists at USD 210/hour.
- Formula.** Blended rate = $\Sigma(\text{count} \times \text{rate}) / \Sigma(\text{count})$.
- Substitution.** $(40 \times 95 + 25 \times 140 + 15 \times 210) / 80 = (3,800 + 3,500 + 3,150) / 80 = 10,450 / 80$.
- Result.** **USD 130.63 per hour** (130.625 exactly; $\approx$ SAR 490 indicatively).
- Interpretation.** The blended rate makes options commensurable: a week of schedule compression staffed from this pool costs a computable amount, and shifting the mix changes the price as much as changing the headcount — adding five specialists moves the blend more than adding five technicians. This is the number that turns Domain 6's "crash it" instruction into an expected-value decision.

7.4.2 Contract models and cost risk

Every contract model is an answer to one question: **who carries cost risk?**

Model	Buyer pays	Cost risk sits with	Suits
Firm fixed price (FFP)	An agreed price, whatever it costs	Seller	Well-defined scope; priced-in risk premium
Fixed price incentive (FPIF)	Target cost/fee with a share ratio, capped by a ceiling	Shared, then seller above the ceiling	Definable scope with real uncertainty
Cost plus fixed fee (CPFF)	Allowable cost + a fixed fee	Buyer	Development, unclear scope
Cost plus incentive fee (CPIF)	Allowable cost + a fee varying with performance	Shared	Scope where effort is uncertain but effort quality matters
Time and materials (T&M)	Rates $\times$ quantities	Buyer	Staff augmentation, short or open-ended work

The leader's discipline: **risk transferred is risk priced**. An FFP contract for ill-defined scope does not remove the risk — it converts it into a premium plus a claims exposure when the scope moves (Domain 10, KA 10.4). Conversely a cost-plus contract on well-defined work pays for uncertainty that no longer exists.

7.4.3 Incentive fees and the point of total assumption

The mechanics. Under an incentive contract, buyer and seller agree a **target cost**, a **target fee** and a **share ratio** (buyer/seller) for over- and under-runs, with a **ceiling price** capping the buyer's exposure.

WORKED EXAMPLE

Worked example 7.4.3 — the incentive that stops incentivising.

1. **Setup.** Auriga's installation subcontract: target cost **USD 2,000,000**, target fee **USD 150,000**, share ratio **70/30** (buyer 70 %, seller 30 %), ceiling price **USD 2,450,000**. The seller finishes at an actual cost of **USD 2,300,000**.
2. **Formula.** Fee = target fee – (overrun × seller share). Buyer pays = actual cost + fee (subject to the ceiling). Point of total assumption $PTA = \text{target cost} + (\text{ceiling} - \text{target price}) / \text{buyer share}$, where target price = target cost + target fee.
3. **Substitution.** Overrun $2,300,000 - 2,000,000 = 300,000$; fee $150,000 - 300,000 \times 0.30 = 150,000 - 90,000$. Buyer pays $2,300,000 + 60,000$. Target price $2,150,000$; $PTA = 2,000,000 + (2,450,000 - 2,150,000) / 0.70$.
4. **Result.** Fee **USD 60,000** (down from 150,000); buyer pays **USD 2,360,000**. $PTA = \text{USD } 2,428,571.43$ — and at that cost the buyer pays exactly the USD 2,450,000 ceiling.
5. **Interpretation.** The share ratio does its job up to the PTA : both parties lose money on overrun, so both want efficiency. **Above the PTA the seller bears 100 %** of further cost — which is the moment the incentive inverts: a seller heading past it has no financial reason to spend more on your project and every reason to argue that the extra cost is your scope change. Knowing where the PTA sits is therefore a **delivery** insight, not an accounting one: it predicts when a commercial relationship will turn adversarial (Domain 10, KA 10.4; Domain 11's negotiation).

7.4.4 Cash flow versus profit

A project can be profitable and still fail for lack of cash — the same truth PFL-AI establishes for financings (its Domain 1). For a delivery leader the mechanism is payment terms: cost is incurred as work happens; cash arrives when invoices are approved and paid. On Auriga, with $AC = \text{USD } 2,120,000$ at week 13 and 60-day terms, roughly **USD 742,000** of incurred cost (about 35 %) is still unpaid — an exposure carried by whoever is funding the work. The leader's obligations are practical: know the terms in both directions (client and subcontractors), front-load nothing you cannot fund, and never let a retention or milestone-payment structure be agreed without someone computing its cash profile. Where the project sits inside a financed asset, this is precisely the CFADS conversation PFL-AI Domain 10 has with lenders.

AI in this KA

Contract analytics — extracting terms, comparing models, flagging inconsistent clauses — is real and useful AI assistance, and it is also decision support rather than legal or commercial advice. Two boundaries: an AI-produced reading of a clause is verified against the clause itself before anyone

relies on it (the document-against-summary check), and commercial and legal positions go to qualified counsel (Domain 10). Fee arithmetic and PTA computations are deterministic — they get the same golden-answer treatment as the earned-value set.

■ KEY TERMS — KA 7.4

Term	Meaning
Blended rate	Weighted average cost per hour of a mixed resource pool.
FFP / FPIF / CPFF / CPIF / T&M	Contract models ordered by who carries cost risk.
Share ratio	The agreed buyer/seller split of over- and under-run.
Ceiling price	The cap on the buyer's total exposure under an incentive contract.
Point of total assumption (PTA)	The cost above which the seller bears 100 % of further overrun.
Retention	A withheld percentage of payment, released on completion criteria.

■ SAMPLE MCQS — KA 7.4

MCQ 7.4-A [7.4.3 · Application] Target cost 2,000,000; target fee 150,000; share 70/30; actual cost 2,300,000. The seller's fee is: - A. USD 150,000 - B. USD 60,000 - C. USD 90,000 - D. USD 45,000

Rationale: The seller absorbs 30 % of the 300,000 overrun: $150,000 - 90,000 = 60,000$. A ignores the incentive; C states the fee reduction rather than the fee; D applies the buyer's share to the fee.

MCQ 7.4-B [7.4.3 · Analysis] Target cost 2,000,000, target fee 150,000, ceiling 2,450,000, buyer share 70 %. The PTA is 2,428,571, and its delivery significance is that above it: - A. the contract becomes void - B. the buyer absorbs all further cost - C. the seller bears 100 % of further cost, so the incentive inverts and cost growth becomes a scope-change argument - D. the fee becomes negative but risk-sharing continues unchanged

Rationale: Beyond the PTA the ceiling binds the buyer, so every further dollar is the seller's — which predictably redirects the seller's effort from efficiency to entitlement. B reverses the exposure; A is fiction; D misses that sharing has stopped.

MCQ 7.4-C [7.4.2 · Analysis] A leader lets an FFP contract for scope that is only 30 % defined. The most likely outcome is: - A. cost risk is genuinely eliminated - B. a priced-in risk premium plus a claims-and-variations exposure as the scope is defined - C. the seller absorbs all scope growth at no cost to the buyer - D. the contract converts automatically to cost-plus

Rationale: Fixed price transfers risk at a price and only for the scope actually specified; undefined scope returns as variations. A and C mistake the contractual form for the underlying uncertainty; D invents a mechanism.

MCQ 7.4-D [7.4.1 · Application] A pool of 40 at USD 95/h, 25 at USD 140/h and 15 at USD 210/h has a blended rate of: - A. USD 148.33 - B. USD 130.63 - C. USD 115.00 - D. USD 140.00

Rationale: $10,450/80 = 130.63$. A averages the three rates unweighted; C weights toward the cheapest grade only; D takes the middle rate as representative.

■ SELF-CHECK — KA 7.4

1. State the one question every contract model answers. — Who carries cost risk.

2. Why is the **PTA** a delivery concern, not just a commercial one? — Above it the seller bears all further cost, so behaviour shifts from efficiency to entitlement.
3. How can a profitable project run out of cash? — Cost is incurred as work happens; cash arrives on payment terms — the gap must be funded.

— Advanced topics — Domain 7

7.A.1 Earned schedule — closing **SPI**'s late-project blind spot

SPI converges on 1.00 as a project finishes, whatever its lateness, because **EV** and **PV** both approach **BAC**. **Earned schedule** restates progress in time: **ES** is the date at which the value now earned was planned to have been earned, and $SPI(t) = ES / AT$. Auriga's week-13 status, with value earned that the baseline expected by week 12, gives $SPI(t) = 12/13 = 0.92$ — and unlike **SPI** it stays meaningful to the last week. This is the bridge flagged in Domain 6 (KA 6.4.3), now with its cost-side companions.

7.A.2 EVM's limitations, stated plainly

Earned value measures conformance to plan, not value to the customer: a project can score 1.00 on both indices while building the wrong thing (Domain 5's acceptance criteria are the defence). It says nothing about quality (Domain 9). It is only as honest as its earning rules and its accruals. And it is blind to risk that has not yet materialised — Domain 8's contingency analysis, not **CPI**, tells you whether what remains is adequately funded. EVM is a measurement system, not a management philosophy; leaders who treat the indices as targets get optimised indices.

7.A.3 The reviewer's cost eye

Invariants a reviewer runs before trusting a cost report: **EV/BAC** equals the claimed percent complete; $CV = EV - AC$ and $SV = EV - PV$ reconcile to the stated indices; $EAC = AC +$ a stated **ETC**, with the assumption named; **TCPI** to **BAC/CPI** equals **CPI**; sum of control-account budgets + contingency equals **BAC**; management reserve sits outside that total; no completed package's budget has moved since last period; accruals are present in the current period; and level-of-effort share is disclosed per control account. Any violation is a defect somewhere — find it before the board builds on it.

— Industry variations — Domain 7

- **Construction and EPC.** Full EVM against a resource-loaded baseline; measured-work valuations and retention dominate cash; FIDIC-style variation and claims machinery makes the **PTA** conversation routine.
- **Government and defence programmes.** Formal EVMS compliance with surveillance, strict LOE caps and control-account discipline; forecasting is auditable rather than discretionary.
- **Technology and product delivery.** Cost is overwhelmingly people, so blended rates and capacity, not materials, drive the number; where cadence replaces a fixed baseline, throughput and cost-per-increment stand in for **CPI** (Domain 13), and hybrid programmes report both.

- **Energy and resources.** Long-lead equipment and commodity exposure put escalation and currency at the centre of the estimate (PFL-AI Domain 3's machinery); contingency is sized quantitatively as a matter of course.
- **Public services and transformation.** Benefits, not cost variance, are the accountability currency (Domain 16); the leader's task is keeping cost reporting honest while the value case is what gets debated.

— Case study — Domain 7: the forecast the board actually needed (utilities)

Situation. Auriga's week-13 review. The cost report shows [CPI](#) 0.91, [SPI](#) 0.92, [CV](#) (200,000), [SV](#) (160,000). The programme director's paper proposes reporting [EAC USD 4,200,000](#) — method (a) — on the grounds that the overrun was the contaminated-ground remediation, a discrete event now closed and remediated (Domain 6's case study).

The challenge. The assurance reviewer asks three questions. Is the cause genuinely closed? Yes for the remediation — but the recovery bought a second civil crew and a fast-track of installation, both of which continue for six more weeks, and neither is in method (a)'s "remaining work at budgeted rate" assumption. What does recovery to [BAC](#) require? [TCPI](#) 1.11 against a demonstrated 0.91. What is the mechanism? There isn't one — the recovery spend increases cost to protect the date.

The outcome. The board is given a **range with named assumptions**: 4,200,000 if the closed event is the whole story; 4,416,667 if current efficiency persists; and a bottom-up (method d) re-estimate of the remaining 52 %, commissioned that week, as the number they will actually manage to. Contingency draw is authorised against the identified risk; management reserve is not touched. The minute records the [TCPI](#) gap and the explicit finding that **no recovery-to-budget plan exists** — so nobody later claims one was promised.

What the domain teaches here. A single-number forecast hides the only thing that matters: the assumption. [TCPI](#) converts optimism into an arithmetic claim someone has to defend, and the honest report is the one that survives the question "what would have to be true?"

— Case study B — Domain 7: past the point of total assumption (technology)

Situation. A systems-integration subcontract ran on an incentive structure: target cost USD 2,000,000, target fee 150,000, 70/30 share, ceiling 2,450,000 — so [PTA](#) USD 2,428,571. By month eight the supplier's internal cost was tracking to 2,600,000, well past the [PTA](#).

What happened. The supplier's behaviour changed abruptly and, in hindsight, rationally: fresh change requests on work previously treated as in-scope, slower responses on defect fixes that earned nothing, and a claim that the buyer's late environment provision had caused the growth. The buyer's team read it as bad faith. It was arithmetic: past the [PTA](#) every additional engineer-hour came out of the supplier's own margin, and the only route back to profit was re-characterising cost as buyer-caused scope.

The outcome. The parties re-set commercially — a re-baselined target cost recognising the genuine environment delay (documented, and partly the buyer's), a revised ceiling, and a time-and-materials envelope for the disputed remainder. Delivery recovered within two months of the reset. The

retrospective's finding: the buyer had never computed the PTA, so a foreseeable inflection was experienced as a betrayal.

What the domain teaches here. Commercial structures create behaviour. Computing the PTA at signature — and watching the supplier's cost trend against it — turns an adversarial surprise into a managed conversation held early, while both parties still have options (Domain 10's supplier governance; Domain 11's negotiation).

— Executive perspective — Domain 7

What a project leader cannot delegate in this domain:

- **The forecast's assumption.** Analysts compute EAC; the leader owns which method and why, in one sentence a board can challenge. Reporting a number without its assumption is the domain's cardinal sin.
- **The TCPI reality test.** Before endorsing any recovery-to-budget plan: what index does the remaining work require, what is being demonstrated, and what specific mechanism closes the gap?
- **Reserve discipline.** Contingency and management reserve kept distinct, spent under stated authority, with consumption trended against progress — the erosion in MCQ 7.1-B is a leadership signal, not a bookkeeping detail.
- **Measurement integrity.** Earning rules fixed in advance, accruals present, level-of-effort share disclosed, no retrospective budget edits. Cost systems fail morally before they fail arithmetically — the same sentence as Domain 6, and the same reason.
- **The commercial shape.** Which model, why, and where the PTA sits. A leader who cannot say who carries cost risk on their largest package is not yet in control of it.

— Calculation exercises — Domain 7

Exercise 7.1 BAC 4,000,000; PV 2,080,000; EV 1,920,000; AC 2,120,000. Compute CV, SV, CPI, SPI, percent complete and percent spent. Solution. CV (200,000); SV (160,000); CPI 0.91; SPI 0.92; complete $1,920,000/4,000,000 = 48.0\%$; spent $2,120,000/4,000,000 = 53.0\%$. Common error: computing percent complete from AC/BAC — that is percent spent, and reporting it as progress overstates the project.

Exercise 7.2 Same data. Compute EAC by methods (a), (b) and (c), and the corresponding VAC. Solution. (a) $2,120,000 + 2,080,000 = 4,200,000$, VAC (200,000). (b) $4,000,000/0.905660 = 4,416,667$, VAC (416,667). (c) $2,120,000 + 2,080,000/(0.905660 \times 0.923077) = 4,608,056$, VAC (608,056). Common error: rounding CPI to 0.91 before dividing — $4,000,000/0.91 = 4,395,604$, USD 21,063 adrift. Indices are display; arithmetic is full precision.

Exercise 7.3 Same data. Compute TCPI to BAC and to the method-(b) EAC, and interpret. Solution. To BAC: $2,080,000/1,880,000 = 1.11$. To (b): $2,080,000/(4,416,667 - 2,120,000) = 2,080,000/2,296,667 = 0.91$. The second equals the current CPI — the identity of 7.3.4: BAC/CPI forecasts are "nothing changes". Common error: using PV rather than AC in the denominator (giving 1.08) and concluding recovery is nearly free.

Exercise 7.4 Incentive subcontract: target cost 2,400,000; target fee 180,000; share 80/20; ceiling 2,900,000. The seller finishes at 2,650,000. Find the fee, what the buyer pays, and the PTA. Solution. Overrun 250,000; fee $180,000 - 250,000 \times 0.20 = \mathbf{130,000}$; buyer pays $2,650,000 + 130,000 = \mathbf{2,780,000}$ (below the 2,900,000 ceiling). Target price 2,580,000; $PTA = 2,400,000 + (2,900,000 - 2,580,000)/0.80 = 2,400,000 + 400,000 = \mathbf{2,800,000}$. Common error: applying the buyer's 80 % share to the fee reduction.

Exercise 7.5 A pool of 30 at USD 105/h, 20 at USD 150/h and 10 at USD 220/h. Compute the blended rate, then the cost of adding four weeks of a 5-person crew at 40 h/week. Solution. $(30 \times 105 + 20 \times 150 + 10 \times 220)/60 = (3,150 + 3,000 + 2,200)/60 = 8,350/60 = \mathbf{USD\ 139.17/h}$. Crew cost $5 \times 40 \times 4 \times 139.17 = \mathbf{USD\ 111,333}$. Common error: using the blended rate for a crew drawn entirely from one grade — the blend applies to a representative mix, not to any arbitrary subset.

— Practitioner's toolkit — Domain 7

Adoption-ready artefacts; adapt headings to your organisation, then keep them stable.

Toolkit 7.T.1 — Cost-report integrity checklist (run before publishing the month)

- [] Earning rules fixed in advance per package and unchanged this period; any change disclosed.
- [] EV claims sampled against physical evidence, not accepted on assertion.
- [] Accruals recognised for work received but not invoiced; open commitments cleansed.
- [] PV read from the current controlled baseline, phased over the approved schedule.
- [] No retro-fitted budgets; transfers between control accounts logged.
- [] EV/BAC reconciles to the reported percent complete; indices recompute from the four numbers.
- [] Every EAC states its method **and its assumption**; TCPI to BAC reported beside it.
- [] Contingency and management reserve shown separately, with consumption vs progress.
- [] Level-of-effort share disclosed per control account.
- [] AI-produced forecasts marked, calibration record attached, human owner named.

Toolkit 7.T.2 — Estimate basis sheet (one per estimate)

Method (analogous/parametric/bottom-up) · definition maturity and accuracy class · range (–x %/+y %) · rate sources with dates and escalation basis · quantities and their source · exclusions stated explicitly · risks feeding contingency (link to the register) · reviewer and date. An estimate whose basis sheet cannot be produced is not releasable.

Toolkit 7.T.3 — Commercial one-pager (per major package)

Model (FFP/FPIF/CPFF/CPIF/T&M) and why · who carries cost risk · target cost, fee, share ratio, ceiling · **PTA and current cost trend against it** · payment terms both directions and the cash profile · retention and release criteria · variation and claims route · escalation contacts. Reviewed at every commercial checkpoint, not filed at signature.

— Exam preparation — Domain 7

The calculation traps. Percent spent reported as percent complete (Exercise 7.1) · rounding **CPI** before dividing (Exercise 7.2) · **PV** instead of **AC** in the **TCPI** denominator (Exercise 7.3) · applying the wrong party's share to an incentive fee (Exercise 7.4) · **EV** valued at actual cost rather than budget · reading a level-of-effort-heavy **SPI** as schedule performance (MCQ 7.3-D) · quoting **SPI** late in a project where it must converge on 1.00 (7.A.1) · confusing commitments with actuals · treating contingency and management reserve as one pot.

Reflection questions. 1. Your cost report shows one **EAC**. What must accompany it before you will sign it? (The method and the assumption about remaining work; **TCPI** to **BAC** beside it.) 2. On your largest package: who carries cost risk, where is the **PTA**, and how close is the supplier's cost trend to it? (7.4.2–7.4.3; toolkit 7.T.3.) 3. Which invariant in 7.A.3 would have caught the last cost surprise you experienced — and why wasn't it running?

— Domain 7 summary

Cost leadership begins with an honest number: a method matched to definition maturity, a range and class always stated, three-point thinking where tails are real, and a budget built through control accounts into a time-phased baseline with contingency inside it and management reserve outside. Measurement then has to be earned — **AC** with accruals, **EV** at budget under earning rules fixed in advance — because everything downstream is arithmetic on those numbers: **CV**, **SV**, **CPI**, **SPI** at Auriga's week 13 (0.91 and 0.92; 48 % complete against 53 % spent), the **EAC** family spanning USD 408,056 on identical data because they encode different assumptions, and **VAC** and **TCPI** turning recovery talk into a defensible index (1.11 required against 0.91 demonstrated). The commercial half is the same discipline pointed outward: blended rates make options commensurable, contract models allocate cost risk at a price, and the point of total assumption predicts when a supplier's incentives invert. Throughout, the leadership rule holds — the number is only as good as the assumption named beside it, and machine-produced forecasts reach a board only with a calibration record and a human owner. Domain 8 quantifies the risk this domain reserves for; Domain 10 takes the commercial relationships into their own depth.

08

DOMAIN 8

Risk, Uncertainty and Resilience (quantitative)

Group: Delivering the work (Part Two). **Target:** ~74 pages. **Binds to:** the PCI Book Pattern Specification and the shared registries ([docs/books/registries/](#)). This domain is the home of **EMV** and supplies the quantification behind the contingency that Domain 7 (KA 7.1.3) reserves and the schedule ranges Domain 6 (KA 6.4.3) forecasts with. British English; USD (+SAR where useful, indicative **USD 1 ≈ SAR 3.75**).

— Why this domain exists

Domains 6 and 7 built plans and budgets; both left an obligation outstanding. Domain 6 forecast with ranges but did not say where the ranges come from. Domain 7 put a contingency reserve inside the baseline and said it is "sized against identified risks" without showing how. This domain closes both gaps and adds the part neither covers: what a leader does when the unexpected happens anyway. It builds risk identification honestly, including opportunities (KA 8.1); works the analysis from qualitative screening through **EMV**, decision trees and aggregated contingency (KA 8.2); sizes and governs reserves (KA 8.3); and turns to resilience, crisis leadership, cognitive bias and AI-enabled risk sensing (KA 8.4). The discipline throughout is that **risk management is a decision-support activity, not a documentation activity** — a register that changes no decision has cost money and bought nothing.

Learning objectives. After this domain a candidate can: distinguish risk from uncertainty and issue; write a risk statement that supports a decision; identify opportunities as rigorously as threats; run qualitative screening and state its limits; compute **EMV** for individual risks and interpret the total; build and solve a decision tree, including the value of information; aggregate independent risks to a mean and a confidence level rather than summing point estimates; size a contingency reserve defensibly and distinguish it from management reserve; select and cost responses; explain resilience as distinct from prediction; recognise the biases that corrupt estimates and reviews; lead in a crisis; and govern AI-produced risk analysis.

The master worked project. Project Auriga continues from Domains 6 and 7 — the 25-week, **BAC USD 4,000,000** control-systems upgrade. Its risk register, used throughout KA 8.2–8.3, carries the ground-conditions risk that actually materialised in Domain 6's case study, so the reader sees the same event before and after the fact.

KNOWLEDGE AREA 8.1

Threats, opportunities and identification

Topics: 8.1.1 risk, uncertainty and issues · 8.1.2 writing a usable risk statement · 8.1.3 identification methods and their blind spots.

8.1.1 Risk, uncertainty and issues

Definitions. A **risk** is an uncertain event that, if it occurs, affects an objective — threats reduce achievement, **opportunities** improve it. **Uncertainty** in the broader sense covers what cannot be enumerated with probabilities at all — the conditions where scenario thinking and resilience substitute for calculation. An **issue** is a risk that has occurred; it is managed, not analysed. Keeping these apart matters practically: registers clogged with issues stop being forward-looking, and treating deep uncertainty as if it were quantifiable produces false precision (KA 8.4.1).

Risk appetite and thresholds convert judgment into something operable: how much exposure the organisation accepts, expressed as thresholds that trigger escalation or response — connecting this domain to Domain 3's decision rights. Appetite that is never expressed numerically cannot be applied consistently by different people, which is the usual reason two projects in one portfolio treat the same exposure differently.

8.1.2 Writing a usable risk statement

Most registers fail at the sentence level. A usable statement has three parts — **cause** → **event** → **consequence** — and names the affected objective:

Because the controller is single-sourced with a volatile lead time (**cause**), it may be that delivery slips beyond the installation window (**event**), resulting in a delay to commissioning and additional preservation cost (**consequence**).

Compare "supplier risk", which supports no decision: it identifies no cause to attack, no event to monitor and no consequence to size. The three-part form is not a formatting preference — each part maps to a different response type (attack the cause, monitor the event, mitigate the consequence), which is why KA 8.3's response selection depends on it.

8.1.3 Identification methods and their blind spots

Methods: structured workshops, checklists and prompt lists from prior projects, assumption analysis (every assumption is a risk in disguise — Domain 2's assumption register), interviews for what people will not say in a group, and lessons from comparable work (Domain 9's lessons-learned). Each has a blind spot, and knowing them is the professional part:

- **Workshops** surface what the group already shares and suppress what a junior member suspects — so run them with an explicit invitation to dissent and a route for private input.
- **Checklists** find known risks and, by their comfort, discourage looking for novel ones.
- **Assumption analysis** is the highest-yield method and the least used, because it requires admitting how much of the plan is assumed.
- **Historical data** encodes the past's risks, not this project's novelties.

Opportunities are systematically under-identified — most registers are 90 % threats — because the process is framed defensively and because nobody is rewarded for one. A leader who wants them must ask for them separately and size them the same way (Auriga's R5 below is one).

AI in this KA

Generating candidate risk lists is a real AI strength: broad, fast, and free of the group's shared blind spot — genuinely useful against the workshop failure mode above. It is equally a machine for **plausible generic risk**, producing items that fit any project and support no decision. The governed workflow: use AI to widen identification (especially by asking it to challenge the plan's assumptions), then require every surviving item to pass the three-part statement test of 8.1.2 with a named owner. **AI proposes; the professional verifies, decides and remains accountable** — and here the verification is specifically that the risk is this project's, not a template's.

■ KEY TERMS — KA 8.1

Term	Meaning
Risk / opportunity	Uncertain event affecting an objective, adversely / favourably.
Uncertainty	Conditions that cannot be meaningfully enumerated with probabilities.
Issue	A risk that has occurred; managed, not analysed.
Risk appetite / threshold	Accepted exposure, expressed operably enough to apply consistently.
Cause → event → consequence	The three-part risk statement; each part maps to a response type.

■ SAMPLE MCQS — KA 8.1

MCQ 8.1-A [8.1.2 · Analysis] Which register entry best supports a decision? - A. "Supplier risk — high" - B. "Because the controller is single-sourced with volatile lead times, delivery may slip beyond the installation window, delaying commissioning and adding preservation cost" - C. "Delay to commissioning — probability 35 %" - D. "Supplier may cause problems — owner: procurement"

Rationale: Only B names a cause to attack, an event to monitor and a consequence to size (8.1.2). C states a consequence with a probability but no cause, so no response can be designed; A and D support nothing at all.

MCQ 8.1-B [8.1.1 · Recall] The ground contamination has been discovered and remediation is under way. In the register this is: - A. a risk with probability 1.0 - B. an issue — it has occurred, and is managed rather than analysed - C. an opportunity, since remediation was funded - D. removed entirely, with no further record

Rationale: Occurred risks become issues (8.1.1). A is the common fudge that clogs registers; D loses the audit trail and the lesson; C is nonsense.

MCQ 8.1-C [8.1.3 · Analysis] A register of 60 items contains 3 opportunities. The soundest inference is: - A. this project genuinely has few opportunities - B. the identification process is framed defensively; opportunities must be asked for separately and sized the same way - C. opportunities do not belong in a risk register - D. the register is too long and should be cut

Rationale: A 95 % threat ratio reflects process framing rather than reality (8.1.3). C contradicts the definition of risk; D addresses a different problem.

■ SELF-CHECK — KA 8.1

1. State the three parts of a usable risk statement. — Cause, event, consequence, against a named objective.
2. Why is assumption analysis the highest-yield identification method? — Every assumption is a risk in disguise, and plans rest on more of them than teams admit.
3. What is the register consequence of treating an occurred risk as a 100 % risk? — It clogs the forward-looking register; occurred risks are issues.

KNOWLEDGE AREA 8.2

Analysis: from screening to quantification

Topics: 8.2.1 qualitative screening and its limits · 8.2.2 expected monetary value · 8.2.3 decision trees and the value of information · 8.2.4 aggregating to a contingency.

8.2.1 Qualitative screening and its limits

Probability-and-impact scoring on a matrix is a **screening** tool: cheap, fast, and adequate for ranking a long list into an attention order. Its limits are structural and must be stated wherever it is used. Ordinal scales cannot be arithmetic — a "4" is not twice a "2", so scores must not be multiplied and summed as if they were money. Band boundaries create false distinctions (a risk at the top of "medium" outranks one at the bottom of "high" in reality but not on the matrix). And the matrix is blind to **correlation**: five risks driven by one supplier are one risk wearing five costumes.

The professional practice: screen qualitatively, then **quantify what matters** — the items that could breach a threshold, plus anything a decision now depends on.

8.2.2 Expected monetary value

$EMV = \text{probability} \times \text{impact}$ (per risk; threats positive, opportunities negative cost)

WORKED EXAMPLE

Worked example 8.2.2 — Auriga's risk register, quantified.

1. **Setup.** Four threats and one opportunity, each with an assessed probability and cost impact (BAC USD 4,000,000).
2. **Formula.** $EMV = p \times \text{impact}$; total exposure = ΣEMV , opportunities carried negative.
3. **Substitution and result.**

ID	Risk	p	Impact (USD)	EMV (USD)
R1	Controller lead-time slip	0.35	240,000	84,000
R2	Ground conditions worse than surveyed	0.50	180,000	90,000
R3	Integration rework	0.25	320,000	80,000
R4	Permit delay	0.15	400,000	60,000
R5	Early-delivery rebate (opportunity)	0.30	(120,000)	(36,000)
Total expected exposure				278,000

1. **Interpretation.** Total EMV is **USD 278,000** — **6.95 %** of BAC, and USD 314,000 if the opportunity is ignored, which is why opportunities must be in the same arithmetic rather than mentioned in prose. Two cautions the table itself teaches. **EMV is an average of outcomes that will not happen:** no single risk will cost 84,000 — R1 costs 240,000 or nothing. So EMV is the right basis for funding a portfolio of risks and the wrong basis for deciding whether one specific risk is survivable. And **ranking by EMV reorders the register:** R4 has the largest impact (400,000) but the smallest threat EMV (60,000), while R3's lower impact carries more expected cost. A leader managing by impact alone would attend to them in the wrong order — though see 8.3.1, because impact still governs survivability.

8.2.3 Decision trees and the value of information

A decision tree makes a choice under uncertainty explicit: decision nodes (what we control), chance nodes (what we do not), and outcomes valued and rolled back to the present decision.

WORKED EXAMPLE

Worked example 8.2.3 — should Auriga survey before excavating?

1. **Setup.** Excavation may hit contamination: probability **0.40**, cost if hit **USD 300,000**. Option A: proceed directly. Option B: commission a **USD 25,000** ground survey first, which reliably reveals the condition; if contamination is present, a planned mitigation costs **USD 90,000** instead of the reactive 300,000.
2. **Formula.** Roll back each branch: $EV(\text{cost}) = \sum \text{probability} \times \text{outcome cost}$. Value of information = $EV(\text{cost without}) - EV(\text{cost with})$.
3. **Substitution.** A: $0.40 \times 300,000$. B: $25,000 + 0.40 \times 90,000 = 25,000 + 36,000$.
4. **Result.** A **USD 120,000** expected cost; B **USD 61,000**. **Value of the survey = USD 59,000**, against its USD 25,000 price.
5. **Interpretation.** The survey is worth buying, and the arithmetic shows why: it converts a reactive 300,000 into a planned 90,000 in the 40 % of futures where the problem exists. This is the general shape of **buying information** — worth it when it changes what you would do, worth nothing when it does not, however reassuring it feels. Note the sensitivity: if mitigation once forewarned still cost 250,000, B's EV would be $25,000 + 100,000 = 125,000$ and the survey would destroy value. **Information is valuable only in proportion to the better action it enables.** (Domain 6's case study is the counterfactual: Auriga did not survey, met the condition reactively, and paid the recovery this tree prices.)

Survey worth buying: it converts a reactive 300,000 into a planned 90,000 in the 40 % of futures where the problem exists




Fig 8.2.1 — The survey decision tree

8.2.4 Aggregating to a contingency

The error to avoid. Summing worst cases produces a number nobody will fund ($240,000 + 180,000 + 320,000 + 400,000 = 1,140,000$, or 28.5 % of BAC); summing EMVs alone (278,000) funds the average, which by construction is exceeded about half the time. Neither is a contingency. What a leader needs is a **confidence level**: an amount that covers the aggregate outcome with stated probability.

WORKED EXAMPLE**Worked example 8.2.4 — Auriga's contingency at P80.**

1. **Setup.** The five register risks of 8.2.2, treated as independent (an assumption examined below).
2. **Formula.** For independent risks: $\text{mean} = \sum p \times \text{impact}$; $\text{variance} = \sum p(1-p) \times \text{impact}^2$; $\text{standard deviation} = \sqrt{\text{variance}}$. A P80 amount $\approx \text{mean} + 0.8416 \times \sigma$ (the 80th percentile of a normal approximation).
3. **Substitution.** Mean **278,000** (from 8.2.2). Variance $0.35 \times 0.65 \times 240,000^2 + 0.50 \times 0.50 \times 180,000^2 + 0.25 \times 0.75 \times 320,000^2 + 0.15 \times 0.85 \times 400,000^2 + 0.30 \times 0.70 \times 120,000^2 = 63,828,000,000$; $\sigma = 252,642$. P80 = $278,000 + 0.8416 \times 252,642$.
4. **Result.** σ **USD 252,642**; **P80 $\approx$ USD 490,624**. Against the worst-case sum of 1,140,000 and the EMV sum of 278,000 — and against the "10 % of BAC" rule of thumb, which would give 400,000 with no reasoning attached.
5. **Interpretation.** The P80 figure is defensible in a way the other three are not: it states what confidence it buys. Three professional caveats belong beside it. **Independence is an assumption, usually optimistic** — if R1 and R3 share a supplier they are correlated, the variance is larger, and the true P80 is higher; correlation is what turns a bad month into a crisis. **The normal approximation is a convenience** for a handful of Bernoulli risks; proper practice runs a Monte Carlo simulation over the register (and over the schedule, extending Domain 6's three-point durations) to produce a distribution rather than a formula. And **the confidence level is a policy choice**, not a technical one — P80 is common for contingency, P50 for a target, and whoever chooses is making a risk-appetite decision (8.1.1) that belongs to governance.

AI in this KA

Simulation, correlation analysis and scenario generation are legitimate machine work, and the outputs are unusually seductive because a distribution looks like evidence. Three checks before any simulated number reaches a board. **Interrogate the inputs:** a simulation is an amplifier of its assumed distributions and correlations, and those are judgments — garbage in, precisely-shaped garbage out. **Recompute a landmark by hand:** the mean should match $\sum \text{EMV}$; if it does not, the model is not doing what you think. **Ask what would change the decision:** a P80 quoted to the dollar with no sensitivity is decoration. The register's inputs remain human judgments with named owners, whatever produced the output.

■ KEY TERMS — KA 8.2

Term	Meaning
EMV	probability $\times$ impact; the average of outcomes that will not individually occur.
Decision tree	Explicit rollback of decisions and chance events to a present choice.
Value of information	The reduction in expected cost that knowing enables; zero if it changes no action.
P50 / P80	Confidence levels — amounts covering the aggregate outcome with stated probability.
Correlation	Shared drivers that raise aggregate variance; independence is usually optimistic.
Monte Carlo simulation	Repeated sampling to produce an outcome distribution rather than a point.

■ SAMPLE MCQS — KA 8.2

MCQ 8.2-A [8.2.2 · Application] A risk has probability 0.15 and impact USD 400,000. Its **EMV** is: - A. USD 400,000 - B. USD 60,000 - C. USD 340,000 - D. USD 26,667

Rationale: $0.15 \times 400,000 = 60,000$. A is the impact; C is impact less **EMV**; D divides instead of multiplying.

MCQ 8.2-B [8.2.2 · Analysis] R4 has the register's largest impact (400,000) but its smallest threat **EMV** (60,000). The correct managerial reading is: - A. R4 should be ignored, having the lowest **EMV** - B. **EMV** sets the funding priority, while impact still governs whether the event is survivable — both readings are needed - C. impact alone should drive priority - D. the assessment must be wrong, since impact and **EMV** disagree

Rationale: **EMV** is the right basis for funding a portfolio of risks; a single large impact may still be existential regardless of probability (8.2.2, 8.3.1). A and C each discard half the information; D misunderstands that the two measure different things.

MCQ 8.2-C [8.2.3 · Application] Proceeding directly has an expected cost of $0.40 \times 300,000$. A USD 25,000 survey reduces the mitigated cost to USD 90,000. The value of the information is: - A. USD 25,000 - B. USD 59,000 - C. USD 84,000 - D. nil — the survey costs more than it saves

Rationale: $120,000 - (25,000 + 36,000) = 59,000$. A is the survey's price, not its value; C omits the survey cost from branch B; D reverses the conclusion.

MCQ 8.2-D [8.2.4 · Analysis] A register's **EMV** sum is 278,000 and its worst-case sum is 1,140,000. Setting contingency at 278,000 means: - A. an appropriately funded reserve - B. funding the average outcome, which by construction is exceeded roughly half the time - C. a conservative reserve, since not all risks will occur - D. the same as a P80 reserve

Rationale: The mean is the ~50th percentile of the aggregate (8.2.4), so it is exceeded about half the time — the reason a confidence level is chosen explicitly. C mistakes the mean for conservatism; D confuses two different statistics (490,624 here).

MCQ 8.2-E [8.2.1 · Analysis] Why must ordinal probability-impact scores not be multiplied and summed as money? - A. the matrix is only for threats - B. ordinal bands are ranks, not quantities — a "4" is not twice a "2", so the arithmetic is meaningless - C. multiplication requires more than five bands - D. scores may be summed provided they are weighted

Rationale: Ordinal scales support ordering, not arithmetic (8.2.1). Weighting (D) does not repair a scale that never carried magnitude.

■ SELF-CHECK — KA 8.2

1. Why is **EMV** wrong for judging whether one risk is survivable? — It averages outcomes that will not occur; the actual event is impact-or-nothing.
2. When is information worth nothing however reassuring? — When it would not change the action taken.
3. Which assumption in the P80 calculation is usually optimistic, and why does it matter? — Independence; correlated risks raise variance and the true confidence amount.

KNOWLEDGE AREA 8.3

Responses, reserves and governance

Topics: 8.3.1 selecting and costing responses · 8.3.2 reserves and their authority · 8.3.3 monitoring and the register that earns its keep.

8.3.1 Selecting and costing responses

The response families — **avoid** (remove the cause or the exposure), **reduce** (lower probability or impact), **transfer** (insurance, contract terms — a price, not a disappearance; Domain 7, KA 7.4.2), **accept** (with or without a fallback), and for opportunities **exploit, enhance, share, ignore**. Two rules make selection professional rather than reflexive.

A response is an investment. It costs money and time and reduces **EMV**; if the reduction is smaller than the cost, the response destroys value and acceptance is the correct answer. Auriga's fast-track decision from Domain 6 is exactly this arithmetic: one week of client bonus (USD 45,000) against an accepted rework risk of $0.20 \times 60,000 = 12,000$ — a net **+USD 33,000**, which is why the recovery plan took it.

Impact governs survivability even at low probability. A 3 % chance of an event the project cannot survive is not managed by its **EMV**; it is avoided, transferred, or the project's viability is reconsidered. **EMV** funds portfolios; existential risks are handled by structure, and confusing the two is how organisations optimise their way into single points of failure.

Secondary risks. Every response creates its own — fast-tracking creates rework risk, a transfer creates counterparty risk, mitigation creates delivery risk on the mitigation itself. A register whose responses have no secondary entries has not been thought through.

8.3.2 Reserves and their authority

Domain 7 (KA 7.1.3) established the structure; this domain supplies the sizing and the governance:

Reserve	Covers	Sized by	Spent by
Contingency	Identified risks in the register	Aggregation to a stated confidence (8.2.4)	Project manager, under a published protocol
Management reserve	Unknown-unknowns and scope change	Judgment and organisational policy	Sponsor / change authority, via change control

Three governance rules. **The draw protocol is published in advance** — which risk, what evidence, what approval — because a reserve released ad hoc is indistinguishable from an overrun. **Consumption is trended against risk retirement**, not against time: burning 60 % of contingency while 20 % of the register has been retired is the signal MCQ 7.1-B describes. And **contingency released by retired risks is returned, not reallocated** to convenient overspends; otherwise the reserve silently becomes a slush fund and the next real risk is unfunded.

8.3.3 Monitoring and the register that earns its keep

A live register has movement: probabilities revised as evidence arrives, risks retired when their window closes, new entries as the project learns, owners who report rather than merely appear. The tests of whether it is working are behavioural, not documentary — **has a decision changed because of it this month?, are the top items the ones the team actually worries about?**, and

did anything that hurt us appear in it beforehand? A register that answers no to all three is theatre, and the honest response is to fix the process rather than reformat the document. **Early-warning indicators** are the register's operational edge: leading measures tied to specific risks (supplier confirmations slipping, defect discovery rate rising, permit queue lengthening) that fire before the event, which is what makes a response affordable (Domain 6's lead-time point).

■ KEY TERMS — KA 8.3

Term	Meaning
Avoid / reduce / transfer / accept	The threat-response families; transfer is a price, not a disappearance.
Secondary risk	Risk created by a response; a register without them is incomplete.
Draw protocol	The published rules for releasing contingency.
Risk retirement	Closing a risk whose window has passed; frees contingency for return.
Early-warning indicator	A leading measure tied to a risk, firing before the event.

■ SAMPLE MCQS — KA 8.3

MCQ 8.3-A [8.3.1 · Application] A mitigation costs USD 50,000 and reduces a risk's **EMV** from 84,000 to 20,000. The decision and its basis are: - A. reject — 50,000 is a large outlay - B. accept — the **EMV** reduction of 64,000 exceeds the 50,000 cost - C. accept — any reduction in **EMV** justifies a response - D. indifferent — cost and benefit are equal

Rationale: Responses are investments: $84,000 - 20,000 = 64,000$ of reduction for 50,000 of cost is value-creating. A prices the outlay without the benefit; C would justify unlimited spend; D miscomputes.

MCQ 8.3-B [8.3.1 · Analysis] A risk has probability 0.03 and an impact the project could not survive. The correct treatment is: - A. accept it — the **EMV** is small - B. treat it structurally: avoid, transfer, or reconsider viability, because **EMV** funds portfolios while survivability is governed by impact - C. fund its **EMV** in contingency and monitor - D. exclude it from the register as improbable

Rationale: Existential exposure is not an averaging problem (8.3.1). A and C both apply portfolio logic to a single point of failure; D removes the entry that most needs governance attention.

MCQ 8.3-C [8.3.2 · Analysis] Contingency freed by a retired risk is used to cover an unrelated overspend. This is: - A. efficient reserve management - B. a governance failure: the reserve silently becomes a slush fund and the next genuine risk is unfunded - C. acceptable if the total baseline is unchanged - D. required, since contingency is inside the baseline

Rationale: Contingency is tied to identified risks; reallocating it to overspends destroys the link between reserve and exposure (8.3.2). C is the reasoning that makes the failure invisible.

■ SELF-CHECK — KA 8.3

1. When is acceptance the professionally correct response? — When the response's cost exceeds the **EMV** reduction it buys — and the risk is survivable.
2. What does transfer actually achieve? — It prices the risk to a counterparty and creates counterparty risk; it does not remove the exposure.
3. What should contingency consumption be trended against? — Risk retirement, not elapsed time.

KNOWLEDGE AREA 8.4

Resilience, bias and crisis leadership

Topics: 8.4.1 resilience versus prediction · 8.4.2 cognitive bias · 8.4.3 crisis leadership · 8.4.4 AI-enabled risk sensing.

8.4.1 Resilience versus prediction

Quantification (KA 8.2) handles risks you can name. **Resilience** is the capability to absorb what you cannot: it assumes the register is incomplete, because it always is. Its levers are structural rather than analytical — **buffers** (float placed deliberately, contingency, capacity headroom), **optionality** (staged commitments, alternative suppliers qualified in advance, reversible decisions — Domain 1, KA 1.3.3), **modularity** (a failure contained rather than propagated across interfaces), **redundancy** where failure is intolerable, and **fast detection** (early-warning indicators, short feedback loops), because response cost rises with delay.

The trade-off is real and should be stated rather than finessed: resilience costs efficiency. A project optimised to the last dollar and week has no capacity to absorb surprise, and one buffered everywhere is uncompetitive. The leadership judgment is where to be resilient — at the points where failure propagates or cannot be tolerated — not whether.

8.4.2 Cognitive bias

Estimates and reviews are produced by people, and the predictable distortions are worth naming because each has a countermeasure:

Bias	Effect	Countermeasure
Optimism bias / planning fallacy	Systematic underestimation of duration and cost	Reference-class forecasting against comparable completed projects
Anchoring	First number dominates subsequent estimates	Independent estimates before any figure is shared
Availability	Recent or vivid events overweighted	Structured checklists and historical base rates
Groupthink / social pressure	Dissent suppressed in workshops	Pre-mortem, private input channels, explicit dissent invitation
Sunk-cost / escalation of commitment	Continuing because of what is spent	Decision framed on remaining cost and benefit only (Domain 2's kill criteria)
Confirmation bias	Evidence sought that supports the plan	Assign someone to argue the opposite; test the disconfirming case

The **pre-mortem** deserves particular mention as the highest-yield, cheapest technique in this domain: before committing, ask the team to imagine the project has failed and to write down why. It licenses the dissent a status meeting suppresses, and it reliably surfaces risks the workshop of 8.1.3 missed.

8.4.3 Crisis leadership

When a serious risk materialises, the leader's job changes shape. The sequence that works: **stabilise** (stop the harm; secure safety, then the situation), **establish facts** (separate what is known from

assumed — crises run on rumour), **decide with a clock** (a good decision now beats an optimal one later, and say explicitly which decisions are reversible), **communicate early and honestly** (bad news does not improve with age; Domain 11's stakeholder work is now executing under compressed time), **protect the team** (crises consume people, and exhausted teams make the second mistake), and **capture the lesson while it is vivid** (Domain 9's lessons-learned). Domain 6's recovery machinery is the schedule expression of the same posture: re-run the passes, price the options, escalate the trade with named authority.

8.4.4 AI-enabled risk sensing

The genuine capability is **detection at scale**: anomaly detection across cost, schedule, quality and supplier data; pattern recognition against historical projects; and sentiment or volume signals in issue logs and correspondence that precede formal escalation. Used well, this shortens the detection half of 8.4.1's fast-detection lever — which is where its value actually lies.

The honest boundaries: models detect **patterns like the past**, so novel risk is invisible to them; they produce **false positives** that erode attention if unfiltered; and a **calibration record is mandatory** before their alerts influence decisions (Domain 6, KA 6.4; Domain 14). The failure mode to watch for is not a wrong alert but **displaced judgment** — a team that stops thinking about risk because a dashboard is watching. The register's owners remain human, the analysis stays challengeable, and the principle holds: AI proposes; the professional verifies, decides and remains accountable.

■ KEY TERMS — KA 8.4

Term	Meaning
Resilience	Capability to absorb unidentified risk; buffers, optionality, modularity, redundancy, fast detection.
Reference-class forecasting	Estimating from comparable completed projects to counter optimism bias.
Pre-mortem	Imagining failure before committing, to license dissent and surface missed risk.
Escalation of commitment	Continuing because of sunk cost; countered by forward-looking framing.
Displaced judgment	Ceasing to think because a tool is watching.

■ SAMPLE MCQS — KA 8.4

MCQ 8.4-A [8.4.2 · Application] A programme's estimates have run 25 % low across four comparable past projects. The most effective countermeasure is: - A. instruct estimators to be more careful - B. reference-class forecasting — estimate from the distribution of comparable completed projects rather than from the plan - C. add a 25 % contingency and continue the same process - D. escalate to the sponsor for a larger budget

Rationale: Optimism bias is systematic, so the fix is methodological (8.4.2). A relies on exhortation against a structural effect; C treats the symptom while preserving the cause; D funds it without correcting it.

MCQ 8.4-B [8.4.1 · Analysis] A project has been optimised to eliminate all float and buffers. Its risk position is: - A. improved — waste has been removed - B. degraded — it now has no capacity to absorb the risks not in the register, and resilience is what covers those - C. unchanged, provided contingency is funded - D. improved, since efficiency reduces exposure time

Rationale: Resilience assumes register incompleteness (8.4.1); removing all absorption capacity maximises fragility. C confuses money with time and capacity — funded contingency cannot buy back a schedule with nowhere to move.

MCQ 8.4-C [8.4.4 · Analysis] An AI monitor has raised 40 alerts this month, 3 of which mattered. The correct response is: - A. disable the monitor - B. tune thresholds and require a calibration record, because unfiltered false positives erode the attention the tool exists to direct - C. investigate all 40 equally - D. accept the ratio as inherent to anomaly detection

Rationale: False positives consume the scarce resource (attention) that detection is meant to focus (8.4.4). A discards real capability, C guarantees the erosion, D abandons the calibration duty.

MCQ 8.4-D [8.4.3 · Recall] In the first hour of a crisis, the leader's priority order is: - A. communicate, then investigate, then stabilise - B. stabilise and secure safety, establish facts, then decide with a clock and communicate early - C. establish blame, then stabilise - D. wait for complete information before acting

Rationale: Stabilisation and factual grounding precede decisions, which are taken against a clock (8.4.3). D is the failure the sequence exists to prevent; C is never a first-hour activity.

■ SELF-CHECK — KA 8.4

1. Why is resilience not a substitute for quantification, and vice versa? — Quantification handles named risks; resilience absorbs the unnamed. Each fails at the other's job.
2. What is a pre-mortem for? — To license dissent before commitment and surface risks a workshop suppresses.
3. What is AI risk sensing's structural blind spot? — Novel risk: it detects patterns like the past.

— Advanced topics — Domain 8

8.A.1 Correlation and why aggregate risk is worse than it looks

Independence made 8.2.4's arithmetic tractable and understated the answer. When risks share a driver — one supplier, one technology, one regulator, one labour market — they move together, and correlated variance is additive in a way that widens the aggregate distribution's tail sharply. The practical implications: identify **common drivers** explicitly and group risks by them; model correlation in simulation rather than assuming it away; and treat a register with many entries and few drivers as concentrated, not diversified. The organisational version of this error is believing a portfolio of projects is diversified when every one depends on the same scarce engineering resource.

8.A.2 Schedule risk analysis and the merge-bias trap

Applying this domain's methods to Domain 6's network produces **quantitative schedule risk analysis**: three-point durations (KA 6.4.3) simulated across the logic to yield a completion distribution and, valuably, a **criticality index** — how often each activity lands on the critical path across simulations — which is more decision-useful than a single deterministic critical path. The trap it exposes is **merge bias**: where paths converge (Domain 6's node E), the merged event waits for the latest predecessor, so the probability of being on time is the product of the paths' probabilities, not the best of them. Two paths each 80 % likely to meet a date give the merge point 64 %. Deterministic CPM cannot see this, which is why convergence points are systematically optimistic and why Domain 6 flagged them for attention.

8.A.3 The reviewer's risk eye

Invariants worth testing in an hour: every entry has cause, event and consequence and a named owner; no occurred risks parked as high-probability entries; opportunities present in credible proportion; Σ EMV reconciles to the register; contingency stated at a named confidence level with its independence assumption disclosed; contingency consumption trended against risk retirement, not time; responses costed against EMV reduction, with secondary risks recorded; existential risks handled structurally rather than by EMV; early-warning indicators defined for the top items; and at least one decision this month traceable to the register. Failure of the last one matters most: it means the process is documentation, not decision support.

— Industry variations — Domain 8

- **Construction and infrastructure.** Ground conditions, weather and permits dominate; quantitative schedule risk analysis is contractually expected on major programmes, and risk allocation is largely executed through contract terms (Domain 10).
- **Energy and resources.** Commodity price and currency exposure sit alongside delivery risk and are managed with financial instruments as well as project responses (PFL-AI Domains 3, 11).
- **Technology and digital.** Requirement volatility and integration risk dominate; resilience is bought through modularity, staged release and reversibility rather than buffers, so 8.4.1's levers apply in a different mix.
- **Pharmaceutical and regulated development.** Low-probability/high-consequence events and regulatory outcomes make structural treatment (8.3.1) the norm; risk files are auditable artefacts, not internal management tools.
- **Public programmes.** Optimism bias is explicitly recognised in appraisal guidance, and reference-class adjustment may be mandated; political and reputational risk carries weight that no EMV captures well.

— Case study — Domain 8: the survey Auriga did not commission (utilities)

Situation. At planning, Auriga's team considered a USD 25,000 ground survey ahead of the civil works. The register carried R2 — ground conditions worse than surveyed, p 0.50, impact USD 180,000 — and the survey was cut in a cost-reduction pass as "nice to have". The decision was not analysed; it was traded against a round-number savings target.

What happened. At week 13 the contamination was found. The reactive response cost more than the planned one would have — a second civil crew at USD 35,000, a fast-track carrying $0.20 \times 60,000 = 12,000$ of expected rework, and the exposure to a USD 45,000-per-week client penalty window that made the whole recovery urgent (Domain 6's case study; Domain 7's forecast consequence, EAC USD 4.2m against BAC 4.0m).

The analysis, done afterwards. On the decision-tree arithmetic of 8.2.3 — 0.40 probability at the time of the decision, USD 300,000 reactive versus USD 90,000 planned — the survey was worth **USD 59,000** against its USD 25,000 price. Cutting it saved 25,000 and cost the project substantially more, and the register had contained the information needed to see that.

What the domain teaches here. Two things, both uncomfortable. **A cost-reduction pass that does not consult the register is a risk decision taken blind** — savings targets and risk exposure must be traded explicitly. And **the register was not the failure; its irrelevance to the decision was** (8.3.3's test: did a decision change because of it?). The corrective adopted afterwards was procedural: any proposed saving above a threshold must state which register entries it affects and by how much.

— Case study B — Domain 8: the diversified portfolio that wasn't (technology programme)

Situation. A transformation programme's board reviewed 22 project-level risks spread across six workstreams and concluded exposure was well diversified — no single item exceeded 8 % of the aggregate EMV, and contingency was set at the aggregate mean plus a modest margin, assuming independence.

What happened. Fourteen of the 22 risks depended on the same six-person integration team. When two members left within a month, eleven risks materialised or worsened together. The aggregate outcome landed far beyond the contingency, which had been sized on an independence assumption nobody had tested — the correlated variance of 8.A.1 in practice. The programme's recovery required a re-baseline and a management-reserve release.

What was done differently. The register was restructured by **common driver** rather than by workstream, immediately revealing the concentration; the integration capability was treated as a single structural risk with a structural response (cross-training, a retained partner, staged scope) rather than as fourteen separate EMV line items; and contingency was re-derived by simulation with correlation modelled explicitly.

What the domain teaches here. Counting entries measures paperwork; counting **drivers** measures exposure. A register with many risks and few drivers is concentrated, and independence is the assumption most likely to be both convenient and wrong.

— Executive perspective — Domain 8

What a project leader cannot delegate in this domain:

- **The confidence level.** Whether contingency is set at P50 or P80 is a risk-appetite decision belonging to governance, and the leader ensures it is chosen, stated and understood — not inherited from a spreadsheet default.
- **The independence question.** Asking, of any aggregate number, which risks share a driver. It is one sentence and it is where Case study B's programme lost control.
- **The structural exposures.** Existential risks handled by structure, never by EMV — and named personally by the leader, because low probability makes them easy to leave in the register.
- **The register's relevance.** Insisting on the 8.3.3 test: which decision changed because of it this month? A register that changes nothing gets fixed or stopped.
- **The bias countermeasures.** Running the pre-mortem, commissioning independent estimates, framing continuation decisions on remaining cost and benefit — because the leader is the only person positioned to make dissent safe.

- **The AI boundary.** Detection is delegable; judgment is not, and a team that has stopped thinking because a dashboard is watching is the outcome to prevent.

— Calculation exercises — Domain 8

Exercise 8.1 Risks: A p 0.40 impact 150,000; B p 0.20 impact 500,000; C p 0.60 impact 80,000; D (opportunity) p 0.25 impact (200,000). Compute each **EMV** and the total exposure. Solution. A **60,000**; B **100,000**; C **48,000**; D **(50,000)**. Total **USD 158,000**. Common error: omitting the opportunity, giving 208,000 and overstating required contingency by 31.6 %.

Exercise 8.2 Using Exercise 8.1's four risks and assuming independence, compute σ and a P80 contingency (mean + 0.8416 σ). Solution. Variances: $0.40 \times 0.60 \times 150,000^2 = 5.40e9$; $0.20 \times 0.80 \times 500,000^2 = 40.0e9$; $0.60 \times 0.40 \times 80,000^2 = 1.536e9$; $0.25 \times 0.75 \times 200,000^2 = 7.50e9$. Total **54.436e9**; $\sigma = \text{USD } 233,315$; $P80 = 158,000 + 0.8416 \times 233,315 = \text{USD } 354,358$. Common error: taking $\sqrt{(\Sigma\sigma)}$ rather than $\sqrt{(\Sigma\text{variance})}$.

Exercise 8.3 A risk (p 0.30, impact 400,000) can be mitigated for USD 70,000, reducing probability to 0.10. Is the mitigation worthwhile? Solution. **EMV** before $0.30 \times 400,000 = 120,000$; after $0.10 \times 400,000 = 40,000$; reduction **80,000** against a cost of 70,000 $\rightarrow$ **net +USD 10,000, accept** (marginally — and the leader should note that a 400,000 impact may warrant response on survivability grounds regardless). Common error: comparing the mitigation cost against the impact rather than the **EMV** reduction.

Exercise 8.4 Two independent paths each have a 0.80 probability of meeting a milestone date and converge on it. What is the probability the milestone is met, and what is the effect called? Solution. $0.80 \times 0.80 = \mathbf{0.64}$. **Merge bias** — the merged event waits for the later path, so convergence points are systematically optimistic in deterministic CPM (8.A.2). Common error: answering 0.80, the best of the two.

— Practitioner's toolkit — Domain 8

Adoption-ready artefacts; adapt headings to your organisation, then keep them stable.

Toolkit 8.T.1 — Risk register columns that earn their keep

ID · **cause** → **event** → **consequence** statement · affected objective · **common driver** (8.A.1) · probability · impact (cost and time) · **EMV** · qualitative band (screening only) · response family and description · **response cost and expected EMV reduction** · **secondary risks created** · owner (a person) · early-warning indicator · review date · status (open/retired/**issue**) · retirement evidence. Omit any column and the register loses a specific capability; the two most commonly missing — common driver and response cost — are the two that make it decision-useful.

Toolkit 8.T.2 — Contingency derivation sheet (one per baseline)

Register version and date · ΣEMV · variance and σ by risk · **independence assumption stated and challenged** (which risks share drivers) · method (formula approximation or simulation, with correlation treatment) · **confidence level chosen and by whom** · resulting contingency ·

management reserve separately, with its basis · draw protocol reference. A contingency figure without this sheet is a number without a defence.

Toolkit 8.T.3 — Pre-mortem session guide

Before commitment, one hour: (1) state the plan and the commitment about to be made; (2) each participant writes privately for ten minutes — "it is twelve months on and this has failed; why?"; (3) round-robin with no challenge permitted during collection; (4) cluster into causes and identify **common drivers**; (5) convert each into a three-part risk statement with an owner; (6) identify the three cheapest early-warning indicators. Rule: the most senior person speaks last, always.

— Exam preparation — Domain 8

The traps. Multiplying and summing ordinal scores as if they were money (8.2.1) · omitting opportunities from total exposure (Exercise 8.1) · setting contingency at ΣEMV and calling it funded (MCQ 8.2-D) · $\sqrt{\Sigma\sigma}$ instead of $\sqrt{\Sigma\text{variance}}$ (Exercise 8.2) · comparing response cost against impact rather than **EMV** reduction (Exercise 8.3) · managing an existential low-probability risk by its **EMV** (8.3.1) · assuming independence without testing common drivers (8.A.1) · answering the best path probability at a merge point (Exercise 8.4) · leaving occurred risks in the register as $p = 1.0$ · treating a transfer as removing exposure.

Reflection questions. 1. Take your project's contingency figure: what confidence level does it represent, who chose it, and what independence assumption sits underneath? 2. Group your register by common driver rather than by workstream. How concentrated is it actually? 3. Which decision changed last month because of your register — and if none, is the process decision support or documentation?

— Domain 8 summary

Risk management earns its place only when it changes decisions. That starts at the sentence level: cause → event → consequence, against a named objective, with opportunities identified as deliberately as threats because defensive framing otherwise buries them. Analysis then progresses from qualitative screening — useful for ordering, invalid as arithmetic — to quantification: **EMV** for each risk (Auriga's register totalling **USD 278,000**, 6.95 % of **BAC**, and USD 314,000 if the opportunity is ignored), decision trees that price choices and the **value of information** (the survey worth **USD 59,000** against its 25,000 price, and worthless if it changed no action), and aggregation to a **stated confidence** — Auriga's P80 of **USD 490,624** from a mean of 278,000 and σ of 252,642, which is defensible precisely because it says what confidence it buys, unlike a worst-case sum, an **EMV** sum or a 10 % rule of thumb. Responses are investments judged against the **EMV** reduction they buy, with secondary risks recorded and existential exposures handled structurally rather than by averaging; reserves are sized by aggregation, separated by authority, and trended against risk retirement rather than time. Beyond what can be named sits resilience — buffers, optionality, modularity, redundancy and fast detection, bought at an efficiency cost that should be chosen rather than stumbled into — together with the bias countermeasures (pre-mortem, reference-class forecasting, independent estimates) and the crisis sequence of stabilise, establish facts, decide against a clock, communicate early. Independence is the assumption most likely to be convenient and wrong: count drivers, not entries. Domain 9 turns to quality and the assurance that catches what risk analysis missed.

Index

A

AC.....	97
Accountability.....	10
Accrual.....	97
Accuracy class.....	94
Adoption.....	13
Alignment decay.....	25
Analogous / parametric / bottom-up.....	94
Assumption register.....	32
Assurance capture.....	55
Assurance map.....	55
Attribution.....	32
Avoid / reduce / transfer / accept.....	119

B

BAC.....	94
Baseline.....	32
Benefits map.....	32
Benefits owner.....	25
Benefits profile.....	27
Blended rate.....	103
Bounded envelope.....	45
Breakeven adoption.....	27

C

Cash-releasing / non-cash-releasing.....	32
Cause → event → consequence.....	112
Ceiling price.....	103
Commitment.....	97
Committee capacity.....	49
Conditional pass.....	55
Constraint.....	78
Contingency reserve.....	94
Control account.....	94
Control schedule.....	72
Convergence point.....	75
Correlation.....	114
Cost of delay.....	13
CPI SPI.....	98
Crashing / fast-tracking.....	82
Critical path.....	75
CV SV.....	98

D

Dangle.....	72
Date constraint.....	72
Daylight test.....	16
Decidability.....	45
Decision record.....	55
Decision right.....	45
Decision tree.....	114
Delegation schedule.....	49
Displaced judgment.....	121
Do-nothing baseline.....	27
Double-counting.....	32
Draw protocol.....	119
Driver.....	25

E

EAC ETC.....	98
Early-warning indicator.....	119
EMV.....	114
Enabling change.....	32
ES EF LS LF.....	75
Escalation class.....	55
Escalation duty.....	10
Escalation lead time.....	55
Escalation of commitment.....	36

F

Feedback and delay.....	13
FFP / FPIF / CPFF / CPIF / T&M.....	103
Forward / backward pass.....	75
Forward-looking NPV.....	36
Free float FF.....	75
FS / SS / FF / SF.....	72

G

Governance.....	45
Governance artefact.....	45
Governance latency.....	49

H

Hallucination.....	16
Hard / soft constraint.....	25
Honest gate.....	36
Honesty asymmetry.....	10

I

Integration milestone.....	82
----------------------------	----

Irreversibility	8
Issue.....	112
K	
Kill criteria	36
L	
Lag / lead	72
Last responsible moment	13
Level of effort	98
Local optimisation	13
M	
Management reserve	94
Meeting interval (M).....	49
Member vs attendee.....	49
Milestone	78
Monte Carlo simulation	114
N	
Named owner	16
O	
Out-of-cycle mechanism	55
Output / outcome / benefit / value.....	13
Over-trust through fluency	16
P	
P50 / P80	114
Paper lead time (L).....	49
Path migration	82
Planning package.....	78
Point of total assumption (PTA)	103
Portfolio.....	8
Portfolio process.....	25
Pre-mortem.....	121
Precedence rule	45
Programme	8
Project.....	8
PV EV AC	98
R	
Re-baselining	97
Recovery plan.....	82

Reference-class forecasting.....	121
Reserved matter	45
Resilience	121
Resource histogram	78
Responsibility	10
Responsible-AI principle.....	8
Retention	103
Risk / opportunity.....	112
Risk appetite / threshold.....	112
Risk retirement.....	119
Rolling wave	78
S	
Scenario analysis	82
Schedule hierarchy (L1-L4)	72
Secondary risk.....	119
Share ratio.....	103
Single-A test	55
Smoothing / levelling	78
Sponsor.....	49
Stage gate	55
Standard of care.....	10
Steering committee	49
Straw options	27
Sunk cost	36
T	
TCPI	98
Temporary organisation	8
Three lines of assurance.....	55
Three-point / PERT estimate	82
Total float TF	75
U	
Uncertainty	112
V	
VAC.....	98
Value of information.....	114
Value per unit of constraint	27
Verification proportionality	16
W	
Weighted scoring	27