PCI — LinkedIn Launch Pack: 20 Posts (institution voice)

How to use. Post 3–4 a week. These are written in the *Institute's* voice, so they don't depend on any personal profile — the teaching carries the authority. Every post is copy-paste ready. Swap the sector (EPC / oil & gas / rail / data-centre) to suit the week. Keep the honesty: never claim accreditation, recognition or outcomes you don't yet have. British English. All numbers below are worked and correct.

Tone rules baked in: teach something concrete, show the arithmetic, land one takeaway, soft CTA, 3–5 tags.

---

1 — The CPI that lies

Your CPI can look great and still be wrong.

Month-end on a control account: earned value **2,200,000**. Invoiced cost so far: **1,850,000**. A subcontractor has done a further **240,000** of work — not yet invoiced.

Skip the accrual and CPI = 2,200,000 ÷ 1,850,000 ≈ **1.19**. Looks brilliant.

Book the accrual and true cost = 1,850,000 + 240,000 = 2,090,000, so CPI = 2,200,000 ÷ 2,090,000 ≈ **1.05**.

That's a fourteen-point swing on one missed entry — and next month, when the invoice lands, the un-accrued version shows a "sudden overrun" that never happened.

Cut-off isn't bookkeeping hygiene. It's the difference between a performance index that means something and one that whipsaws with invoice timing. Earned value is only as honest as the accruals beneath it.

Free 13-domain course outline in the comments.

#ProjectControls #EarnedValue #CostEngineering #EPC #ProjectManagement

---

2 — EAC is a decision, not a formula

Three teams, same data, three different forecasts at completion. Who's right?

- BAC/CPI (performance persists): one number
- AC + (BAC − EV) (the overrun was a one-off): a lower number
- CPI × SPI (cost and schedule both bite): a higher number

They're all "correct." The formula was never the hard part.

The professional's job is to choose the assumption the numbers rest on — and defend it to a board. Is the variance structural or a blip? Does the remaining scope look like the work done so far?

A forecast you can't defend is a guess with decimals. Pick the method for a reason, and say the reason out loud.

This is what our Body of Knowledge trains: judgement on top of arithmetic.

#EarnedValue #Forecasting #ProjectControls #EAC #ProjectManagement

---

3 — Disruption vs delay (the measured mile)

"We lost money but finished on time." That's disruption, not delay — and most claims for it fail.

Why? Because they assert a lump of lost money with no causal chain. The fix is the **measured mile**: compare your *own* productivity in a clean period against the impacted one.

Cable-pulling, same crews:

- Clean period: 1,200 m in 3,000 hours → 2.5 h/m
- Impacted (piecemeal access): 800 m in 2,600 hours → 3.25 h/m

Excess: 0.75 h/m × 800 m = 600 hours. At 85/hour → **51,000**, priced entirely from your own records.

Cause (evidenced), effect (measured), quantum (from records). No theoretical norms — your own performance is the benchmark, which is exactly why it persuades.

Records make or break this. Code timesheets to areas and periods a year before anyone knows a claim is coming.

#Claims #ContractsManagement #EPC #ProjectControls #Commercial

---

4 — Correlation moves the tail, not the mean

Three risks. Each 30% likely, each 300,000 impact. How much contingency?

Expected value: 3 × 0.30 × 300,000 = **270,000**. That's the same whether the risks are related or not.

But watch the P80 (the level you actually fund):

- **Independent** — outcomes follow a binomial; one hit only reaches 78% cumulative, so P80 = two hits = **600,000**.
- **Perfectly correlated** (one shared driver — say a single supplier) — either nothing happens (70%) or all three land (30%, costing 900,000). P80 = **900,000**.

Same 270,000 mean. Contingency jumps by half.

A register that prices risks one by one and ignores shared drivers under-funds the appetite test. Correlation doesn't change the average — it fattens the tail, and the tail is what breaks projects.

#RiskManagement #ProjectControls #Contingency #MonteCarlo #ProjectManagement

---

5 — The number that moves before CPI does

On labour-driven work, the cost ledger is the last to know. The field knows first.

The **productivity factor**: PF = earned hours ÷ burned hours. It's CPI in labour-hours — and it arrives weekly from timesheets and quantity reports, long before the month closes.

Piping scope, 10,000 budget hours. By week 12: work worth 4,200 earned hours; timesheets show 4,800 burned.

PF = 4,200 ÷ 4,800 = **0.875**. Forecast hours at completion = 10,000 ÷ 0.875 ≈ **11,429** — a ~1,429-hour overrun, roughly **121,000** of labour, visible in week 12 from the field alone.

The ledger won't show it until accruals land. PF already did.

The discipline: hold earned hours to the same rules as EV — no credit for hours merely worked.

#ProjectControls #CostControl #Productivity #EarnedValue #Construction

---

6 — The headroom illusion

The budget screen says you're fine. You're not.

Control account, budget 2,000,000. At the data date: POs raised 1,450,000; invoiced 830,000; a further 190,000 of work done but not invoiced; and 500,000 of scope planned but not yet ordered.

Actual cost = 830,000 + 190,000 = 1,020,000. Open commitment = 1,450,000 − 830,000 = 620,000. Total exposure = 1,640,000. Apparent headroom = 2,000,000 − 1,640,000 = **360,000**. Comfortable, right?

Now add the scope you haven't ordered: forecast final = 1,640,000 + 500,000 = 2,140,000 — a **140,000 overrun**.

The illusion is reading headroom against commitments *made* instead of against the full remaining scope. The budget screen flatters you until the day it doesn't.

#CostControl #ProjectControls #Forecasting #EPC #ProjectManagement

---

7 — Three values of one physical progress

The same wall, built once, produces three different numbers. Confusing them is where projects lose control.

1. **Earned value** — what performance has earned against budget (Domain 6).
2. **Certified value** — what the client's QS has agreed to pay this period (the BoQ valuation).
3. **Recognised revenue** — what accounting can book under IFRS 15.

They rarely match, and they shouldn't be forced to. The professional's skill is to *reconcile* them — to explain the gaps (retention, unbilled work, over-measurement risk), not paper over them.

When EV runs ahead of certified value, that gap is a contract asset — and a question: will the client certify it, or have we over-measured?

One progress, three lenses, three different questions. Master the reconciliation and you master the commercials.

#ProjectControls #EarnedValue #Commercial #IFRS15 #ProjectFinance

---

8 — AI proposes, the professional disposes

We built an AI domain into our credential. Then we built a fence around it.

AI is genuinely useful in project controls: extracting contract clauses, drafting variance narratives, flagging anomalies, first-pass cost coding. Used well, it gives hours back.

But it never decides. Every AI-assisted output passes through a gate: verify against source, and a named professional signs. "AI proposes, the professional disposes" isn't a slogan — it's an operating rule: a threshold, a review step, an audit sample.

The failure mode isn't AI being wrong. It's a human trusting it without checking, then owning a number they never understood.

Governed AI is a competency. Ungoverned AI is a liability. Our Body of Knowledge treats it as the former.

#AI #ProjectControls #ResponsibleAI #ProjectManagement #Governance

---

9 — How do you actually evaluate an AI tool?

"Our AI codes invoices with 97% accuracy." Okay — measured how, on what, and what did it miss?

Evaluating AI in controls needs machinery, not vibes. Two numbers, chosen by the cost of being wrong:

- **Precision** — of what it flagged, how much was right (the cost of false alarms).
- **Recall** — of what was truly there, how much it found (the cost of misses).

A fraud monitor lives on recall (a miss is expensive). An auto-coder lives on precision (a false code pollutes the ledger). Which matters more is an *economic* question, not a technical one.

And you measure both against a **golden set** — a sample whose correct answers professionals established, version-controlled, and never used to train the model it tests.

Building and refreshing golden sets is unglamorous work. It's also exactly what separates governed AI from vibes.

#AI #ProjectControls #DataQuality #ResponsibleAI #Analytics

---

10 — "90% complete" for three months running

If your schedule has said 90% for a quarter, that's not bad luck. It's a measurement failure.

The culprit is usually **cost percent complete** (AC ÷ EAC) standing in for **physical percent complete** (quantities installed, milestones passed, systems turned over).

They diverge on real projects — cost runs ahead of physical when money burns without progress (rework, standing time, escalation). And if you earn value from cost consumed, CPI collapses to 1.0 by construction and the index goes blind.

Earn value from *physical* measures. When the two percentages disagree materially, that gap belongs in the month-end narrative — it's the story, not a nuisance.

The 90% plateau is optimism wearing a hard hat.

#Scheduling #ProjectControls #EarnedValue #ProgressMeasurement #EPC

---

11 — The man-hour rate trap (MDMH)

Man-day / man-hour contracts are everywhere in the Gulf — commissioning, secondment, EPC site services. They're also the easiest place to fund presence instead of progress.

2 senior engineers and 5 technicians, 22 approved days each, at 720 and 360 per day:

- Senior: 2 × 22 × 720 = 31,680
- Technician: 5 × 22 × 360 = 39,600
- Month invoice = **71,280**, payable on verified timesheets.

The invoice check is attendance arithmetic. But the *control* question is different: what did those 154 man-days actually buy?

Under an MDMH schedule you're buying **input, not output** — productivity risk sits entirely with the client. If commissioning progress doesn't move in step with man-days consumed, you're paying for people to be there.

Someone has to own that comparison every single month. On many projects, nobody does.

#Contracts #Commercial #EPC #OilAndGas #ProjectControls

---

12 — The most abused lever in working capital

Every day of DSO is a day of your cash sitting in someone else's account.

The cash-conversion cycle: CCC = DSO + DIO − DPO.

- DSO — days your revenue sits in receivables
- DIO — days cost sits in inventory
- DPO — days you take to pay suppliers

DPO is the tempting lever — it's someone else's cash. It's also the most abused. Extending payment terms by *agreement* is treasury. Extending them by *arrears* is just paying late — and it returns as your project risk when a squeezed supplier fails.

The clean win is DSO. Cut it 5 days on 30,000/day of revenue and you free **150,000** — cash you already earned, just trapped in billing.

Working capital is a controls discipline, not an accounting afterthought.

#ProjectFinance #WorkingCapital #ProjectControls #CostEngineering #Cash

---

13 — Contingency isn't padding

"Trim the contingency to sharpen the bid." That sentence has sunk more projects than bad estimating.

Contingency isn't fat. It's analysed risk, converted to money. When a register's EMV sums to 185,000 and a Monte Carlo run puts the P80 at, say, 415,000, that number isn't negotiable comfort — it's the funding your own analysis says the risk requires.

Cut it to "look competitive" and you haven't removed the risk. You've removed the *funding* for the risk. The risk is still there, now naked.

The professional move is the opposite: set contingency from the quantified register, protect it, and draw it down against actual risk events — with a log, so a reviewer can see where it went.

Appetite is only real when it's arithmetic: a ceiling, a measured exposure, a decision rule.

#RiskManagement #ProjectControls #Contingency #Estimating #ProjectManagement

---

14 — The forecast that only goes up

Six monthly EACs, each a little higher than the last. Respectable accuracy? Look again.

A package's EAC over six months (000s): 9,800 · 9,900 · 10,100 · 10,300 · 10,450 · 10,500 — final out-turn 10,500. Errors: −700, −600, −400, −200, −50, 0. Mean absolute error ≈ 3.7%. Not bad.

But every error is on the *same side*.

Six upward revisions aren't six independent surprises. They're one systematic under-forecast surfacing slowly — the **ratchet**. Accuracy measures how far off; **bias** measures whether you're wrong in one direction. Bias is the more dangerous failure, because it's correctable and usually cultural: optimism, or sandbagging.

Measure your forecasts against out-turn. A function that never checks its own bias repeats it forever.

#Forecasting #ProjectControls #EarnedValue #EAC #ProjectManagement

---

15 — The cost of a slow decision

Governance prices the risk of deciding *wrongly*. It rarely prices the risk of deciding *slowly*.

A programme burns 1,200,000 a month. The investment committee meets every eight weeks. A gate pack misses one cut-off; the decision lands six weeks late.

Cost of that latency: 1.5 months × 1,200,000 = **1,800,000** of burn under a direction the gate might be about to change. In the worst case — a kill decision — that's 1.8 million spent on a programme the organisation had, in substance, already decided to stop.

Decision latency is a governable number. It's set by committee cadence, pack cut-offs and delegation thresholds — all design choices, not facts of nature.

The gate protected you from a bad decision and charged you 1.8 million for a slow one. Put latency in the governance design, not the calendar.

#Governance #ProjectManagement #ProjectControls #PMO #DecisionMaking

---

16 — Time, in a schedule that speaks money

Classical SPI drifts to 1.0 at completion even on a late project — it compares value, and value catches up. That's why it flatters finishes.

Earned schedule fixes it by speaking in *time*. Find the point on the baseline where planned value equals today's earned value; that's your earned schedule, ES.

Baseline PV by month (000s): 200, 500, 900, 1,400, 2,000. At the end of month 4, EV = 1,150.

ES interpolates between months 3 and 4: 3 + (1,150 − 900)/(1,400 − 900) = 3 + 0.5 = **3.5 months**. SPI(t) = 3.5 ÷ 4 = **0.875**. Forecast duration = 10 ÷ 0.875 ≈ **11.4 months** — about six weeks late, on today's tempo.

Classical SPI told you "0.82 and closing." Earned schedule told you the truth: you're half a month behind and it will cost you six weeks.

#EarnedValue #Scheduling #ProjectControls #EarnedSchedule #Forecasting

---

17 — How a fair pass mark is set (building in public)

A credential is only as trustworthy as the way its pass mark is set. So here's how ours *will* be — before a single exam is sat.

Not by picking a round number. Not by curving to a target pass rate. By a **modified-Angoff standard-setting study**: a panel of experienced practitioners judges, question by question, how a "just-competent" candidate would perform, and the pass mark is built up from those judgements. Criterion-referenced — you're measured against the standard, not against other candidates.

We're saying this *now*, in public, before launch, because rigour you can see beats a slick claim you can't check.

We describe every standard we reference (IFRS, ISO principles, agile frameworks) in our own words, reproduce none of them, and imply no accreditation we don't hold. Honest is the whole brand.

#Certification #ProjectControls #Assessment #Psychometrics #ProfessionalDevelopment

---

18 — Why we wrote our own Body of Knowledge

You can't build a serious credential on borrowed material. So we wrote a complete one — in our own words.

Thirteen domains. Sixty-one knowledge areas. Weighted 40% project accounting & finance, 40% project management principles, 20% the governed use of AI. Hundreds of worked examples in a five-step format — setup, formula, substitution, result, interpretation — because a controls professional is made by working the arithmetic, not reading around it.

Two sector case studies per domain. Calculation exercises with full solutions. An examination-style question bank with rationales for every option.

It's a first edition, under continuous review, pending subject-matter-expert sign-off before it feeds a live exam — and we say so plainly.

The free course outline is in the comments. Judge the substance for yourself.

#ProjectControls #BodyOfKnowledge #CostEngineering #ProjectManagement #PCPAI

---

19 — Free study resources (lead magnet)

We've published a set of free candidate resources — no payment, no catch. Take what's useful.

- **Course Outline & Study Guide** — all 13 domains and 61 knowledge areas with their topics
- **Master Formula Sheet** — every formula, grouped, with where it's taught

- **Glossary** — the field's terms, defined once
- **Sample Questions** — scenario-based items with full rationales

Whether or not you ever sit the exam, these are solid revision material for any cost engineer, planner or controls professional.

Download links in the comments. If they help, pass them to a colleague.

#ProjectControls #FreeResources #CostEngineering #Scheduling #ProfessionalDevelopment

---

20 — A place for the discipline to talk

Project controls is a small world that rarely gets a room of its own. So we built one.

Our community forum is open to everyone — practitioners, candidates, the curious. No account needed to read; pick a display name to post. Ask how others baseline EVM on hybrid projects, how they handle a disputed prolongation claim, how they're actually using AI without letting it decide.

No sales, no gatekeeping. Just the conversation the discipline deserves — and a standard being built in the open, with the people who'll use it.

Come and start a thread. Link in the comments.

#ProjectControls #Community #ProjectManagement #EPC #EarnedValue

---

Posting notes

- **First line is everything** on LinkedIn — the hook is what shows before "...see more". Each post above leads with one.
- **Put links in the first comment**, not the post body (LinkedIn suppresses reach on posts with outbound links).
- **Rotate the three pillars**: teach (1–7, 10–16), point-of-view/AI (8, 9), build-in-public/credibility (17, 18), community/resources (19, 20). Roughly 60% teach, 20% POV, 20% credibility/community.
- **Localise freely**: swap the worked-example sector to oil & gas, rail, data-centre or infrastructure to match the week's audience — the arithmetic stays the same.
- **Stay honest**: never add "accredited", "recognised", "guaranteed", pass rates or salary claims. The honesty is the differentiator — don't trade it for a stronger-sounding line.