PCB-05 · BODY OF KNOWLEDGE — EXECUTIVE SUMMARY

How the Body of Knowledge was built

The authoring order, the evidence rule, the arithmetic discipline and the gates a chapter must survive

VERSION 1.0 · DRAFT
2026-08-04

PROJECT CONTROLS INSTITUTE GLOBAL, INC.

AI proposes. The professional disposes.

Summary

A body of knowledge is only worth as much as its method. This document sets out how the PCL-AI Body of Knowledge was authored: the style spine that fixed notation before content, the sequence in which the thirteen domains were written and why, the per-knowledge-area completion checklist, the four-way classification of every claim, the arithmetic verification rule, and the review gates a chapter must pass — including what is still unfinished.

About this document

Document	PCB-05
Series	Body of Knowledge — Executive Summary
Version	1.0
Status	draft
Date	2026-08-04
Unresolved placeholders	2

Contents

How the Body of Knowledge was built

In one paragraph. A body of knowledge is only worth as much as its method. This document sets out how the PCL-AI Body of Knowledge was authored: the style spine that fixed notation before content, the sequence in which the thirteen domains were written and why, the per-knowledge-area completion checklist, the four-way classification of every claim, the arithmetic verification rule, and the review gates a chapter must pass — including what is still unfinished.

Who this is for. Practitioners deciding how seriously to take the framework, subject-matter experts considering reviewing it, and anyone building a body of knowledge of their own who would rather not rediscover the failure modes.

---

1. The problem a body of knowledge has to solve

Authority is the whole problem. A framework that decides whether someone passes an examination cannot borrow its authority from the confidence of its prose. It has to be built so that a hostile expert can find the joins: where a claim is a standard's requirement, where it is the Institute's recommendation, where competent practitioners legitimately differ, and where a number came from.

A second problem sits underneath. A reference this size cannot be written in one pass by one person, and material written in parallel drifts — the same symbol means two things by chapter nine, the same concept is defined three times in slightly different words. A candidate who meets `PV` as Planned Value in one domain and present value in another has been set up to fail an item testing neither. Both problems are solved before writing, not after.

2. The style spine — fixing the conventions before the content

The first document authored was not a domain. It was the style spine, which fixes the conventions and binds every domain to them.

Terminology. A term is defined once and used identically everywhere; silent redefinition in a later chapter is a defect, not an inconsistency. Terms a domain introduces join the shared glossary in a consolidation pass rather than living locally.

The symbol table. One table of symbols with units, seeded up front even for notation first used deeply in later domains, so the classical and agile treatments of earned value share one notation. Genuine clashes are resolved by explicit rule: `PV` denotes Planned Value in earned-value contexts and present value in discounting contexts, and discounting contexts must write "present value" in words or `PV(x)` so the two never collide on a page.

The worked-example format. Five lines, every time: setup with the given data; the formula with each variable named and its unit; the substitution shown explicitly; the result with unit and rounding; and an interpretation of what the number means for the decision. Where a computation is inherently tabular — a network forward and backward pass, a multi-scenario comparison — a table replaces the first four lines. The interpretation is never optional, because a number a reader cannot act on has not been taught.

Figure specifications. Every figure carries a number on the Domain.KA pattern, a caption, the underlying data so it is reproducible, and a render-ready description. A figure whose data is not stated cannot be checked.

The item format. Four options, one key, and a rationale saying why the key is right and why each distractor is wrong, tagged with topic number and cognitive level. Distractors are the results of common errors, never filler.

3. The authoring order, and why it is not one to thirteen

The domains were written in a sequence dictated by dependency, not numbering. The four finance domains came first, locking the notation and worked-example style on material where precision is least forgiving; the project management domains inherited it. The agile domain was authored **after** earned value, because agile cost forecasting reuses the earned-value machinery on variable scope and would otherwise have invented a parallel notation for the same arithmetic. The AI domain was authored **last**, so its workflows could cross-reference the domains they operate on rather than describing AI in the abstract — which is why Domain 13's largest knowledge area is a systematic pass across the earlier domains rather than a survey of tools.

4. The evidence rule — four kinds of statement, each marked

Every claim is one of four things, and the reader is told which. A **fact** is verifiable and, where it matters, attributed to a named framework. A **recommended practice** is what the Institute advises, stated as advice. A **professional judgement** is a point where competent practitioners legitimately differ — said so explicitly, with the grounds for choosing given rather than withheld. A **PCI interpretation** is the Institute's position, labelled as the Institute's.

Two absolute rules govern sources. **Name real frameworks; never fabricate.** IFRS 15, IAS 1, IAS 2, IAS 16, IAS 23, IAS 37 and IFRS 16; the PMBOK Guide and AACE's Total Cost Management framework; ISO 31000 and ISO/IEC 17024; the Agile Manifesto, the Scrum Guide, Kanban and Lean are named where relevant, and no clause number, page reference or quotation is invented. **Never reproduce protected text.** A standard's principle is explained in the framework's own words — "under IAS 37 a provision is recognised when..." — and its wording, tables, diagrams and question banks are never pasted. Every example, table, figure and item is original.

A third rule governs the AI material: describe real capabilities and their limits — hallucination, data quality, bias, confidentiality, auditability — mark evolving capabilities as evolving, and never hype. What makes all of it enforceable is the blanket prohibition on fabrication: no invented quotations, case studies, company names, project data, survey results or statistics. A plausible-looking citation that was not consulted is a conduct matter, not a style error.

5. Arithmetic, jurisdiction and language

Every calculation is independently recomputed by someone other than its author, and numerical items get two reviewers. Units, currency, period and rounding are stated; the substitution is shown, not only the result; and where an answer depends on an assumption, the assumption is named in the same breath as the answer. Rounding is fixed by convention — money to the whole currency unit unless the topic needs cents, ratios to two decimal places, percentages to one. A chapter containing an unverified calculation is rejected at gate.

No jurisdiction's legal, tax or accounting treatment is presented as universal. Where treatment varies, the framework says so and teaches the principle; where an example needs a legal frame, the frame is stated as an assumption of the example. **British English throughout**, acronyms expanded on first use.

6. The per-knowledge-area completion checklist

A knowledge area is not finished when it is long. It is finished when every applicable line of a published checklist is satisfied: a precise definition and purpose for each topic with the real standard named; the underlying principle; every formula stated with each variable and unit defined; at least one fully worked example in the five-line format with the numbers re-checked; a second example or mini-case for any non-trivial topic; at least one numbered figure specification; common pitfalls; a treatment of how AI assists that area and where it must not be trusted; a key-terms box; three to six sample items with rationales and level tags; two to four self-check questions with answers; and cross-references by number.

That is what makes depth auditable. "Does this knowledge area have a worked example whose numbers add up?" is a question with an answer; "is this chapter deep enough?" is not.

7. The gates, and what fails at them

Every chapter passes a fixed sequence of reviews before it is treated as content: pattern conformance, competency alignment, technical accuracy, calculation verification, source verification, originality, style, legal and company-mention review, diagram and table review, accessibility, cross-reference integrity, item validation, and final editorial approval. A reviewer never reviews their own draft; one chapter has one owner and one file.

Four things fail at gate regardless of quality elsewhere: a placeholder citation, an unresolved calculation, duplicated prose, and generic filler. Length is explicitly not an acceptance criterion — page count is an output check, never a writing method, and no phase may be skipped to meet one.

8. What is not yet done

Three things are outstanding, and the framework is published saying so.

Subject-matter-expert review. Every domain is a **first authored draft**: it has passed the internal gates in §7 but not independent subject-matter-expert review, and until it has, it is not final certification content. Panel: [CONFIRM: the number of subject-matter experts on the first-edition review panel and their disciplines]. Completion: [CONFIRM: date subject-matter-expert review of the first edition completes].

The job-task analysis. The blueprint's sampling has not been validated against practitioners' actual work. That study will confirm item counts — which is why no item count is published anywhere — and will test the group weightings defended in `PCB-03 — Why 40/40/20`.

Standard setting. The pass mark configured on the platform is sixty-five per cent; the definitive standard will be set by a modified-Angoff study, a documented judgement about what a minimally competent professional must demonstrate.

Publishing a method with its gaps visible invites the objection that the framework is unfinished, which it is. The alternative — a confident framework whose gaps surface after the first cohort has paid — is the reason certification bodies are distrusted. A reader who can see which parts are settled can challenge the parts that are not, and that challenge is worth more than the appearance of completeness.

Related

- `PCB-01 — The Project Controls Body of Knowledge — executive summary` — what the material built this way claims
- `PCB-02 — The thirteen domains at a glance` — the domain and knowledge-area structure this method produced
- `PCB-04 — What a project controls professional must know` — the propositions the framework was built to support
- `EXB-04 — How items are written, reviewed and retired` — the item-development standard that continues from §2
- `EXB-06 — Job-task analysis: where the blueprint gets its authority` — the outstanding study named in §8

Sources and standards

- The Style Spine — Conventions of This Reference (`docs/bok/00-style-spine.md`), August 2026: §1 scope and weighting, §3 the seed glossary, §4 the master symbol table and notation-clash rule, §5 the worked-example format, §6 figure specifications, §7 language and rounding conventions, §8 the item format, §9 citation rules, §10 the per-knowledge-area checklist.
- PCL-AI Body of Knowledge, first authored draft and its README (`docs/bok/`), August 2026: the authoring order and the per-domain draft status.
- PCI Editorial Charter (`docs/books/EDITORIAL_CHARTER.md`), August 2026: §3 the non-negotiable quality rules, §6 concurrent-authoring and file-ownership rules, §7 the chapter quality gates, §8 phase gates, §10 the definition of done.
- Examination Blueprint (`docs/downloads/examination-blueprint.md`), August 2026: §5, the job-task analysis and modified-Angoff commitments.

The standards named in §4 are named as frameworks the Body of Knowledge references and explains in its own words. No edition, clause or wording is cited here, because none was verified for this document.

Status and version

Founding-stage document · Version 1.0 · draft · Reviewed under PCI governance. PCI makes no claims of accreditation or recognition beyond what is true today.