

CMP-03 · COMPETENCY FRAMEWORKS

PCL-AI: the fourteen competencies

Each competency of the PCI AI Project Controls Leader —
defined, levelled, evidenced, and separated from its look-alike

VERSION 1.0 · DRAFT
2026-08-04

PROJECT CONTROLS INSTITUTE GLOBAL, INC.
AI proposes. The professional disposes.

— Summary

The PCI AI Project Controls Leader credential publishes fourteen competencies. This document takes each one in turn and states what it is in terms of decisions, what it looks like at each of the four levels, the named artefacts that evidence it, and the common false positive — the thing that is routinely mistaken for the competency and is not it.

— About this document

Document	CMP-03
Series	Competency Frameworks
Version	1.0
Status	draft
Date	2026-08-04
Unresolved placeholders	1

Notice. This document is an educational publication of Project Controls Institute Global, Inc., and does not constitute accounting, tax, legal, financial or other professional advice. It must not be relied upon as such. Where a topic depends on the contract or the governing law, entitlement and treatment vary by jurisdiction; take local advice.

Standards. Standards and frameworks are referred to by name and described in this publication's own words. No standard's text, tables, diagrams or question banks are reproduced. All trademarks remain the property of their respective owners, and no endorsement by any standards body, employer or vendor is implied or should be inferred.

Status. Project Controls Institute Global, Inc. is a Delaware Non-Stock Corporation. The Institute has been developed with reference to ISO/IEC 17024 personnel-certification principles and is **not accredited** by ANAB, IAS or any other accreditation body. It intends to seek 501(c)(3) recognition, which has not been granted. No governmental approval, academic equivalence, or guaranteed employment, promotion or salary outcome is claimed or implied.

Examples. All worked examples are illustrative and fictional. Figures are chosen to demonstrate a method, not to describe any real project, organisation or market. Sample examination items, where they appear, are study material maintained separately from any live examination bank.

Completeness. Passages marked **[CONFIRM: ...]** are operational or market facts that have not yet been confirmed from a cited source. A document carrying them is a working draft and is not approved for publication.

© 2026 Project Controls Institute Global, Inc. · All rights reserved.

Contents

1. How to read this document.....	4
2. Governance, cost, schedule and value	4
2.1 Project controls governance.....	4
2.2 Cost management	5
2.3 Planning and scheduling	5
2.4 Earned value management.....	6
2.5 Forecasting and EAC.....	6
2.6 Performance measurement.....	7
2.7 Project risk.....	7
2.8 Commercial and contract controls	8
3. The data, digital and AI group	8
3.1 Predictive analytics.....	8
3.2 AI-enabled project controls	9
3.3 Digital reporting.....	9
3.4 Automation.....	10
3.5 Responsible AI	10
3.6 Human validation	11
4. Where the boundaries actually lie	11
5. How this goes wrong	12
6. Checklist	12
Related	13
Sources and standards.....	13
Status and version	13

PCL-AI: the fourteen competencies

from its look-alike.

In one paragraph. The PCI AI Project Controls Leader credential publishes fourteen competencies. This document takes each one in turn and states what it is in terms of decisions, what it looks like at each of the four levels, the named artefacts that evidence it, and the common false positive — the thing that is routinely mistaken for the competency and is not it. It is the working reference for anyone assembling evidence, writing a role specification, or assessing a claim against the PCL-AI standard.

Who this is for. Candidates for PCL-AI; project controls managers and heads of function designing roles and development plans; assessors applying the standard to evidence they did not collect.

— 1. How to read this document

The fourteen competencies below are the published PCL-AI set, in the published order. They are not reordered, renamed, merged or extended here. Where two of them appear to overlap — and several do, by design — the boundary is stated explicitly, because the boundary is where assessment decisions are actually made.

Each entry has four parts, per CMP-01 §3: **definition** (the decisions the competency governs), **by level** (foundation, practitioner, professional, leader — the scale is defined in CMP-02 §2 and is not restated here), **evidence** (named artefacts), and **the false positive** (what looks like it and is not).

Three clusters are treated in greater depth elsewhere and are cross-referenced rather than repeated: cost and commercial in CMP-06, planning, scheduling and risk in CMP-07, and the data, digital and AI group in CMP-08.

One boundary of scope is worth stating up front. The published examination blueprint for PCL-AI carries the Body of Knowledge domain weighting; it does not publish a mapping from these fourteen competencies to those domains [\[CONFIRM: published competency-to-domain mapping for the PCL-AI examination blueprint\]](#). The competency set states what the credential claims; the blueprint states how the examination samples knowledge. They are complementary and neither is derived from the other.

— 2. Governance, cost, schedule and value

2.1 Project controls governance

Definition. Deciding how the controls function itself operates: what is baselined and by whom, which changes require which authority, the calendar on which data is cut and reported, the segregation between the person who forecasts and the person who approves, and what happens when a

rule is broken. Governance is the competency that makes the other thirteen mean the same thing twice.

By level. Foundation — follows the controls procedure and knows which approvals a change requires. Practitioner — operates the governance calendar for a scope: runs the cut-off, maintains the change log, prepares change papers to the required authority. Professional — designs the governance for a project: sets thresholds, defines the delegation of authority in controls terms, and refuses to process a change that has not been authorised, including when the requester is senior. Leader — owns the controls governance framework across a portfolio, decides where projects may deviate from it, and rebuilds it after a failure rather than reissuing the same document with a stern covering note.

Evidence. A project controls execution plan you authored; a change authority matrix in force with your name against it; a cut-off calendar with adherence recorded; a documented refusal or escalation of an unauthorised change; a governance review finding you raised and the change that followed.

The false positive. Owning the procedure document. A well-written procedure that nobody follows is evidence of authorship, not governance. Ask what has been stopped by the governance in the last six months; if the answer is nothing, the governance is decoration.

2.2 Cost management

Definition. Controlling the flow from estimate to budget to commitment to accrual to actual cost to forecast, keeping each step reconcilable to the last, and knowing at any date what has been spent, what has been committed but not spent, and what has been incurred but not yet invoiced. Depth in CMP-06.

By level. Foundation — codes and processes cost transactions against the cost breakdown structure, and queries a misallocation. Practitioner — owns control accounts: maintains commitments, raises and substantiates accruals, reconciles to the ledger, explains a variance in terms of its driver rather than its size. Professional — sets the cost breakdown structure and the code of accounts for a project, decides accrual policy at the boundary cases, and defends a cost position to a finance function that computes it differently. Leader — governs the cost control environment across projects: the structures, the cut-off discipline, the interface with corporate finance, and the standard of substantiation required.

Evidence. A cost breakdown structure and code of accounts you designed; a signed month-end accrual pack with your substantiation; a documented reconciliation between the project cost report and the general ledger, including the explained differences; a variance analysis that identified a driver later confirmed.

The false positive. Cost reporting. Producing an accurate report of what a system contains is administration. Cost management is the competency that determines whether what the system contains is right — that the accrual reflects work done rather than invoices received, that a commitment is a commitment and not an intention.

2.3 Planning and scheduling

Definition. Building and owning a schedule whose logic represents how the work will actually be done, maintaining it as reality diverges from it, and being able to state what is driving the completion date and what would change it. Depth in CMP-07.

By level. Foundation — updates progress against an existing schedule and flags logic that has become impossible. Practitioner — owns a discipline or area schedule: builds the network, applies calendars and resources, runs the update cycle, and explains a movement in the critical path in terms of the activities that caused it. Professional — decides the schedule architecture for a project — level structure, coding, calendar strategy, how contractor schedules integrate — and defends a completion forecast against a delivery team that wants an earlier date. Leader — sets scheduling standards across projects, decides where constraint-driven planning is legitimate and where it is a disguised instruction, and owns schedule credibility with the client or board.

Evidence. A schedule you built, with its basis-of-schedule document; a documented critical path analysis before and after an event; a schedule quality review you performed on someone else's programme with the findings and the resolution; a completion forecast you held that was later borne out, with the contemporaneous note.

The false positive. Tool fluency. Rapid competence in scheduling software with no view on whether the logic is credible is data entry at speed. The test is not what the software shows but whether the person can say why the current critical path is the critical path — and whether they have ever concluded that it is not.

2.4 Earned value management

Definition. Operating the measurement system that relates work performed to work planned and to cost incurred: the scope-to-budget structure, the rules of credit that decide how progress is claimed, the integrity of the resulting earned value, and the variance analysis that follows from it. Expanded in CMP-06 §3.

By level. Foundation — applies rules of credit correctly and calculates the standard indices. Practitioner — owns the performance measurement baseline for a scope, applies and defends rules of credit at the awkward boundaries (procured but not installed, fabricated off site, engineering substantially complete), and produces variance analysis with causes attached. Professional — designs the earned value system for a project: work package sizing, technique selection per work type, the treatment of level-of-effort scope, and the rules that prevent progress being claimed for work that has not happened. Detects and corrects earned value that has drifted from physical reality. Leader — sets the standard for earned value across projects and decides where it should not be used at all, because the scope will not support honest measurement.

Evidence. A rules-of-credit sheet you wrote and applied; a performance measurement baseline with your work-package structure; a documented instance where you rejected a progress claim and the reason; variance analyses with root causes that were subsequently confirmed by events.

The false positive. Producing indices. A cost performance index calculated to three decimal places from a system nobody has validated is arithmetic. Earned value competence is the ability to say whether the earned value is real — which usually means having walked the work, or having found the cases where the claim and the physical position disagreed.

2.5 Forecasting and EAC

Definition. Producing a forward view of final cost and date — the estimate at completion (EAC) and the forecast completion — choosing the method deliberately, stating the assumptions the answer depends on, and defending the result when it is unwelcome. Depth in CMP-06 §4.

By level. Foundation — calculates an estimate at completion using the prescribed formula and identifies where the inputs look wrong. Practitioner — produces the periodic forecast for a scope,

selecting from established methods according to procedure, and explains the movement since last period. Professional — decides which method fits the situation where procedure is silent — performance-index based, remaining-work re-estimate, or a hybrid — states the range and its drivers, and holds the forecast when it is challenged, changing it only on evidence. Leader — governs forecasting across a portfolio: the method standard, the review cadence, the treatment of undisclosed exposure, and the escalation route when a forecast is being suppressed.

Evidence. Dated forecast notes recording method, alternatives rejected and key assumptions; a forecast you revised upward before external events forced it, with the record; a back-test of your own forecasts against outturn; a range or scenario forecast with the drivers of the spread named.

The false positive. A single number produced by the tool's default formula. The competency lives in the method selection and the assumption statement, not in the arithmetic. A person who cannot say what would have to be true for their forecast to be right does not hold this competency at professional level, however sophisticated the model behind it.

2.6 Performance measurement

Definition. Deciding what to measure so that the measures describe the project rather than flatter it: physical progress against claimed progress, productivity, quality and rework, interface and dependency health, and the leading indicators that move before the lagging ones. Broader than earned value, which is one measurement system among several.

By level. Foundation — collects measurement data accurately and to the defined cadence. Practitioner — operates the measurement set for a scope, validates data at source, and notices when a measure has stopped tracking reality. Professional — designs the measurement set for a project: chooses the measures, defines each one operationally (what counts, who counts it, at what frequency, with what tolerance), and removes measures that have become performative. Leader — sets what the organisation measures, and accepts the consequence that people optimise what is measured; designs measure sets that survive that pressure.

Evidence. A measurement definition sheet you authored, with operational definitions per measure; an instance where you retired or replaced a measure and why; a documented case where a leading indicator you introduced anticipated an outcome; a data-quality audit at source with the corrections made.

The false positive. Dashboard richness. Twenty-four indicators, six of which are actionable and none of which is defined operationally, is measurement theatre. The competency test is whether the person can name a decision that a measure changed.

2.7 Project risk

Definition. Identifying what could change the outcome, quantifying it where quantification is meaningful, assigning ownership that is real, and connecting the analysis to the contingency and the decisions it should inform. Depth in CMP-07 §6.

By level. Foundation — records risks accurately, maintains the register, chases owners for updates. Practitioner — facilitates risk identification for a scope, assesses probability and impact with a defined scale, tracks mitigation actions to completion and closes risks that have expired. Professional — decides the analysis approach — qualitative, quantitative cost risk analysis, quantitative schedule risk analysis (QSRA) — designs the modelling, interprets the output honestly including its sensitivity to correlation assumptions, and sets contingency on that basis. Challenges risk scores that are political rather than analytical. Leader — governs risk across projects, owns the relation-

ship between risk analysis and contingency policy, and decides when a risk position must be escalated regardless of who is uncomfortable.

Evidence. A quantitative risk analysis you ran, with the input basis, correlation assumptions and the contingency recommendation that followed; a register entry you closed as expired with reasoning; a risk you escalated that materialised; a documented challenge to an impact assessment you judged understated.

The false positive. Register maintenance. A tidy register with current dates and complete fields demonstrates administration. Risk competence shows up in whether the register contains the risks that actually matter — and specifically whether it contains the one nobody wants written down.

2.8 Commercial and contract controls

Definition. Controlling the commercial position: what has been instructed, what has been valued, what has been agreed, what is claimed and unresolved, and what exposure sits outside the current forecast. Includes change and variation control, the substantiation of entitlement, and the interface with the contract's notice and time-bar mechanics. Depth in CMP-06 §5.

By level. Foundation — logs instructions and variations, and knows which contract clause governs a notice. Practitioner — owns the change and variation register for a scope, values changes to procedure, maintains the record required to substantiate an entitlement, and raises notices within their period. Professional — decides the commercial position on contested change: what to price, what to reserve, what to disclose as exposure rather than book as forecast; assembles the factual record supporting an extension of time (EOT) or additional cost claim; and states the commercial risk to the project board in terms they can act on. Leader — sets the standard for commercial control across projects, including the boundary between the controls function and the commercial or legal function, and owns the organisation's exposure reporting.

Evidence. A change control log you owned, with the audit trail from instruction to agreed value; a contemporaneous records regime you set up and the claim it later supported; an exposure report distinguishing forecast from unquantified exposure; a notice issued on time under a time-bar provision.

The false positive. Contract knowledge. Being able to cite the clause is necessary and not sufficient. The competency is in the discipline of the record — whether the daily allocation sheets, the delay records and the correspondence exist and are organised well enough that entitlement can be demonstrated eighteen months later by someone who was not there.

— 3. The data, digital and AI group

These six competencies form a chain, not a set of alternatives: data and analytics produce a proposal, the controls design decides where that proposal enters the work, reporting and automation carry it, responsible AI governs it, and human validation makes a person answerable for it. CMP-08 treats the chain in depth; the entries here define each link and its boundary.

3.1 Predictive analytics

Definition. Using quantitative technique on project data to say something about what has not happened yet — trend extrapolation, regression against drivers, benchmark comparison, statistical modelling of outcome ranges — and knowing the conditions under which each technique is valid.

By level. Foundation — reads and correctly interprets an analytical output produced by someone else, including its stated uncertainty. Practitioner — applies established analytical techniques to project data, prepares and cleans the data, and reports the result with its confidence and its limitations. Professional — selects the technique for the question, tests the assumptions that make it valid, holds out data to back-test, and reports the error characteristics of their own model rather than only its central answer. Leader — decides what the organisation will and will not predict analytically, sets the data standard that makes prediction possible, and owns the consequence when a model is used outside its range.

Evidence. An analysis with its data basis, method rationale and validation results; a back-test of your own predictions against actuals; a documented decision that the data would not support the analysis requested, and what you did instead.

The false positive. Curve fitting to date. A trend line through eight months of cost data with a high coefficient of determination is a description of the past, not a prediction, and the coefficient says nothing about whether the relationship will hold. The competency test is whether the person has ever reported that their own model was wrong, and by how much.

3.2 AI-enabled project controls

Definition. Designing how the controls function actually uses artificial intelligence: which tasks are delegated to a model, what the model receives, what it returns, where its output enters the process, what check stands between the output and reliance, and what happens when the model is unavailable or wrong.

By level. Foundation — uses approved AI-assisted tools within the defined process and knows which outputs must be checked before use. Practitioner — operates an AI-assisted workflow end to end, prepares inputs to the required standard, and recognises output that is implausible for the project. Professional — designs the workflow: selects the task, specifies the input and output contract, defines the check, and decides which tasks are not delegated at all. Can state the failure mode of each delegated task. Leader — sets the controls function's AI operating model, evaluates and admits tools, decides where the function will not use AI, and owns the outcome of those decisions.

Evidence. A process map showing AI-assisted steps with their input, output and check points; a tool evaluation you conducted with the criteria and the decision; a task you removed from an AI-assisted workflow, and why; an incident review after an AI-derived output reached a report.

The false positive. Tool adoption. Buying and using a capable product is procurement plus practice. The competency is the design — knowing precisely where in the month-end the model's output enters, and what would catch it if it were confidently wrong. AIG-11 covers tool evaluation in depth.

3.3 Digital reporting

Definition. Building the path from source data to the decision it informs so that the number in the report can be traced to its origin, is produced the same way every period, and is available in time to change something.

By level. Foundation — produces reports from the established pipeline and identifies data that looks wrong. Practitioner — owns a reporting product: manages its data sources, its refresh cycle and its definitions, and resolves discrepancies between it and the underlying systems. Professional — designs the reporting architecture for a project: source of truth per data item, lineage, definition

ownership, the audience-by-audience content decision, and the deliberate removal of what nobody uses. Leader — governs reporting across the portfolio, including the hard part: retiring reports, and enforcing one definition of a term across functions that each prefer their own.

Evidence. A reporting architecture or data-lineage document you produced; a reconciliation regime between the report and its sources; a definition dispute you resolved and the decision record; a report you decommissioned.

The false positive. Visual quality. A clear, well-designed dashboard drawing on an unreconciled extract is an attractive route to a wrong decision. The competency is in the lineage and the definitions, not the rendering. BPG-15 covers the visualisation craft.

3.4 Automation

Definition. Removing manual steps from the controls cycle without removing the controls: deciding what to automate, designing the automation so its failures are visible rather than silent, and maintaining it as the surrounding systems change.

By level. Foundation — runs automated routines and reports when one fails. Practitioner — builds and maintains automation for a defined task, with error handling and a reconciliation that proves the output. Professional — decides what should be automated and what should not, on the basis of frequency, stability and the cost of a silent error; builds in the break-detection that makes failure loud; and documents it so it survives the author. Leader — sets the automation standard across the function, including ownership, change control and the requirement that no automation is a single-person dependency.

Evidence. An automated process with its documentation, error handling and reconciliation control; a break you designed for and that later caught a real failure; a decision not to automate a task, with reasoning; a handover in which someone else took over your automation without you.

The false positive. Volume of scripts. Automation that works only while its author is present, or that fails silently and produces a plausible wrong number, has increased risk while appearing to increase efficiency. The question that separates the levels is: when this breaks, how will anyone know?

3.5 Responsible AI

Definition. Governing the use of AI so that its outputs are explainable, its data use is lawful and proportionate, its known biases are considered before reliance, its use is disclosed where a reader would reasonably want to know, and there are named cases where it is not used at all.

By level. Foundation — follows the organisation's AI use policy, including on confidential project data. Practitioner — applies the policy in practice: knows what may be entered into which tool, discloses AI assistance where required, and raises outputs that look biased or unexplainable. Professional — decides the responsible-use position for a project: what may be delegated, what must be disclosed, what data may be used, and where explainability is a precondition of reliance rather than a nicety. Leader — owns the AI governance framework, including procurement criteria, incident handling, and the decision to prohibit a use case the organisation would find commercially convenient.

Evidence. An AI use policy or project-level position you authored; a disclosure convention you introduced into reporting; a documented refusal of a use case, with reasoning; an incident review after an AI-related error, with the control change that followed.

The false positive. Policy possession. Having a policy, or having attended AI ethics training, is not the competency. The competency shows in a decision that cost something — a tool not adopted, a use case declined, a disclosure made that invited an awkward question. AIG-08 and AIG-09 own the underlying framework.

3.6 Human validation

Definition. The specific, personal act of checking an AI-derived or automated output before it is relied on, and taking ownership of it as your own work. Includes designing the check so it is capable of detecting the error that matters, performing it, recording it, and being answerable for what you signed.

By level. Foundation — performs the prescribed check and escalates a failure rather than passing it on. Practitioner — validates outputs within their scope against independent sources, and records what was checked and what was found. Professional — designs the validation for a workflow: what an independent check is for this output, what sample size or coverage is enough, which errors the check would miss, and what evidence the check leaves behind. Rejects output that cannot be validated. Leader — sets the validation standard for the function, ensures that validation cannot be discharged by a tick, and owns the failures of validated output as failures of the standard rather than of the individual.

Evidence. A validation protocol you designed, naming the independent source and the tolerance; signed validation records with findings, including findings that stopped a release; a rejection of an unvalidatable output and what replaced it; a post-incident revision of a check that had missed something.

The false positive. The review tick. A signature confirming that output "has been reviewed", with no statement of what was checked against what, is a record of a signature. The house position is that AI proposes and the professional disposes; that is only true where the disposal is a real, documented, independent check performed by someone who could have caught the error.

— 4. Where the boundaries actually lie

Four pairs are confused often enough to be worth stating flatly.

Earned value management and performance measurement. Earned value is one measurement system, concerned with work performed against plan and cost. Performance measurement is the decision about what the project measures at all, including everything earned value cannot see — productivity, rework, interface health, data quality.

Predictive analytics and AI-enabled project controls. Predictive analytics is a class of technique applied to data. AI-enabled project controls is an operating-model competency: where any such technique enters the process, under what control. A person can be strong in one and absent in the other.

Automation and digital reporting. Automation removes manual steps; digital reporting establishes what the number means and where it came from. Automating an unreconciled report produces wrong numbers faster.

Responsible AI and human validation. Responsible AI is the governing position — what may be used, by whom, disclosed how. Human validation is the individual act at the point of reliance. An

organisation can have exemplary policy and no validation, and the failures will look like model failures when they are control failures.

— 5. How this goes wrong

The list is treated as fourteen boxes to tick, all at leader level. Nobody has that profile, and a candidate who claims it is telling the assessor something unintended. Real evidence portfolios are uneven; say which competencies are your depth and which are your working knowledge.

Overlapping competencies are evidenced with the same artefact four times. One month-end pack offered as evidence for cost management, earned value, performance measurement and digital reporting. It may legitimately support two, if the artefact shows different decisions. Reusing it across four means the candidate has not distinguished the competencies, and the assessor will notice.

The AI competencies are evidenced by tools rather than decisions. A list of platforms used answers a question nobody asked. Every one of the six data and AI competencies is evidenced by a decision: where the output enters, what check stands between it and reliance, what was declined.

The false positive is defended rather than addressed. A candidate whose risk evidence is register maintenance argues that maintaining the register is important. It is. It is also not the competency, and arguing the point burns the assessor's goodwill on the competency where the candidate is weakest.

Governance is evidenced by documents nobody follows. The quickest test an assessor applies to a governance claim is to ask what the governance has prevented recently. A blank answer converts a claimed leader-level competency into evidence of authorship.

The set is quietly reordered or renamed to fit an in-house framework. It then no longer maps to the credential, and a candidate prepares against a set that does not exist. Map onto the published set and keep the published names.

Recency is not stated. A strong demonstration from a project that closed five years ago is offered without a date, and the omission reads as concealment even when it is not. Date every artefact, and say plainly where a competency has not been exercised lately.

— 6. Checklist

For a candidate assembling evidence, or an assessor reviewing it:

- [] Have all fourteen competencies been addressed, including the ones where the honest answer is "practitioner, and here is why"?
- [] Does each competency have at least one **dated artefact**, with an identifiable contribution?
- [] Is any single artefact doing duty for more than two competencies?
- [] For each professional-level claim, does the evidence show a decision made where procedure was silent?
- [] For each leader-level claim, is there a rule others applied, and a failure that was owned?
- [] Has each competency been tested against its **false positive** in §2 and §3?
- [] Are the six data and AI competencies evidenced by decisions rather than by tool names?
- [] Is human validation evidenced by a protocol and findings, rather than by a review signature?

- [] Is the boundary between overlapping competencies (§4) reflected in the evidence chosen?
- [] Is every artefact dated, and is anything older than three years flagged as historical?

— Related

- [CMP-01 – The PCI competency framework – overview](#) — what the set is for and how each entry is structured
- [CMP-02 – Competency levels and how they are evidenced](#) — the four-level scale used throughout this document
- [CMP-06 – Cost and commercial competencies in depth](#) — cost management, earned value, forecasting and commercial control at working depth
- [CMP-07 – Planning, scheduling and risk competencies in depth](#) — schedule and risk competence, including quantitative analysis
- [CMP-08 – Data, digital and AI competencies in depth](#) — the six-competency chain from analytics to human validation
- [CMP-10 – Assessing competence – evidence, rubrics and moderation](#) — how these claims are actually tested

— Sources and standards

- PCI platform certification catalogue ([backend/Data/MultiCert.cs](#)), verified August 2026 — the PCL-AI seeded competency set, fourteen items, reproduced here in the published order and wording.
- PCI Body of Knowledge domains 4, 5, 6, 7, 10, 12 and 13 ([docs/bok/](#)) — the knowledge base underlying these competencies. Cited by domain; not reproduced.

— Status and version

Founding-stage document · Version 1.0 — effective date to be confirmed · Reviewed under PCI governance. PCI makes no claims of accreditation or recognition beyond what is true today.