

EXB-06 · EXAM BLUEPRINT

Job-task analysis: where the blueprint gets its authority

The study that will fix the item counts, test the 40/40/20 weighting against evidence, and replace expert opinion with practitioner data

VERSION 1.0 · DRAFT
2026-08-04

PROJECT CONTROLS INSTITUTE GLOBAL, INC.
AI proposes. The professional disposes.

— Summary

A blueprint cannot validate itself. A job-task analysis is the study that asks practitioners what the work actually consists of, how often each task is performed and how badly it goes when performed wrongly, and then links those tasks back to the examination's content. This document explains what a job-task analysis is, what the PCL-AI study will settle — including the item counts, which do not exist today — how the method works step by step, and how the resulting data converts into sampling weights. It works the linkage arithmetic through with illustrative figures and states what stands until the study reports. It applies to the PCL-AI only; no examination blueprint exists yet for PFL-AI or PML-AI.

— About this document

Document	EXB-06
Series	Exam Blueprint
Version	1.0
Status	draft
Date	2026-08-04
Unresolved placeholders	2

Notice. This document is an educational publication of Project Controls Institute Global, Inc., and does not constitute accounting, tax, legal, financial or other professional advice. It must not be relied upon as such. Where a topic depends on the contract or the governing law, entitlement and treatment vary by jurisdiction; take local advice.

Standards. Standards and frameworks are referred to by name and described in this publication's own words. No standard's text, tables, diagrams or question banks are reproduced. All trademarks remain the property of their respective owners, and no endorsement by any standards body, employer or vendor is implied or should be inferred.

Status. Project Controls Institute Global, Inc. is a Delaware Non-Stock Corporation. The Institute has been developed with reference to ISO/IEC 17024 personnel-certification principles and is **not accredited** by ANAB, IAS or any other accreditation body. It intends to seek 501(c)(3) recognition, which has not been granted. No governmental approval, academic equivalence, or guaranteed employment, promotion or salary outcome is claimed or implied.

Examples. All worked examples are illustrative and fictional. Figures are chosen to demonstrate a method, not to describe any real project, organisation or market. Sample examination items, where they appear, are study material maintained separately from any live examination bank.

Completeness. Passages marked **[CONFIRM: ...]** are operational or market facts that have not yet been confirmed from a cited source. A document carrying them is a working draft and is not approved for publication.

© 2026 Project Controls Institute Global, Inc. · All rights reserved.

Contents

1. The question a blueprint cannot answer about itself.....	4
2. What a job-task analysis is	4
3. What the PCL-AI study will settle.....	5
3.1 The item counts	5
3.2 Whether the 40/40/20 weighting survives contact with evidence.....	5
3.3 The sampling density within groups	5
3.4 The description of the minimally competent candidate	6
3.5 Whether a per-group hurdle is warranted.....	6
4. The method	6
4.1 Define the role and its boundary.....	6
4.2 Build the task inventory	6
4.3 The validation survey	6
4.4 Sampling.....	7
4.5 Analysis: from ratings to weights.....	7
4.6 Worked example of the linkage arithmetic.....	7
4.7 Review and adoption	8
5. What stands until the study reports	8
6. How a job-task analysis goes wrong.....	9
7. Cadence	9
8. Checklist	10
Related	10
Sources and standards.....	11
Status and version	11

Job-task analysis: where the blueprint gets its authority

In one paragraph. A blueprint cannot validate itself. A job-task analysis is the study that asks practitioners what the work actually consists of, how often each task is performed and how badly it goes when performed wrongly, and then links those tasks back to the examination's content. This document explains what a job-task analysis is, what the PCL-AI study will settle — including the item counts, which do not exist today — how the method works step by step, and how the resulting data converts into sampling weights. It works the linkage arithmetic through with illustrative figures and states what stands until the study reports. It applies to the PCL-AI only; no examination blueprint exists yet for PFL-AI or PML-AI.

Who this is for. Practitioners who may be asked to take part; employers deciding whether the credential measures the work their people actually do; and assessment professionals auditing the basis of the published blueprint.

— 1. The question a blueprint cannot answer about itself

EXB-01 – Examination blueprint – PCL-AI states what is examined and in what proportion. It is a careful document, built by experienced practitioners from a body of knowledge that took a great deal of work to assemble. And it cannot, on its own, answer the question that matters most:

How do you know the examination measures the job?

Expert judgement is the right starting point and it is not sufficient, for a reason that has nothing to do with the quality of the experts. Experts are unrepresentative by definition. They are more senior than the candidate population, they have worked in a particular set of sectors, they remember the discipline as it was when they learned it, and they systematically over-value the parts they find interesting. Ask ten excellent controls professionals what a credential should test and you will get ten defensible answers, no two of which allocate weight the same way.

A **job-task analysis** is the instrument that replaces that argument with data. It is the difference between a syllabus that experts believe in and a syllabus that can be shown to correspond to the work.

— 2. What a job-task analysis is

A job-task analysis is a structured study of a role, with three components.

A task inventory — a comprehensive list of the discrete tasks that make up the role, written at a consistent level of granularity and in the language practitioners actually use.

A validation survey — the inventory put to a large, representative sample of practitioners, who rate each task on defined scales. The two scales that matter are **frequency** (how often do you perform this?) and **criticality** (how serious are the consequences when it is performed badly?).

A linkage — an explicit mapping from tasks to the knowledge and skill required to perform them, and from there to the examination's content structure.

The output is not an opinion about what should be examined. It is a defensible statement, traceable to practitioner responses, about what the role consists of and which parts of it carry the most weight.

— 3. What the PCL-AI study will settle

Five things, of which the first is the one this series repeats most often.

3.1 The item counts

No item count exists for the PCL-AI. Not published, not configured, not privately assumed. Anyone quoting one is quoting something the Institute has not produced.

Item counts will be confirmed by the job-task analysis, in combination with a reliability target. The logic runs in that order and not the other way round: the study establishes what proportion of the role each content area represents; the reliability requirement establishes how many items are needed to measure each area with acceptable precision; the total follows from both. A sub-score computed from a handful of items carries a wide margin of error, which is why an area's item count is not simply its share of the blueprint multiplied by a number chosen in advance.

Setting the total first and reverse-engineering a justification is the failure mode. It is also, in the Institute's judgement, how a large number of professional examinations arrived at their item counts.

3.2 Whether the 40/40/20 weighting survives contact with evidence

The current weighting is a considered judgement, defended in [EXB-02 – Domain weightings and the 40/40/20 rule](#) §4 on grounds the Institute is willing to argue for. It is not yet an evidenced one.

The study will produce weighted task data that maps onto the three groups. If that data supports the current weighting, the weighting stands and is thereafter evidenced rather than merely argued. If it does not, the weighting changes — published as a major version of [EXB-01](#), with an effective date and a notice period. The Institute has committed to that in advance, in public, which is the only point at which such a commitment costs anything.

3.3 The sampling density within groups

[EXB-01](#) §4 states that weights attach to groups, not to individual domains. The study will show which domains within each group carry the most weighted task activity, and therefore where sampling should be denser. That governs form assembly and it governs where the item bank needs to be deep — the commissioning input described in [EXB-04 – How items are written, reviewed and retired](#) §3.

3.4 The description of the minimally competent candidate

The standard-setting study in [EXB-05 – Standard setting: modified Angoff and the pass mark](#) depends entirely on a written description of the just-competent professional. That description should be derived from the task data, not composed by the standard-setting panel from memory. A panel writing its own description of the borderline candidate is doing two studies at once and doing neither properly.

3.5 Whether a per-group hurdle is warranted

[EXB-02 §6](#) sets out the honest limitation of a compensatory scoring model, and says the question of a per-group minimum will be put to the job-task analysis. If the data shows that competence in one group cannot in practice substitute for competence in another — that they are distinct capabilities rather than facets of one — that is an argument for a hurdle. If it shows they covary strongly, it is an argument against. This is exactly the sort of question that opinion cannot settle and data can.

— 4. The method

4.1 Define the role and its boundary

The study begins by defining who is in scope, because a task inventory for "project controls" that includes everyone from a document controller to a portfolio director describes no role at all.

The PCL-AI attests competence at Leader level in the integrated discipline of project controls, cost engineering and project finance, and it is open to anyone with three years of professional experience in any field. The boundary must accommodate both facts: the study describes the role the credential attests, across sectors and regions, without narrowing to any single organisation's job description.

4.2 Build the task inventory

Three inputs, in this order:

Document review — the existing body of knowledge, published role descriptions, competency frameworks, and the artefacts controls professionals actually produce. This generates a first draft, and it is only a draft, because documents describe what organisations say the job is.

Practitioner workshops — facilitated sessions with working professionals across sectors and regions, who add what the documents missed, delete what no longer happens, and correct the language. This is where most of the value is created, and it is the stage most often skipped.

Editing to a consistent grain — the hardest editorial task in the study. "Manage cost" is too coarse to rate. "Open the cost report spreadsheet" is too fine. Tasks must sit at a level where a practitioner can answer both rating questions meaningfully, and every task in the inventory must sit at roughly the same level, or the ratings are not comparable.

The inventory must be **comprehensive and non-overlapping**. A task that appears twice in different words gets double-counted in the weighting, which quietly inflates whichever area it sits in.

4.3 The validation survey

The inventory is put to a large sample of practitioners. Each task is rated on defined scales.

Frequency — how often the respondent performs the task, on a scale from never to daily.

Criticality — how serious the consequences are when the task is performed incorrectly, from negligible to severe.

A third question is worth asking and is often omitted: **where the capability was acquired** — on the job, through formal study, or not yet at all. That answer tells the Institute where an examination adds most value, because assessment adds least where experience reliably teaches the skill and most where it does not.

Respondents also supply role, sector, region and years of experience, so that responses can be analysed by segment. A weighting that holds across sectors is a stronger result than an average that conceals two opposing patterns.

Scale design and survey construction: [\[CONFIRM: job-task analysis panel composition and survey design\]](#).

4.4 Sampling

The survey's authority rests entirely on who answered it, and this is where job-task analyses most often fail quietly.

A sample dominated by one sector, one region or one seniority band produces a weighting for that group and presents it as a weighting for the discipline. The Institute's requirements are that the sample spans sectors, spans regions, spans experience levels within the population the credential serves, and is large enough that segment-level analysis is possible rather than merely reported.

Response rates and any known non-response bias will be reported with the results. A study that publishes its weighting and not its sample is asking to be taken on trust in the one respect where trust is least warranted.

4.5 Analysis: from ratings to weights

Frequency alone is the wrong model, and it is the model most often used.

A task performed daily with trivial consequences matters less to a credential than a task performed twice a year that destroys a project when it goes wrong. Certification exists to protect against the second. So the weighting combines both scales — typically as a product, so that a task must score on both dimensions to carry weight, and a task rated zero on criticality drops out however often it is performed.

4.6 Worked example of the linkage arithmetic

Illustrative figures. Five tasks, mean ratings on 0–4 scales, standing in for a full inventory of several hundred. Products are shown to two decimal places; shares are of the total weighted inventory and are rounded to two decimal places.

Task (illustrative)	Frequency	Criticality	Weight = $F \times C$	Share
Reconcile the project cost ledger to the general ledger at period end	3.4	3.8	12.92	27.73 %
Compute and interpret cost and schedule performance indices	3.6	3.5	12.60	27.04 %
Prepare an estimate at completion and justify the method chosen	2.8	3.9	10.92	23.44 %
Validate a machine-generated cost forecast before it is issued	1.9	3.7	7.03	15.09 %
Price a bill of quantities remeasurement	1.2	2.6	3.12	6.70 %
Total			46.59	100.00 %

Substitutions: $3.4 \times 3.8 = 12.92$; $12.92 \div 46.59 = 0.2773 = 27.73 \%$. The remaining rows follow the same two steps.

Now read the fourth row, because it is the reason criticality is in the model at all.

Validating a machine-generated forecast has the lowest frequency of the four substantive tasks (1.9) and the second-highest criticality (3.7). On frequency alone it would rank near the bottom of the inventory and attract minimal sampling. Weighted by criticality it takes 15.09 per cent — because it is not yet a daily task for most practitioners, and when it is done badly a fluent, confident, wrong forecast goes to a board.

That single row is the empirical form of the Institute's position on AI: the weight belongs to the consequence, not to the frequency. It is also a concrete illustration of how a defensible 20 per cent group weight could emerge from data rather than from conviction — which is precisely what the study exists to establish or refute.

The final step is **linkage**: each task is mapped to the knowledge areas required to perform it, task weights are aggregated up through knowledge areas to domains to groups, and the resulting group shares are compared against the published 40/40/20. Where a task requires knowledge from more than one area, its weight is apportioned, and the apportionment rule is fixed before the analysis is run.

4.7 Review and adoption

The analysis goes to a review panel drawn from outside the study team, then to the Standards Committee and the Certification & Impartiality Committee for adoption. The full study — inventory, scales, sample composition, response rates, analysis and linkage — is published alongside whatever blueprint change it produces, including the case where it produces none.

Study completion: [\[CONFIRM: job-task analysis completion date\]](#).

— 5. What stands until the study reports

The Institute is founding-stage, and the honest position is this.

The published blueprint is **expert-derived and internally coherent**: thirteen domains, sixty-one knowledge areas, three groups, 40/40/20, built from a body of knowledge developed to handbook depth. It is a defensible specification, it is published in full, and it is the specification every form is assembled to today.

It is **not yet practitioner-validated**. That is a real limitation and it is stated here rather than left for a reviewer to discover. The certification framework is developed with reference to ISO/IEC 17024 personnel-certification principles, of which a validated relationship between the examination and the role is one; PCI is **not accredited** by any ISO/IEC 17024 accreditation body, and completing this study is one of the things that would have to be true before accreditation could sensibly be sought.

Until it reports: the weighting stands as published, the item counts remain unstated, and nobody at the Institute quotes a number for them.

— 6. How a job-task analysis goes wrong

A sample that represents one sector. The most common and most consequential defect. The weighting that results is real, precisely measured, and about the wrong population.

Ratings from experts instead of practitioners. A job-task analysis surveyed to a panel of senior experts reproduces the expert opinion it was supposed to replace, now wearing the authority of a study.

Frequency-only weighting. Under-serves exactly the low-frequency, high-consequence tasks a credential exists to protect against — the fourth row of §4.6.

An inventory at inconsistent grain. Coarse tasks absorb weight that fine tasks fragment. The area of the inventory written most carefully ends up looking least important.

Overlapping tasks. Double-counts weight into whichever area the duplication sits in.

Running the study and then ignoring it. The failure that destroys the credibility of every other stage: a study that produces an inconvenient weighting, followed by a blueprint that does not change and a report that is not published. The public commitment in §3.2 exists to make that outcome costly.

Treating it as a one-off. A task inventory describes the work at a moment. The Institute's discipline is changing fastest in exactly the area weighted at 20 per cent, so a study left unrepeated for a decade describes a job nobody has any more.

— 7. Cadence

A job-task analysis is repeated on a defined cycle rather than when someone notices it is stale. Two pressures set the interval.

The **structural content** of project controls — accounting, earned value, contracts, scheduling, risk — changes slowly, over years, and does not need frequent re-study.

Domain 13 content changes inside a single certification cycle. A candidate certified for three years may hold a credential whose AI content was validated against practice that has since moved substantially. The Institute's approach is therefore a full study on the long cycle, with a lighter tar-

geted refresh of the AI-related task inventory between full studies — which is the same reasoning that puts a mandatory AI-currency component into recertification.

— 8. Checklist

For a practitioner invited to take part:

- ☐ Rate frequency against what you actually do, not what your job description says.
- ☐ Rate criticality against the consequence of doing it badly, not the consequence of not doing it.
- ☐ Rate the whole inventory, including tasks you never perform — a genuine zero is data.
- ☐ Answer the segment questions honestly; they are what allow the analysis to detect sector bias.

For anyone auditing a job-task analysis — ours or another body's:

- ☐ Is the task inventory published, not just its conclusions?
- ☐ Is the sample composition published, by sector, region and experience?
- ☐ Are response rates and known non-response bias reported?
- ☐ Does the weighting model include criticality, or only frequency?
- ☐ Is the task-to-content linkage published, with the apportionment rule for multi-area tasks?
- ☐ Were the weights compared against the existing blueprint, and was the comparison published even where it disagreed?
- ☐ Is there a stated cadence for repeating the study?

Publishing a blueprint before the study that validates it is a choice with a cost: it invites exactly the question this document opens with, and the Institute cannot yet answer it with data. The alternative — to publish nothing until everything is validated — would leave candidates preparing against a syllabus nobody had seen. The position taken here is to publish the specification, name precisely what it does not yet rest on, and commit in advance to changing it if the evidence says so. That commitment is only worth anything because it was made before the evidence arrived.

— Related

- [EXB-01 – Examination blueprint – PCL-AI](#) — the specification this study will validate or change
- [EXB-02 – Domain weightings and the 40/40/20 rule](#) — the weighting that will be tested, and the hurdle question the study will settle
- [EXB-05 – Standard setting: modified Angoff and the pass mark](#) — the study that depends on this one's description of the borderline candidate
- [CMP-03 – PCL-AI: the fourteen competencies](#) — the competency set the task inventory must be reconcilable with

— Sources and standards

The method described here is standard job-task analysis practice in professional certification, set out in the Institute's own words; no published text, instrument or scale is reproduced. The figures in §4.6 are illustrative, invented for teaching, and independently recomputed before submission — they are not data from any survey, practitioner or organisation. The certification framework is developed with reference to ISO/IEC 17024 personnel-certification principles; PCI is not accredited by any ISO/IEC 17024 accreditation body.

No external citation is claimed for this document.

— Status and version

Founding-stage document · Version 1.0 · draft · Reviewed under PCI governance. PCI makes no claims of accreditation or recognition beyond what is true today.