SAL-01 · SALARY AND SKILLS REPORT

Salary and skills report — framework and methodology

The measurement rules that decide whether a number is fit to publish

VERSION 1.0 · DRAFT
2026-08-04

PROJECT CONTROLS INSTITUTE GLOBAL, INC.

AI proposes. The professional disposes.

Summary

This document sets out how the Institute will measure reward and skills in project controls: what the survey asks, who it asks, how bias is mitigated, the minimum number of responses a cell must contain before any figure is published, how outliers and currencies are handled, and what the report will never claim. It contains no salary figures, because the survey has not been run. The methodology is published first so that the numbers, when they exist, can be checked against the rules that produced them.

About this document

Document	SAL-01
Series	Salary and Skills Report
Version	1.0
Status	draft
Date	2026-08-04
Unresolved placeholders	6

Contents

Salary and skills report — framework and methodology

In one paragraph. This document sets out how the Institute will measure reward and skills in project controls: what the survey asks, who it asks, how bias is mitigated, the minimum number of responses a cell must contain before any figure is published, how outliers and currencies are handled, and what the report will never claim. It contains no salary figures, because the survey has not been run. The methodology is published first so that the numbers, when they exist, can be checked against the rules that produced them.

Who this is for. Project controls managers and heads of function who are asked to justify a pay position; human resources and reward specialists who buy or build benchmark data; and researchers or subject-matter experts who will review, run or criticise the Institute's survey.

1. What practitioners are currently forced to use

Ask a planner what their role pays and the honest answer is a stack of things that are not measurements. A recruiter's verbal range, which is a negotiating position. A job advert, which is an offer to a market and not a record of what anyone was paid. A colleague's disclosure, which is one observation with no context. A published "salary guide" that names no sample, no definition of the role, no reference date and no method — and which, on inspection, is often a summary of the publisher's own placements or of scraped advertisements.

Each fails in the same specific way: it reports a number without reporting what the number is a number of. A figure attached to the word "planner" is meaningless until you know whether that planner carries a portfolio or a single package, works on a rotation or from home, is paid a base salary or a day rate, in what currency, gross or net, and in which year. Project controls professionals would reject a cost report built to that standard in a heartbeat. Market data deserves the standard we hold our own reporting to.

2. Why this edition contains no numbers

The Institute has not run the survey. Therefore this series publishes the instrument, the taxonomy, the publication rules and the report shell — and no figures at all.

This is a deliberate sequence, not a gap waiting to be filled quietly. Three reasons.

First, a number you cannot defend is worse than silence. A published figure travels: into pay reviews, into business cases, into other people's articles, and eventually back to us stripped of its caveats. Once it is loose, the caveat cannot catch it. The only point at which we control that is before publication.

Second, publishing the method first is a falsifiable commitment. Anyone can write a methodology appendix after the data is in, tuned to justify what the data happened to say. A methodology published before fielding can be checked against the report that follows it: if the minimum cell size in §6 is relaxed for one attractive table, that is visible. We are making it easy to catch ourselves.

Third, it is the discipline we certify. A controls professional who reports an estimate at completion (EAC) they cannot substantiate has failed at the core of the job. A professional body that publishes a methodology before it publishes a number is demonstrating, at its own cost, the behaviour it examines candidates on. If we were willing to invent a salary range for the sake of a press cycle, no reader should trust anything else we publish.

This is not an apology for an absence. Everything needed to run an honest annual report is here: a fielded instrument (SAL-02), a role taxonomy and levelling rubric that makes comparison possible (SAL-03), a skills vocabulary (SAL-04), and a report template whose tables are structurally complete and empty (SAL-05). The missing element is field data, and field data is the one thing that cannot be written at a desk.

3. What the report measures

Unit of analysis. One individual, in one role, over one defined reference period. Not a job, not a vacancy, not an employer's grade.

Reference period. A single stated period — the twelve months ending on a declared reference date — so that every respondent answers about the same window [CONFIRM: fielding window and reference date for the first cycle — to be set by the Institute's research lead].

In scope: individuals whose substantive work is project controls as defined by the role taxonomy in SAL-03, in any sector and any country, employed or self-employed, whether or not they hold any certification.

Out of scope: aspirational or advertised pay; pay for roles the respondent no longer holds outside the reference period; employer-level pay policy; equity in privately held companies where no defensible valuation exists; and any figure the respondent is contractually barred from disclosing.

Remuneration is decomposed, never asked as one number. Base pay, fixed allowances, variable pay, overtime, day rate and employer retirement contributions are separate questions with separate bases. This is the single most common defect in market pay data: one respondent reports base, the next reports total cash, and the resulting "median" is a summary of two different quantities. SAL-02 §7 specifies the decomposition and the wording.

4. Sampling frame and recruitment

There is no register of the world's project controls professionals. That is a fact about the profession, not a difficulty to be papered over, and it constrains everything downstream.

The frame is therefore constructed, not enumerated. Recruitment runs through the Institute's own channels (members, candidates, newsletter subscribers), professional social networks, employer and consultancy contacts who agree to circulate the link without seeing responses, and open publication of the survey URL so that participation does not require any relationship with the Institute.

This produces a non-probability, volunteer sample. No sample size converts a volunteer sample into a probability sample. The report will say so on the same page as its first table, in those words, every year. What we can do is characterise the sample precisely, mitigate the biases we can name, and refuse to publish where the sample is too thin to bear a claim.

Open participation is a deliberate control. If only members could respond, the sample would be a study of PCI members' pay reported as a study of the profession. Membership and certification status are recorded as respondent characteristics — never as eligibility.

5. Self-selection and the other biases we can name

Bias	How it enters	Mitigation
Self-selection	People with an unusual position — very well or very badly paid — have more reason to answer	Recruit through channels with no reward angle; keep the questionnaire short enough that indifferent respondents finish; publish composition against every characteristic collected
Seniority skew	Senior practitioners are more visible and more often reached	Report every cut by level; never present an all-respondent figure as "the profession"; disclose the level distribution beside every headline
Geographic and language skew	Recruitment travels furthest in the languages and networks we already occupy	Translate the instrument [CONFIRM: translation languages for the first cycle] (see `SAL-02` §12); suppress regions below threshold rather than merging them into a flattering aggregate
Non-response within the instrument	Pay questions are abandoned more often than any other	"Prefer not to say" on every remuneration question, so a respondent can skip one figure without abandoning the survey; report item-level response counts, not just the overall count
Anchoring	Showing a figure in the questionnaire moves the answers towards it	No figure, band or "typical" value appears anywhere in the instrument before the respondent's own answer
Social desirability	Reported skills drift upward; reported dissatisfaction drifts downward	Skills are asked against behavioural anchors (`SAL-04`), not adjectives; satisfaction items are separated from identity items
Survivorship	People who left the profession are absent by construction	Stated as a limitation; the report never infers anything about retention or attrition

Two of these deserve emphasis because they are the ones that quietly destroy a pay survey. **Anchoring** is why the instrument shows nothing. **Composition** is why every table carries its own denominator: a "median for cost engineers" computed from a sample that happens to be four-fifths one region is a regional figure wearing a global label.

6. The publication threshold — minimum cell size and suppression

A cell is any combination of reporting dimensions for which a figure might be published: role × level, role × level × region, sector × contract type, and so on. The rules below are absolute. They are applied by the analyst before anyone sees the output, and they are not waived for a table someone wants.

#	Rule	Test	Consequence
1	Minimum cell size	Fewer than 30 valid responses in the cell	No statistic of any kind is published for that cell. Not a median, not a range, not an "indicative" figure, not a sentence describing it in words
2	Distribution detail	30 or more but fewer than 100 responses	Median and quartiles only. No deciles, no minimum, no maximum, no tail commentary
3	Employer dominance	One employer contributes more than a quarter of the cell, or the three largest contribute more than half	Cell suppressed. A figure that is really one organisation's pay policy is not a market figure
4	Complementary suppression	A suppressed cell can be recovered by subtracting published cells from a published total	A second cell is suppressed so the first cannot be reconstructed
5	Aggregation discipline	A thin cell is combined with another to reach the threshold	Permitted only along the pre-declared hierarchy in `SAL-03` §6, and the combined cell is labelled as combined
6	Comparison across cycles	Both cycles meet rules 1–3 on the same definition, and the relevant question wording was unchanged	Otherwise no year-on-year statement is made, in figures or in prose
7	Denominator disclosure	Every table, chart and sentence containing a figure	Carries its unweighted response count
8	Suppression log	Every cell suppressed under rules 1–4	Listed in an annex with the rule invoked, so suppression cannot be applied selectively to shape a story

Rule 1 has one property worth stating plainly: **it will suppress the cuts people most want.** Niche roles in smaller markets are exactly where individual practitioners have least information and most need it, and exactly where the sample will be thinnest. We will not fill that gap with a number we cannot stand behind. The honest output is a visible hole with a stated reason, and an invitation to help the next cycle reach the threshold.

7. Weighting — the method, and why the first cycle will not use it

Post-stratification corrects a sample whose composition differs from a known population. For stratum h, the weight is the population share divided by the sample share:

```
w_h = (N_h ÷ N) ÷ (n_h ÷ n)
```

where N_h is the population count in stratum h, N the population total, n_h the responses in stratum h, and n the total responses. Weights are trimmed at bounds declared before weighting is applied, because a single respondent carrying an enormous weight is a single respondent. Precision after weighting is reported through the effective sample size (the Kish formula), never through the raw count, which overstates what the weighted estimate can support.

The method requires N_h, and we do not have it. No authoritative population distribution of project controls professionals by role, level, sector and region is known to the Institute [CONFIRM: whether any nameable, dated population distribution exists that could serve as a weighting frame; weighting is not applied without one]. Weighting to a guessed population is worse than not weighting, because it converts a visible composition problem into an invisible one and attaches a scientific-sounding label to it.

Accordingly: **the first cycle is reported unweighted, with full composition disclosed.** If a defensible frame becomes available, weighting will be introduced as a clearly labelled addition, both weighted and unweighted results will be published side by side for the first cycle in which it applies, and the frame will be named and dated in the report.

8. Outliers, validation and unit errors

Three different problems get called "outliers", and conflating them is how good data gets deleted.

Entry errors are wrong values: a monthly amount typed into an annual field, a currency selected incorrectly, a digit added. They are detected by ratio and plausibility checks against bounds set per currency and basis before fielding and recorded in the fielding log. A record is corrected only where the intended value is unambiguous from the rest of the response; otherwise the affected item — not the whole response — is excluded.

Ineligible responses fail a screening rule, duplicate one respondent, or show patterns inconsistent with genuine completion: a straight-lined matrix, or a completion time below a pre-set floor. These are removed and counted.

Genuine extremes are real values far from the centre. The rule for handling them — exclusion above a stated distance from the quartiles, winsorising to a stated percentile, or retention — is declared before the data is examined and applied identically to every cell, never chosen after seeing which choice produces the tidier result.

Three disclosure rules apply to all three: the count of records affected is published; nothing is removed silently; and no record is ever removed because it makes a series less smooth.

9. Currency, conversion and the purchasing-power caveat

Respondents answer in the currency they are paid in, and state the basis (gross or net, annual, monthly, daily or hourly). Nothing is converted at the point of entry.

For reporting, amounts are converted to a single base currency at a single reference date using a named, published rate source [CONFIRM: foreign-exchange rate source and the single reference date used for conversion]. One date for the whole dataset, disclosed, so that conversion is a stated transformation rather than a moving target.

The caveat that must travel with every converted figure: **conversion compares nominal amounts, not living standards.** A converted figure says nothing about what the money buys, about tax and social contributions, about mandatory retirement contributions, about employer-provided healthcare, or about which of those a jurisdiction treats as pay at all. Employment, tax and pay-transparency law differ by country, and the report will not present any one country's treatment as the general case. A purchasing-power adjustment, if published, is a separate and separately labelled table using a named index [CONFIRM: whether a purchasing-power adjustment is published for the first cycle, and which named index would be used] — never silently folded into the headline.

10. Disclosure — what accompanies every figure

Every figure the report publishes is accompanied by six things, without exception: the definition of the population it describes; the unweighted response count; the reference period; the remuneration component it measures; the currency and basis; and the rule under which any neighbouring cell was suppressed. A figure lacking any of the six is not published, and a figure quoted onward without them has been misquoted — which is the subject of `SAL-06`.

The instrument, the role taxonomy, the skills codebook and the suppression log are published as annexes in every edition, so a reader can reconstruct what was asked and reproduce the coding decisions.

11. Publication rules and governance

Three rules bind the report's production.

No commercial influence over content. No sponsor, advertiser, employer or recruitment business sees results before publication, chooses which cuts are published, or has any role in the methodology. Where a partner circulates the survey, that is disclosed and confers nothing else.

No individual or employer identification. Free-text answers are screened for identifying detail before coding, and that detail is removed rather than paraphrased. No employer is named in any output. Personal data handling is governed by the consent and privacy provisions in `SAL-02` §2 [CONFIRM: data controller, lawful basis, retention period and erasure procedure — pending legal review, see SAL-02 §2].

Corrections are public. If a published figure is wrong, the correction is posted to the channels that carried the original, with a dated note in the report itself. A quietly replaced file is a second error on top of the first.

12. Anti-goals — what this report will never claim

Stated as prohibitions, so that a future edition drifting towards any of them is visibly in breach.

1. It will never claim to state the market rate, the going rate, or what a role "should" pay. It reports what a defined set of respondents said they received.
2. It will never claim that holding a PCI credential — or any credential — causes a pay outcome. The Institute makes no claim of guaranteed employment or salary outcomes, and a cross-sectional survey cannot establish causation in either direction.
3. It will never publish a figure for a cell below the threshold in §6, however often it is requested.
4. It will never present advertised pay as paid pay, or blend advertisement data into survey results.
5. It will never forecast pay, or describe a movement between two cycles as a trend.
6. It will never present one jurisdiction's employment, tax or pay-transparency law as universal.
7. It will never publish anything from which an individual or an employer can be identified.
8. It will never restate a prior cycle's figure after a definition change without labelling the break.
9. It will never be sold, licensed or presented as a job-matching or pay-setting service for an individual.
10. It will never imply accreditation, endorsement or recognition the Institute does not hold.

13. How this goes wrong

The seniority spiral. Year one skews senior; the headline figure is high; the survey acquires a reputation as the place where well-paid people report; year two skews further. The defence is refusing to publish an all-respondent headline at all, and reporting only within level.

The title trap. Respondents self-classify by their employer's title, and "senior planner" at one company is "planning manager" at another. Left alone, this produces cells that are internally incoherent. `SAL-03` exists to break it: respondents are classified by accountability and autonomy, with the title recorded separately as free text.

The single-employer cell. A large organisation circulates the survey enthusiastically, and a regional cell becomes that organisation's pay bands with a market label on it. Rule 3 in §6 catches it — but only if employer is collected at a level that makes the test possible, which is why `SAL-02` asks for employer type and, optionally, an employer identifier used solely for this test and never reported.

The silent definition change. A question is improved between cycles, and the improved series is compared with the old one because both are called "base pay". Rule 6 forbids it; the change log in `SAL-05` makes the breach visible.

The press-release cut. Someone finds the one comparison that makes a headline, and it is published without its denominator — or a thin cell is merged upward until it clears the threshold and is then discussed as though it were the original. The rules that every figure carries its response count, and that combination follows a declared hierarchy, exist precisely for the moment when they are inconvenient.

14. Worked example — the arithmetic, stated symbolically

No illustrative figures are used in this series, by design. Elsewhere in the Institute's publication programme numeric worked examples are the backbone of the teaching. Here they would be a hazard: an invented respondent distribution, however clearly labelled, is exactly the kind of number that gets screenshotted without its label. So the arithmetic is shown in symbols only.

Post-stratification weight for stratum h, with population and sample counts as defined in §7:

```
w_h = (N_h ÷ N) ÷ (n_h ÷ n)
```

Read it as: a stratum that is under-represented in the sample relative to the population receives a weight above one, and vice versa. Applied to the first cycle, the expression is undefined, because N_h has no value we can name. That is the whole argument of §7 in one line.

Effective sample size for a weighted estimate, over respondents i with weights w_i:

```
n_eff = (Σ w_i)² ÷ Σ (w_i²)
```

Read it as: the more unequal the weights, the fewer independent observations the estimate is really worth. When weighting is used, precision claims are made against n_eff, never against the raw count.

The suppression test for a cell c with response count n_c and largest-employer share s_c: publish only if $n_c \geq 30$ **and** $s_c \leq ¼$ **and** the sum of the three largest employer shares ≤ ½. All three conditions, not any of them.

15. Checklist — the pre-publication gate

Run before any edition of the report is released. Any unticked line stops publication of the affected table, not of the report.

- [ ] Every published cell meets the minimum response count in §6, rule 1
- [ ] Employer-dominance test run on every published cell, and its result recorded
- [ ] Complementary suppression applied where a suppressed cell is recoverable by subtraction
- [ ] Suppression log complete, with the rule invoked for each suppressed cell
- [ ] Every table, chart and figure-bearing sentence carries its unweighted response count
- [ ] Currency, basis, reference period and remuneration component stated for every figure
- [ ] Conversion source and reference date stated; no purchasing-power claim without a named index
- [ ] Outlier and exclusion rules were declared before the data was examined, and the affected counts published
- [ ] No year-on-year statement without both cycles meeting the threshold on unchanged wording
- [ ] Composition annex published against every respondent characteristic collected
- [ ] Instrument, taxonomy and codebook published as annexes
- [ ] Anti-goals in §12 re-read against the final draft, including the headline and the social posts
- [ ] No sponsor or partner has seen results before publication
- [ ] Free text screened for identifying detail before coding

---

Related

- `SAL-02 — The survey instrument` — the questionnaire this methodology governs, question by question
- `SAL-03 — Role taxonomy and levelling` — the classification that makes any cross-employer comparison possible
- `SAL-05 — Report template and data tables` — the empty shell these rules will fill
- `SAL-06 — Using market data honestly` — the reader-facing companion, including how to read a median
- `EXB-06 — Job-task analysis: where the blueprint gets its authority` — the same evidential discipline applied to the examination

Sources and standards

No external market source is cited, because none has been consulted and none would be cited without consultation. The internal artefacts this document depends on are the Institute's canonical facts register (`docs/publication-framework/00-framework/CANONICAL-FACTS.md`), the editorial standard governing evidence and claims (`00-framework/EDITORIAL-STANDARD.md` §4), and the unverified-fact rule in the framework README §4. Statistical methods named in §7 and §14 — post-stratification, weight trimming, the Kish effective sample size, winsorising — are standard survey-methodology techniques described here in our own words; the report's methodology annex will cite the specific texts relied on once the research lead has selected them.

Status and version

Founding-stage document · Version 1.0 · draft · Reviewed under PCI governance. PCI makes no claims of accreditation or recognition beyond what is true today.