

BPG-19 · BEST PRACTICE GUIDES

Project controls assurance and health checks

Testing whether the numbers can be relied on, early enough for the answer to matter

VERSION 1.0 · DRAFT
2026-08-04

PROJECT CONTROLS INSTITUTE GLOBAL, INC.
AI proposes. The professional disposes.

— Summary

A project controls health check tests whether a project's control system can produce a number a decision can rest on. This guide sets out what such a review examines, how to score it on evidence rather than impression, how much independence the reviewer needs and what to do without it, the difference between assurance and audit, why timing determines whether findings can be acted on at all, and how to write findings that are accepted rather than resented. The worked example scores a review and then demonstrates, arithmetically, why a composite score cannot express a blocking weakness.

— About this document

Document	BPG-19
Series	Best Practice Guides
Version	1.0
Status	draft
Date	2026-08-04
Unresolved placeholders	0

Notice. This document is an educational publication of Project Controls Institute Global, Inc., and does not constitute accounting, tax, legal, financial or other professional advice. It must not be relied upon as such. Where a topic depends on the contract or the governing law, entitlement and treatment vary by jurisdiction; take local advice.

Standards. Standards and frameworks are referred to by name and described in this publication's own words. No standard's text, tables, diagrams or question banks are reproduced. All trademarks remain the property of their respective owners, and no endorsement by any standards body, employer or vendor is implied or should be inferred.

Status. Project Controls Institute Global, Inc. is a Delaware Non-Stock Corporation. The Institute has been developed with reference to ISO/IEC 17024 personnel-certification principles and is **not accredited** by ANAB, IAS or any other accreditation body. It intends to seek 501(c)(3) recognition, which has not been granted. No governmental approval, academic equivalence, or guaranteed employment, promotion or salary outcome is claimed or implied.

Examples. All worked examples are illustrative and fictional. Figures are chosen to demonstrate a method, not to describe any real project, organisation or market. Sample examination items, where they appear, are study material maintained separately from any live examination bank.

Completeness. Passages marked **[CONFIRM: ...]** are operational or market facts that have not yet been confirmed from a cited source. A document carrying them is a working draft and is not approved for publication.

© 2026 Project Controls Institute Global, Inc. · All rights reserved.

Contents

1. The question a health check is actually asking	4
2. Assurance and audit are different instruments	5
3. Evidence-based scoring	5
4. Independence, and what to do without it.....	6
5. Timing: early enough to matter	7
6. Writing findings that get acted on	7
7. How this goes wrong	8
8. Worked example	8
8.1 The sample and its coverage.....	9
8.2 The scores	9
8.3 Why 55 per cent is not the finding	9
8.4 Testing whether a dependency rule rescues the composite	10
8.5 What the deliverable should therefore be	10
8.6 Closure of the previous review	10
9. Checklist	11
Related	11
Sources and standards.....	12
Status and version	12

Project controls assurance and health checks

In one paragraph. A project controls health check tests whether a project's control system can produce a number a decision can rest on. This guide sets out what such a review examines, how to score it on evidence rather than impression, how much independence the reviewer needs and what to do without it, the difference between assurance and audit, why timing determines whether findings can be acted on at all, and how to write findings that are accepted rather than resented. The worked example scores a review and then demonstrates, arithmetically, why a composite score cannot express a blocking weakness.

Who this is for. PMO leads and heads of project controls who commission or run reviews; controls managers who are reviewed; and sponsors and assurance committees who have to decide what a review's output entitles them to conclude.

— 1. The question a health check is actually asking

A project controls health check is not an inspection of whether procedures are being followed. It asks one question: **can this project's control system produce a number that a decision can safely rest on, and if not, exactly where does the chain break?**

That framing rules out the compliance sweep — a checklist of documents that exist, signed by people who exist — because a project can hold every artefact the framework requires and still be incapable of saying what it will cost. It also rules out the opinion piece, in which a reviewer walks the project for two days and reports what they felt. Neither can be challenged, defended or acted on with precision.

The chain being tested runs in one direction, each link depending on those before it:

- **Scope is defined and structured** — one breakdown shared by budget, schedule and cost ledger. Without it nothing downstream reconciles and every comparison is manual.
- **A baseline exists and is controlled** — approved, dated, with a change process that is used. Without it, variance analysis is comparison with the most recent opinion.
- **Cost is captured completely and on time** — commitments, actuals and accruals, at the cut-off, coded to the structure. Without it, every efficiency measure has an unreliable denominator.
- **Physical progress is measured against rules** — documented rules of credit, applied consistently, verified independently on a sample. Without it, earned value is a claim.
- **Forecasting is method-based** — a method chosen for a reason connected to the cause of the variance, with the assumption written next to the answer.
- **Risk and change are live** — a register that changes decisions, and a change process that captures scope movement before it is executed.
- **Reporting reaches the decision** — the right measure, at the right cadence, to the person who acts.

A review that walks this chain in order produces self-prioritising findings, because a break early in the chain invalidates everything after it. §8 shows why a scoring scheme cannot substitute for that property.

— 2. Assurance and audit are different instruments

The two words are used interchangeably and should not be. The distinction changes what the reviewer may conclude and what the project is obliged to do.

	Audit	Assurance / health check
Question	Was the defined requirement complied with?	Will this control system produce a reliable number?
Criterion	An existing rule — procedure, contract, standard, policy	Fitness for the decisions the project must support
Output	Exceptions against the rule	Capability, and the weaknesses that limit it
A project can fail by	Not following the process	Following the process, where the process is inadequate
Relationship to the team	Independent testing; not negotiable in substance	Advisory; negotiable in how, not in whether
Typical timing	After the fact	Early enough to change the outcome

The consequential row is failing while compliant. An audit against a controls procedure can find full compliance on a project whose forecast is meaningless, because the procedure never required the forecast method to be justified. Assurance catches that only if its criterion is fitness rather than conformance.

Projects need both. What causes damage is running one and reporting it as the other: an audit dressed as assurance produces a clean report on an uncontrolled project, and assurance dressed as an audit produces findings the team treats as non-negotiable when they are recommendations.

— 3. Evidence-based scoring

A score that cannot be traced to something the reviewer inspected is an opinion with a number attached, and it will be argued with successfully, because there is nothing behind it. Three requirements make a score defensible.

A published rubric, written before the review. Each element has descriptors for each score, expressed as observable conditions rather than adjectives. "0 poor, 1 weak, 2 adequate, 3 good, 4 excellent" is not a rubric. This is, for progress measurement:

Score	Condition
0	No documented rules of credit; progress is a judgement recorded as a percentage.
1	Rules of credit exist as a document; in practice progress is claimed by percentage judgement.
2	Rules documented and applied on major packages; no independent verification.
3	Applied across the measured scope, with independent verification on a defined sample.
4	As 3, and progress, valuation and earned value reconciled each period with differences explained.

That version can be evidenced, disputed on facts, and re-tested at the next review by a different person who reaches the same answer.

Evidence recorded per score: the artefact, its date, the test applied. "Progress measurement scored 2" is an assertion. "Progress measurement scored 2: rules of credit document rev C dated 14 March covers packages 100–400; the April progress return for package 500 records 62 per cent with no supporting quantity record; no evidence of independent verification in the last three periods" is a finding, and it survives the reviewer leaving the organisation.

A stated sample, with coverage expressed two ways — which accounts, why those, how many of how many, and what proportion of value. §8 shows a case where the two figures are 8.6 per cent and 65.3 per cent, and reporting either alone misrepresents the review.

Two disciplines protect the scoring from itself. **Score the element, not the person:** a low score on a project whose measurement was never resourced is a finding about resourcing. And **agree the condition before writing the conclusion** — show the project the factual observations, correct any that are wrong, then write cause and consequence. Almost every dispute about a review is a dispute about a fact that could have been corrected in ten minutes.

— 4. Independence, and what to do without it

The reviewer must not be assessing their own work, and must not report through the person whose project is being reviewed. Both conditions are frequently breached, usually with good intentions: the most knowledgeable person available is the one who helped build the thing. Four threats are worth naming. **Self-review** — the reviewer designed the process, template or baseline being examined. **Reporting line** — their next assignment, appraisal or invoice depends on the person receiving the findings. **Familiarity** — a long relationship makes the uncomfortable finding costly to write and makes it easy to accept an explanation that would not be accepted from a stranger. **Advocacy** — the reviewer has argued publicly for the approach in use.

Full independence is often unavailable, particularly in smaller organisations. The professional answer is neither to pretend nor to refuse. It is to **declare the impairment on the face of the report, restrict the conclusions accordingly, and add a compensating control:** pair the internal reviewer with someone from another project, have a second person re-perform the scoring, and exclude the elements they built. A report stating "the reviewer designed the reporting template; element 6 was scored by a reviewer from the adjacent programme" is more credible than one claiming an independence nobody believes.

The reviewer also needs access, and access is granted, not assumed — to the cost system rather than extracts prepared for the review, to the schedule file rather than a printed bar chart, to

the change log and register, and to people without their manager in the room. Access limitations are themselves a finding.

— 5. Timing: early enough to matter

The value of a health check falls sharply with elapsed project time, for a structural reason. Most controls weaknesses are foundational: a coding structure that does not support the reporting the project needs, rules of credit never written, a baseline never properly approved. The cost of fixing one scales with everything built on top of it. Correcting a coding structure in month two means recoding two months of transactions; correcting it in month fourteen means recoding fourteen months, restating every report issued from them, and explaining to a board why the trend it has been watching has changed shape. The second is not fourteen times harder — it is often simply not done, and the project runs to completion on a structure everyone knows is wrong.

Three timing points earn their cost:

Before baseline acceptance. The highest-value review by a distance. Everything is still cheap to change, nothing has been reported externally, and findings can become conditions of baseline approval — the only mechanism that reliably gets them fixed.

At the first reporting cycle after mobilisation, when the design meets reality. This catches the gap between the controls plan and what the team can execute, while the correction is a process change rather than a restatement.

At each gate, and on trigger — a change of controls manager, a re-baseline, the first month a forecast moves beyond a defined tolerance, the mobilisation of a major subcontract package.

A review at sixty per cent complete can still be worth doing, but its honest purpose is different: not improving the control system, but establishing what the current numbers can and cannot support. Say which of the two jobs you are doing, because a late review reported as an improvement exercise generates recommendations nobody has the appetite to implement.

— 6. Writing findings that get acted on

The difference between a finding that is implemented and one that is resented is mostly structure. Five components, in this order:

Condition — what is, factually, with the evidence and its date. No adjectives. This is the part the project must be given the chance to correct before anything else is written.

Criterion — what should be, and whose rule that is. "The project's own controls execution plan section 4.2 requires..." is far stronger than "good practice requires...", because it removes the argument about whose standard applies. Where the criterion is the reviewer's judgement, say so and give the reasoning.

Cause — why the gap exists, at the level of process, resourcing or design rather than person. "The progress measurement role was not filled until month four" is a cause. "The cost engineer did not prioritise this" is an accusation, and it converts a finding into a personnel dispute in which the substance is lost.

Consequence — what it costs in decisions, not in abstract risk. "The forecast presented to the September board cannot be traced to a stated method, so its sensitivity to the steelwork productiv-

ity assumption is unknown" tells a sponsor what they are exposed to. "A significant risk to project delivery" tells them nothing.

Correction — the specific act, its owner and its date. If the reviewer cannot name the act the finding is not ready; a correction reading "improve the process" will be closed by an assertion that it has been.

Four rules decide whether the report changes anything. **Rank by consequence, not by severity adjective:** three findings ranked by what they prevent the project from knowing beats fifteen ranked high, medium and low. **Cap the count** — a report with forty findings will have none of them fixed. **Route each finding to the person who can fix it**, because some are not the project's to fix — a corporate coding structure, an unresourced role, a PMO template — and sending those to the project guarantees they stay open. **Track to closure with evidence:** a finding is closed when the reviewer sees the changed artefact.

— 7. How this goes wrong

The compliance sweep sold as assurance. Every required document exists, so the project scores well, and the forecast is still unusable. The tell is a report whose findings are all about documents and none about numbers.

The score without the rubric. Numbers assigned on impression, disputed by the project, defended by seniority. The next review reaches different scores on an unchanged project, and the trend is meaningless.

The composite score as the deliverable. A single percentage travels well and hides everything the review found. §8 shows arithmetically why it cannot represent a blocking weakness.

No sample stated. "We reviewed the cost reports" — how many, which, chosen how, covering what value. Without those, the review cannot be repeated or extended.

Findings written as accusations. Cause stated at the level of a person. The project's response becomes a defence of individuals and the substance is never discussed.

Corrections that cannot be verified. "Strengthen the change control process" — closed at the next review on an assertion, because there is no artefact to inspect.

Too late to matter, and not said. A review at sixty per cent complete reported as though its findings could still be implemented. The project agrees to everything and does none of it.

The review that is never closed out. Findings raised, actions agreed, nobody tracks them, and the next review raises the same findings. Publishing the previous review's closure rate is the cheapest correction available and the most often omitted.

— 8. Worked example

Illustrative figures. One project controls health check at a single review date. Scores are on a 0–4 scale against a published rubric. Weights sum to 100 per cent. All arithmetic shown; percentages rounded to whole numbers where stated.

8.1 The sample and its coverage

The project has **140 control accounts** with a total budget of **9,500,000** (USD, illustrative). The review inspected the **12 largest by budget**, carrying **6,200,000**.

coverage by count = $12 \div 140 = 0.0857 \rightarrow 8.6 \%$
 coverage by value = $6,200,000 \div 9,500,000 = 0.6526 \rightarrow 65.3 \%$

Both belong in the report. Count alone reads as a thin review; value alone conceals that 128 accounts — every small account, where control weaknesses concentrate — were not examined. State the sampling rule too, because its bias is knowable: "the twelve largest by budget" systematically excludes the small, unglamorous accounts, and a future review should sample differently for that reason.

8.2 The scores

#	Element	Weight	Score (0-4)	Weighted contribution
1	Baseline integrity and change control	20 %	3	$0.20 \times 3 = 0.60$
2	Schedule quality	20 %	2	$0.20 \times 2 = 0.40$
3	Cost and commitment capture	15 %	3	$0.15 \times 3 = 0.45$
4	Progress measurement	15 %	1	$0.15 \times 1 = 0.15$
5	Forecasting	15 %	2	$0.15 \times 2 = 0.30$
6	Risk and change management	15 %	2	$0.15 \times 2 = 0.30$
Total		100 %		2.20

weight check = $20 + 20 + 15 + 15 + 15 + 15 = 100 \%$
 composite = $0.60 + 0.40 + 0.45 + 0.15 + 0.30 + 0.30 = 2.20$
 as a percentage of the 4.00 maximum = $2.20 \div 4.00 = 0.550 \rightarrow 55 \%$

8.3 Why 55 per cent is not the finding

Element 4, progress measurement, scored **1**: rules of credit exist as a document but progress is claimed by percentage judgement, with no independent verification in the last three periods.

Every subsequent number on this project is computed from that element. Earned value derives from measured progress; the cost and schedule performance indices derive from earned value; the forecast derives from those indices. A score of 1 on element 4 does not mean "one element is weak" — it means the project's cost performance measurement, its schedule performance measurement and its forecast all rest on unverified percentage judgements.

The composite cannot express this, because a weighted average treats the elements as independent and additive. They are neither.

8.4 Testing whether a dependency rule rescues the composite

A common remedy is a rule: **no element may be scored above the lowest score of the elements it depends on**. Forecasting (element 5) depends on progress measurement (element 4, scored 1) and cost capture (element 3, scored 3), so it is capped at 1 and falls from 2 to 1.

```
revised composite = 0.60 + 0.40 + 0.45 + 0.15 + (0.15 × 1) + 0.30
                  = 0.60 + 0.40 + 0.45 + 0.15 + 0.15 + 0.30
                  = 2.05

as a percentage   = 2.05 ÷ 4.00 = 0.5125 → 51 %
```

The composite moves from 55 per cent to 51 per cent — four percentage points to represent a project whose entire performance measurement chain is unverified.

That is the point of the calculation, and it is not a failure of the rule. A single weighted average is structurally incapable of expressing a blocking weakness, because averaging compensates: strength elsewhere always offsets weakness here. Improving the rule cannot fix a property of the arithmetic.

8.5 What the deliverable should therefore be

Three things, in this order.

The blocking finding, stated first. Element 4 scored 1. Earned value, both performance indices and the forecast presented to the board derive from unverified percentage judgements. Until independent verification is in place on a defined sample, no performance figure from this project should be relied on for a funding decision.

The element profile, as a labelled chart or table with the rubric descriptor next to each score — never colour alone, never a single dial.

The composite, last, and only as a trend line for this project against itself. Comparing 55 per cent here with 55 per cent elsewhere is meaningless unless both used the same rubric, weights and a comparable sample, which is rarely true and almost never checked.

8.6 Closure of the previous review

The previous review raised **14 findings**. At this review: **5** closed with evidence, **6** open with revised dates, **3** open with no movement.

```
check: 5 + 6 + 3 = 14
closure rate = 5 ÷ 14 = 0.357 → 35.7 %
```

A closure rate of 35.7 per cent with three findings showing no movement is itself a finding, and it belongs ahead of any new observation — because a project that does not close findings will not close the ones raised today either, and that changes what this review can reasonably expect to achieve.

— 9. Checklist

Commissioning the review

- [] The question stated as fitness for decisions, not conformance — or, if an audit is wanted, called one.
- [] Rubric published before the review, with observable descriptors for every score.
- [] Weights published and summing to 100 per cent.
- [] Sampling rule agreed in advance, coverage to be reported by count and by value.
- [] Independence analysed against self-review, reporting line, familiarity and advocacy; any impairment declared in the report.
- [] Access agreed: source systems, not extracts; people, not only managers.
- [] Timing justified — before baseline, first cycle, gate or trigger — and the purpose stated accordingly.
- [] Previous review's findings retrieved and their closure status tested first.

Conducting it

- [] The chain walked in order: structure, baseline, cost capture, progress, forecast, risk and change, reporting.
- [] Each score evidenced by artefact, date and test applied.
- [] Factual conditions agreed with the project before cause and consequence are written.
- [] Access limitations recorded as findings.
- [] Findings capped, ranked by consequence, each routed to the person able to fix it.

Reporting it

- [] Blocking finding stated first, before any score.
- [] Element profile published with descriptors; no meaning carried by colour alone.
- [] Composite reported last, and only as a trend against this project's own history.
- [] Every finding carries condition, criterion, cause, consequence and correction with an owner and date.
- [] No correction written that cannot be verified by inspecting a changed artefact.
- [] Closure rate from the previous review published.

A review run this way produces fewer findings, and the project argues with almost none of them — because the facts were agreed before the conclusions were written and every correction names something specific. The measure of whether it worked is not the score. It is whether the next review can start somewhere new.

— Related

- [TPL-15 – Project controls health check](#) — the review instrument, rubric structure and scoring sheet
- [BPG-05 – Schedule quality – a practical review](#) — the schedule element of the chain, in depth
- [TPL-14 – Schedule quality review checklist](#) — the schedule checks in usable form

- [BPG-04 – Baselineing and baseline change control](#) — the baseline integrity element and the pre-baseline review
- [BPG-14 – Monthly reporting that gets read](#) — what "reaches the decision" means
- [BPG-16 – Risk registers that work](#) — the register health metrics a reviewer can compute in minutes

— Sources and standards

- PCL-AI Body of Knowledge ([docs/bok/](#)), Domain 10 — Project Scheduling, first authored draft, August 2026: schedule health-check metrics and the schedule basis document.
- PCL-AI Body of Knowledge ([docs/bok/](#)), Domain 8 — Project Management Lifecycle, first authored draft, August 2026: stage-gate design, gate conditions tracked to closure, and the PMO's assurance role.
- PCL-AI Body of Knowledge ([docs/bok/](#)), Domain 4 — Performance Management, Variance Analysis and Management Reporting, first authored draft, August 2026: measure design, gaming patterns and auditing the measure rather than only the result.
- PCI Canonical Facts ([docs/publication-framework/00-framework/CANONICAL-FACTS.md](#)), verified August 2026: naming, status and claims policy.

No maturity model, assurance framework or published review methodology is reproduced or cited here, and no scoring threshold is attributed to any external body. The weights, scale and rubric in §8 are illustrative; an organisation adopting them should set its own and publish them before its first review.

— Status and version

Founding-stage document · Version 1.0 · draft · Reviewed under PCI governance. PCI makes no claims of accreditation or recognition beyond what is true today.